



# AIRFLOW - DATA FLOW ENGINE FROM AIRBNB

By Walter Liu 2016/01/28

# Who am I?


- ☒ Archer Architect of TrendMicro Coretech backend
- ☐ A new user in Airflow




**TREND**  
**M I C R O™**

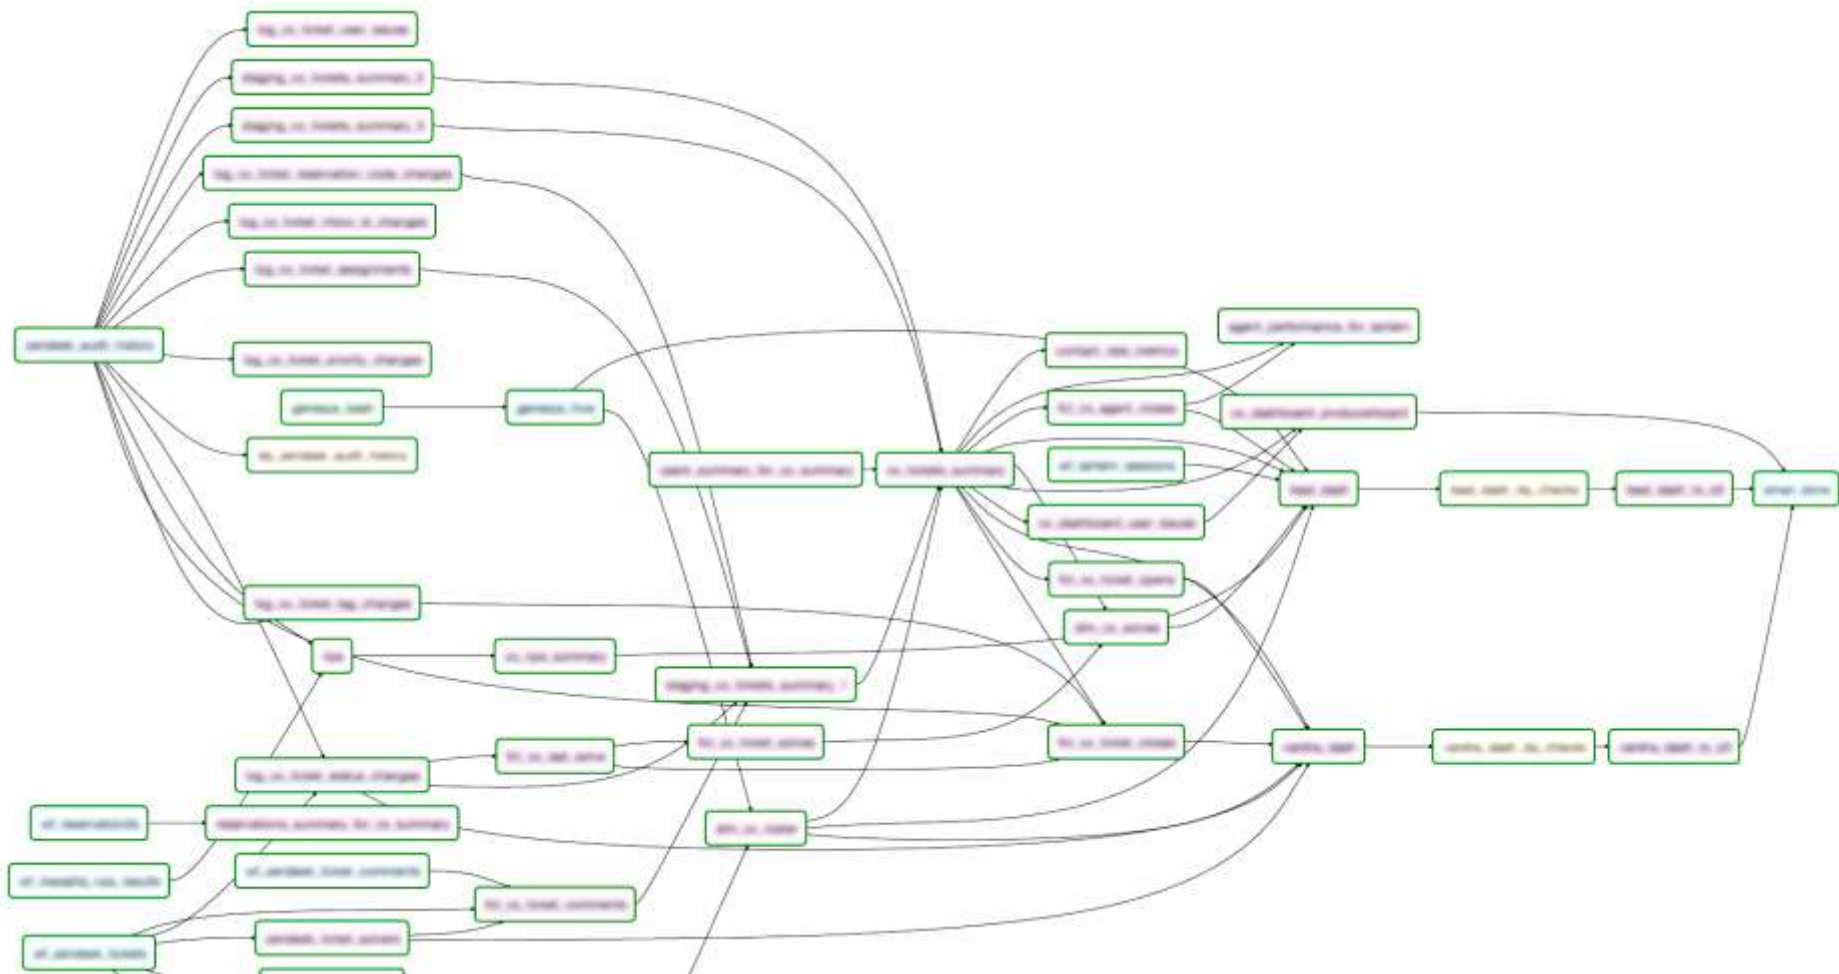
# **Why Data Flow Engine?**

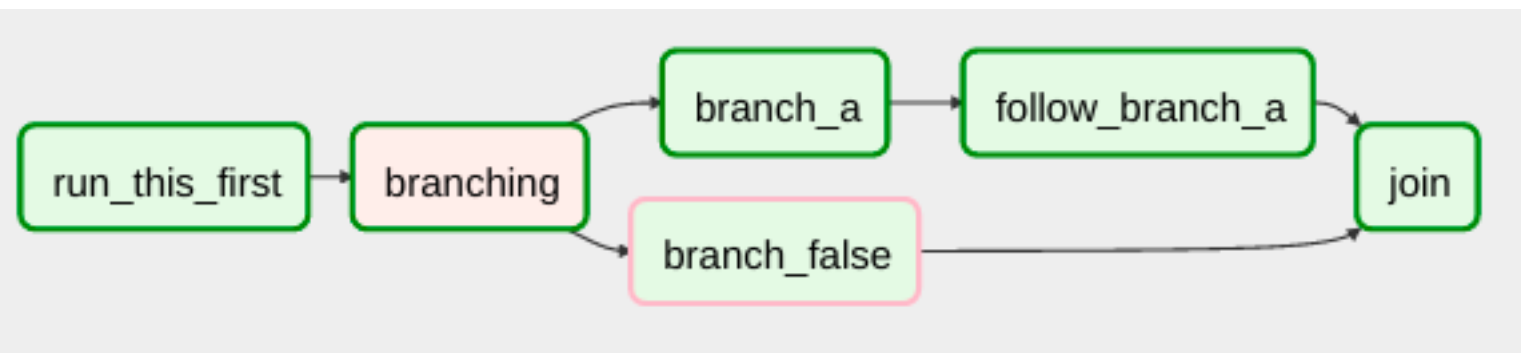
# Cron Job



# Direct Acyclic Graph (DAG)







# Data relationships

---

- Data availability
  - ▣ if the data is not there, trigger the process to generate the data.
- Data dependency
  - ▣ Some data relies on other data to generate.



# Operability

---

- ❑ Job failed and resume
- ❑ Job monitor
- ❑ Backfill

# Airflow



**DAG structure as code**

```
default_args = {
    'email': ['airflow@airflow.com'],
    'email_on_failure': False,
    'email_on_retry': False,
    'retries': 1,
    'retry_delay': timedelta(minutes=5),
}
```

```
dag = DAG('tutorial', default_args=default_args)
```

```
# t1, t2 and t3 are examples of tasks created by instantiating operators
```

```
t1 = BashOperator(task_id='print_date', bash_command='date', dag=dag)
```

```
t2 = BashOperator(task_id='sleep', bash_command='sleep 5', retries=3, dag=dag)
```

```
templated_command = """
```

```
    {% for i in range(5) %}
        echo "{{ ds }}"
        echo "{{ macros.ds_add(ds, 7)}}"
        echo "{{ params.my_param }}"
    {% endfor %}
```

```
"""
```

```
t3 = BashOperator(
    task_id='templated',
    bash_command=templated_command,
    params={'my_param': 'Parameter I passed in'},
    dag=dag)
```

```
t2.set_upstream(t1)
```

```
t3.set_upstream(t1)
```

```
default_args = {
    'owner': 'airflow',
    'depends_on_past': False,
    'start_date': datetime(2015, 6, 1),
    'email': ['airflow@airflow.com'],
    'email_on_failure': False,
    'email_on_retry': False,
    'retries': 1,
    'retry_delay': timedelta(minutes=5),
    # 'queue': 'bash_queue',
    # 'pool': 'backfill',
    # 'priority_weight': 10,
    # 'end_date': datetime(2016, 1, 1),
}
```

```
templated_command = ""
    {% for i in range(5) %}
        echo "{{ ds }}"
        echo "{{ macros.ds_add(ds, 7) }}"
        echo "{{ params.my_param }}"
    {% endfor %}
""
```

```
{{ ds }} => now YYYY-MM-DD
{{ yesterday_ds }} => yesterday YYYY-MM-DD
{{ tomorrow_ds }} => tomorrow YYYY-MM-DD
...
```

```
dag = DAG('tutorial', default_args=default_args)
```

```
# t1, t2 and t3 are examples of tasks created by instatiating operators
```

```
t1 = BashOperator(  
    task_id='print_date',  
    bash_command='date',  
    dag=dag)
```

```
t2 = BashOperator(  
    task_id='sleep',  
    bash_command='sleep 5',  
    retries=3,  
    dag=dag)
```

```
t3 = BashOperator(  
    task_id='templated',  
    bash_command=templated_command,  
    params={'my_param': 'Parameter I passed in'},  
    dag=dag)
```

```
t2.set_upstream(t1)
```

```
t3.set_upstream(t1)
```


# Demo

- ❑ python tutorial.py
- ❑ airflow list\_dags
- ❑ airflow list\_tasks tutorial
- ❑ airflow list\_tasks tutorial --tree
- ❑ airflow test tutorial print\_date 2015-06-01
- ❑ airflow test tutorial sleep 2015-06-01
- ❑ airflow run tutorial templated 2015-06-01
- ❑ Backfill: airflow backfill tutorial -s 2015-06-07 -e 2015-06-10
  - ▣ Run again
  - ▣ Run another date range

Site: <http://localhost:8080/admin/>



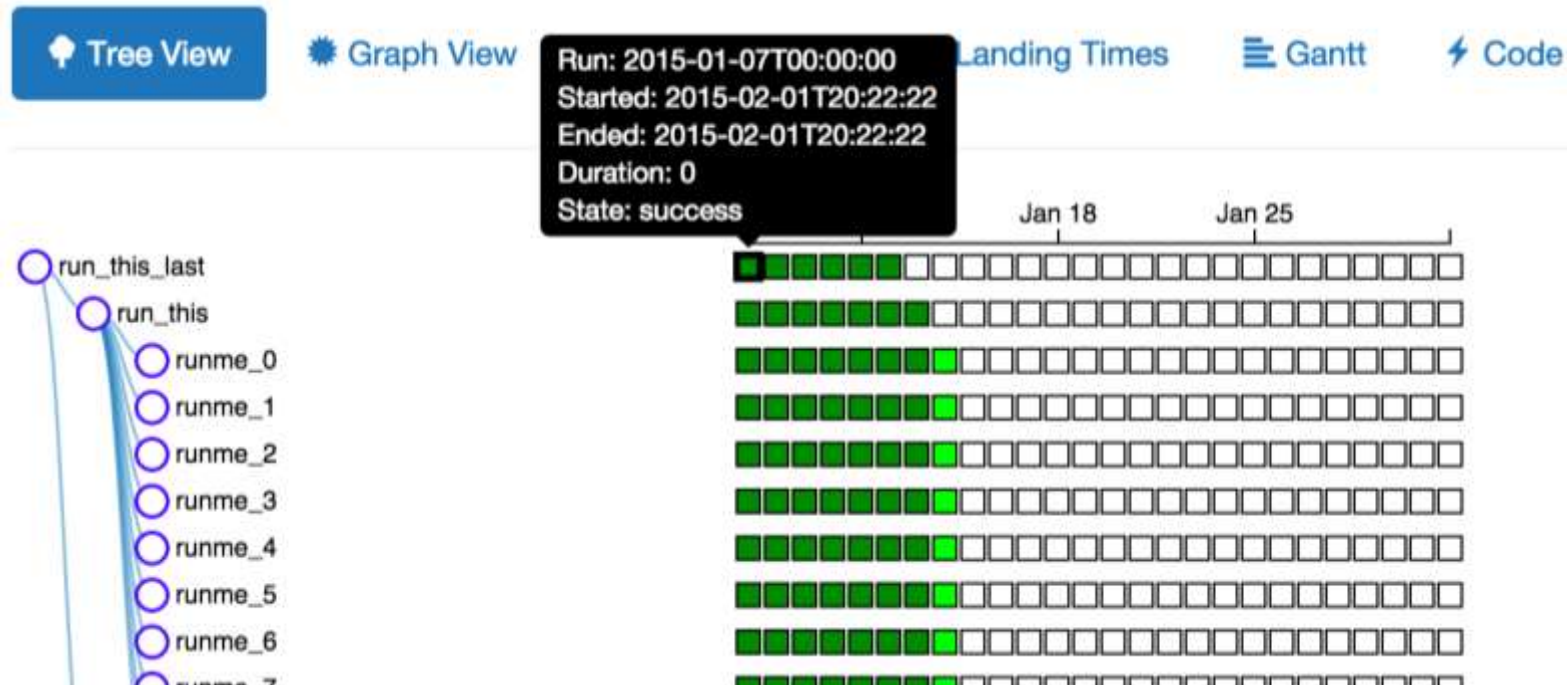
# UI – DAG view



## DAGs

DAG	Filepath	Owner	Task by State	Links
example1	example_dags/example1.py	airflow	<div><div>80</div><div>1</div><div>0</div></div>	      
example2	example_dags/example2.py	airflow	<div><div>128</div><div>10</div><div>0</div></div>	      
example3	example_dags/example3.py	airflow	<div><div>138</div><div>5</div><div>0</div></div>	      

# UI – Tree View






# UI - Gantt

Run: 2015-02-01 00:00:00


Go



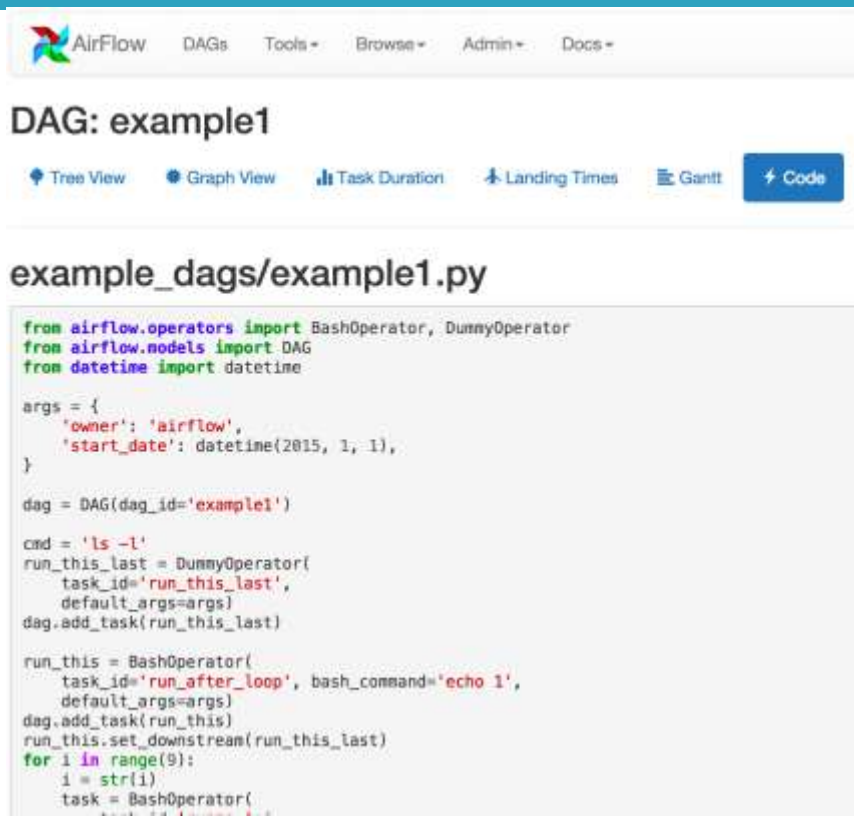
# UI – Task duration

DAG: core\_cx

Tree View Graph View Task Duration Landing Times Gantt Code



# UI – Code



The screenshot displays the Apache AirFlow web interface. At the top, there is a navigation bar with the AirFlow logo and links for DAGs, Tools, Browse, Admin, and Docs. Below this, the title "DAG: example1" is shown. A row of view options includes Tree View, Graph View, Task Duration, Landing Times, Gantt, and a Code button. The main content area displays the Python code for "example\_dags/example1.py".

```
from airflow.operators import BashOperator, DummyOperator
from airflow.models import DAG
from datetime import datetime

args = {
    'owner': 'airflow',
    'start_date': datetime(2015, 1, 1),
}

dag = DAG(dag_id='example1')

cmd = 'ls -l'
run_this_last = DummyOperator(
    task_id='run_this_last',
    default_args=args)
dag.add_task(run_this_last)

run_this = BashOperator(
    task_id='run_after_loop', bash_command='echo 1',
    default_args=args)
dag.add_task(run_this)
run_this.set_downstream(run_this_last)
for i in range(9):
    i = str(i)
    task = BashOperator(
```

# Airflow - Pros

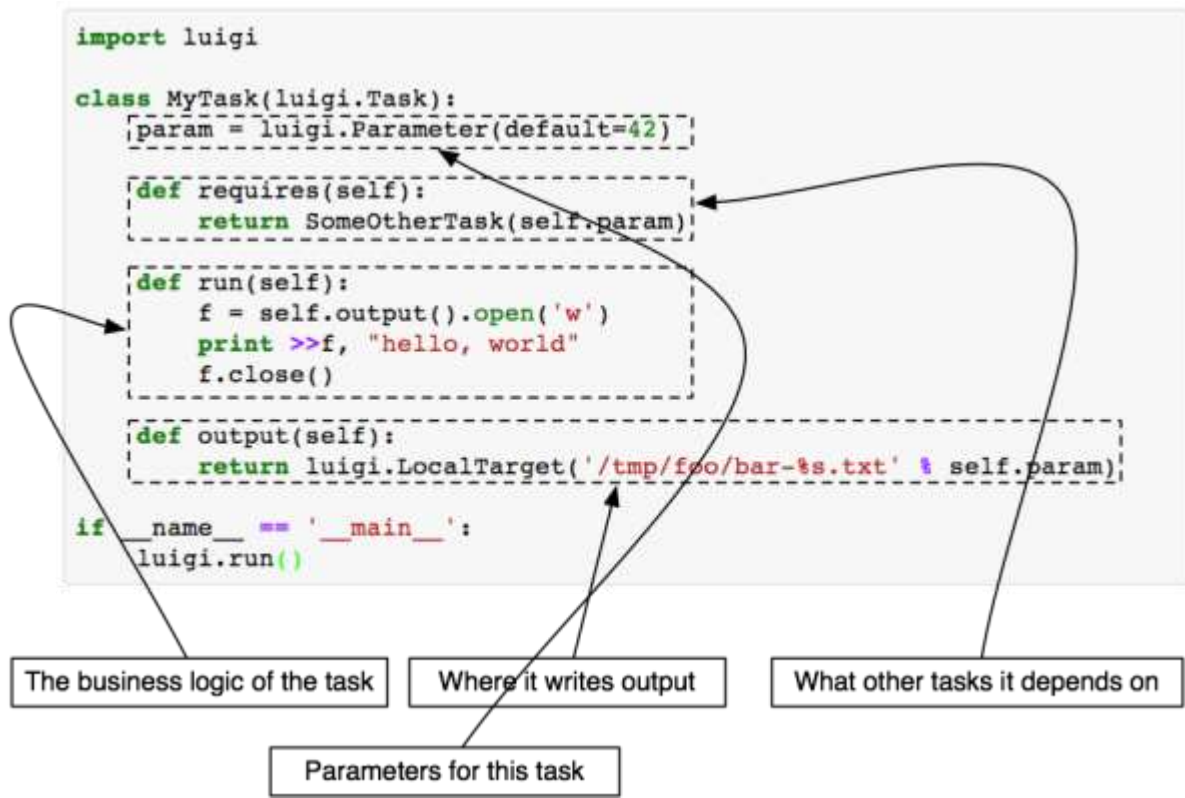
- ❑ Dynamic generating path
- ❑ Have both Time scheduler and Command line trigger
- ❑ Has Master/Worker model (automatically distribute tasks)
- ❑ Scale if you have many tasks in a chain.
  - ▣ But not be so useful to most of our tasks. Maybe useful for full dump.
- ❑ Fancy UI
  - ▣ Dependencies of tasks
  - ▣ Task success/failure
  - ▣ Scheduled tasks status
- ❑ Has utility lib to wait\_data for S3, Hadoop ...

# Airflow - Cons

- ❑ Additional DB/Redis or Rabbitmq for Celery
  - ▣ HA design: Use RDBMS/redis-cache in AWS
- ❑ Require python 2.7 and many other libraries.
- ❑ Not dependent on data. Just task dependency.  
(Not big cons)
  - ▣ Write check data file code.



# A snippet of the task of Luigi



# Recap

- Solving DAG job management problem
- Features to improve daily operation
  - ▣ Monitor/Visualization
  - ▣ Job failure and resume
  - ▣ Backfill



Backup slides

# Other Selections?

- ❑ Oozie (on Hadoop)
- ❑ Luigi (by Spotify)
  - ▣ Mario (Luigi in Scala)
  - ▣ Ruigi (Luigi in R)
- ❑ Airflow (by Airbnb)
- ❑ Mistral (by OpenStack)
- ❑ Pinball (by Pinterest)
- ❑ ...

# Github statistics

	<b>Start Date</b>	<b>Star</b>	<b>Commits in recent 2 months</b>
Airflow	2014/10/05	1516	362
Luigi	2011/11/13	3777	105
Pinball	2015/05/01	476	0
Oozie	2011/08/28	167	14

# Oozie

## □ Pros

- ▣ Mature
- ▣ Native support of Hadoop ECO system
- ▣ Python is not required.

## □ Cons

- ▣ Big and complex XML to define the flow and config.
- ▣ Control flow is somehow restrictive
- ▣ Hadoop ECO system is required.
  - Unix box, Java, Hadoop, Pig, ExtJS library

# Oozie XML Example

```
<!--
Copyright (c) 2011 NAVTEQ Inc. All rights reserved.
NGMB IPS ingestor Oozie Script
-->
<workflow-app xmlns='uri:oozie:workflow:0.1' name='NGMB-IPS-ingestion'>
  <start to='ingestor' />
  <action name='ingestor'>
    <java>
      <job-tracker>${jobTracker}</job-tracker>
      <name-node>${nameNode}</name-node>
      <configuration>
        <property>
          <name>mapred.job.queue.name</name>
          <value>default</value>
        </property>
      </configuration>
      <main-class>com.navteq.assetmgmt.MapReduce.ips.IPSLoader</main-cla
      <java-opts>-Xmx2048m</java-opts>
      <arg>${driveID}</arg>
    </java>
    <ok to="merging"/>
    <error to="fail"/>
  </action>
  <fork name="merging">
    <path start="mergeLidar"/>
    <path start="mergeSignage"/>
  </fork>
  <action name='mergeLidar'>
    <java>
```

```
      <name-node>${nameNode}</name-node>
      <configuration>
        <property>
          <name>mapred.job.queue.name</name>
          <value>default</value>
        </property>
      </configuration>
      <main-class>com.navteq.assetmgmt.hdfs.merge.MergerLoader</main-clas
      <java-opts>-Xmx2048m</java-opts>
      <arg>-drive</arg>
      <arg>${driveID}</arg>
      <arg>-type</arg>
      <arg>lidar</arg>
      <arg>-chunk</arg>
      <arg>${lidarChunk}</arg>
    </java>
    <ok to="completed"/>
    <error to="fail"/>
  </action>
  <action name='mergeSignage'>
    <java>
      <job-tracker>${jobTracker}</job-tracker>
      <name-node>${nameNode}</name-node>
      <configuration>
        <property>
          <name>mapred.job.queue.name</name>
          <value>default</value>
        </property>
      </configuration>
      <main-class>com.navteq.assetmgmt.hdfs.merge.MergerLoader</main-clas
      <java-opts>-Xmx2048m</java-opts>
      <arg>-drive</arg>
```

# Luigi - Pros

- ❑ Makefile like, dependencies on data (Upward instead of downward)
- ❑ Command line trigger
- ❑ Not only Hadoop
- ❑ Support target: HDFS/S3/RDBMS/.... (not limited)
  - ▣ Write your own code to support Casandra, etc.
- ❑ Dependencies are decentralized, no big XML.
- ❑ UI
- ❑ Dynamic dependency/task is supported. (Don't know it is better than airflow or not)
- ❑ [util] Date algebra



# Luigi - Cons

- ❑ No built-in triggering
  - ▣ No crontab like things (could borrow Chronos)
  - ▣ Use cron to run tasks on specific nodes.  
(normally)

# Luigi - Notes

- ❑ Local execution
  - ▣ Pros: Easy to debug
  - ▣ Cons:
    - need access to other system. (Well, other system didn't achieve this either.)
- ❑ No scalability if you have many tasks in a chain.
- ❑ Centralized scheduler servers (optional, recommended for production)
  - ▣ Make sure two instances of the same task are not running simultaneously
  - ▣ Provide visualization of everything that's going on.
- ❑ HA option
  - ▣ Run 2+ tasks on 2+ machines at the same time

# References

---

- <http://erikbern.com/2015/07/02/more-luigi-alternatives/>