



# Apache Airflow @ Agari

Sid Anand (@r39132)

Apache Airflow Meet-up

June 2016



# Agari

What We Do!

# Agari : What We Do



# Agari : What We Do



# Agari : What We Do



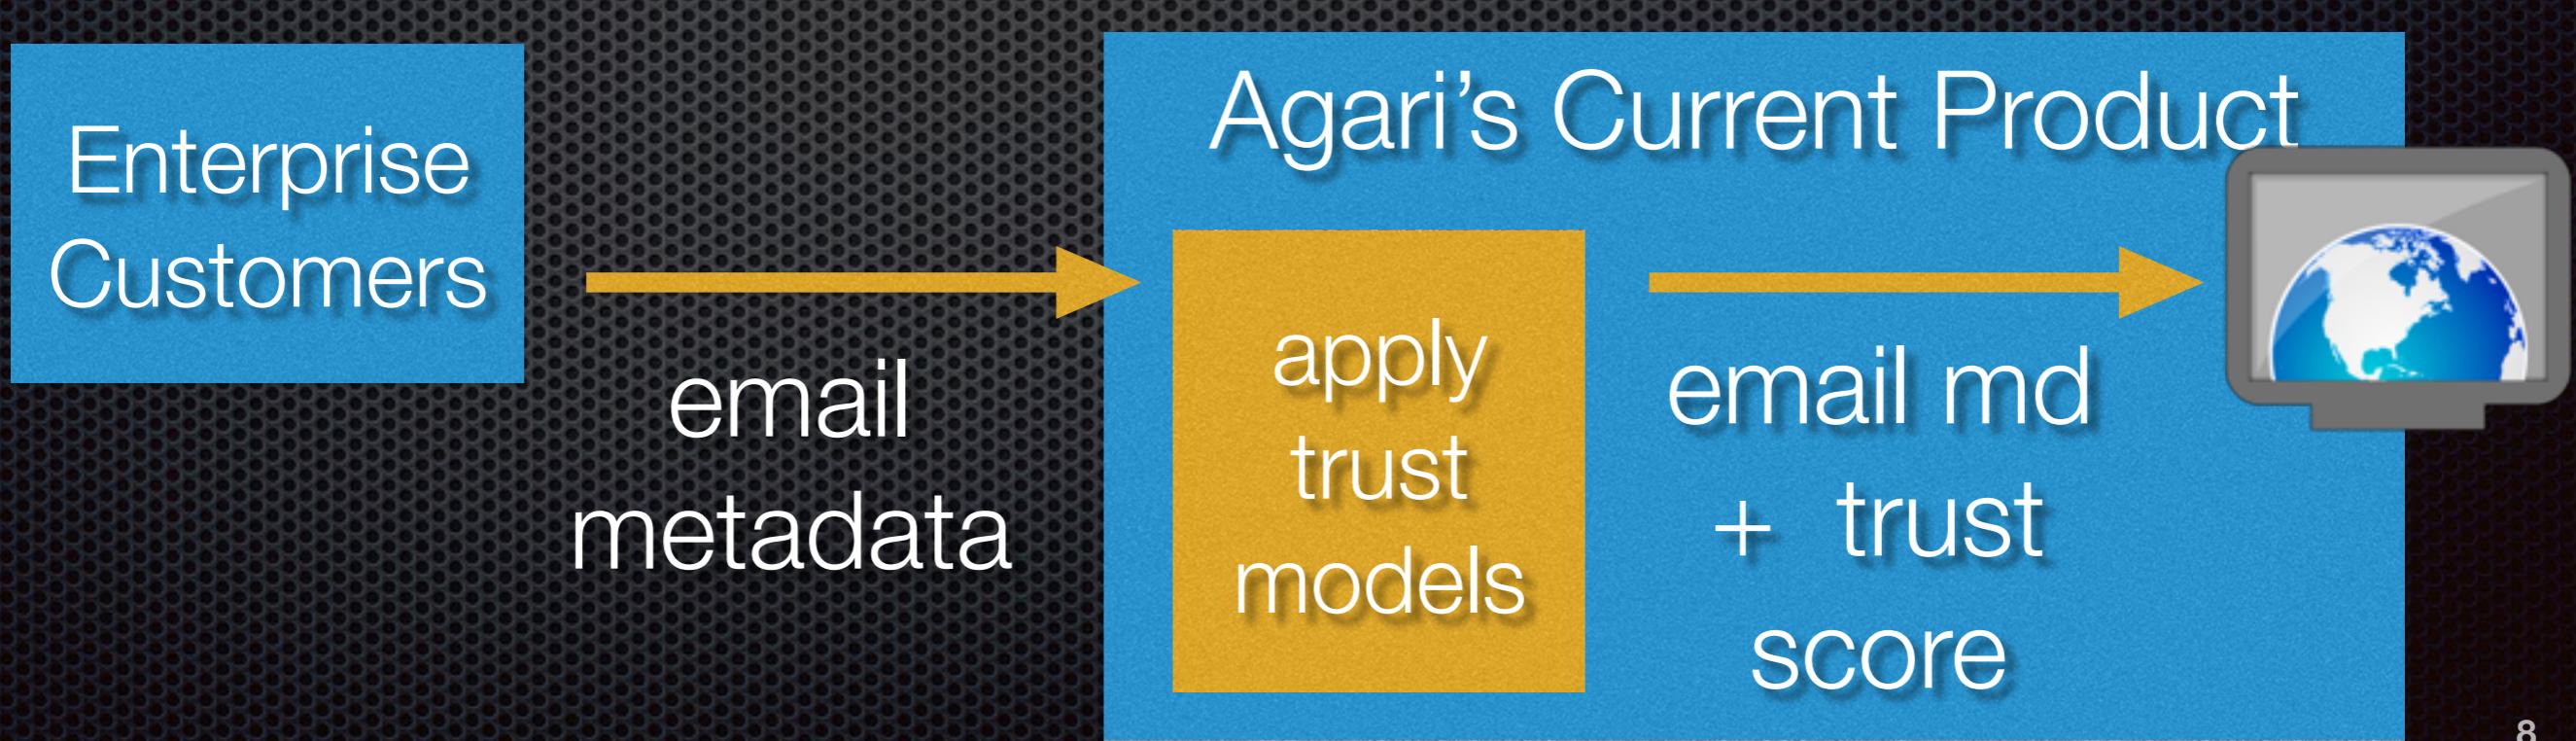
# Agari : What We Do



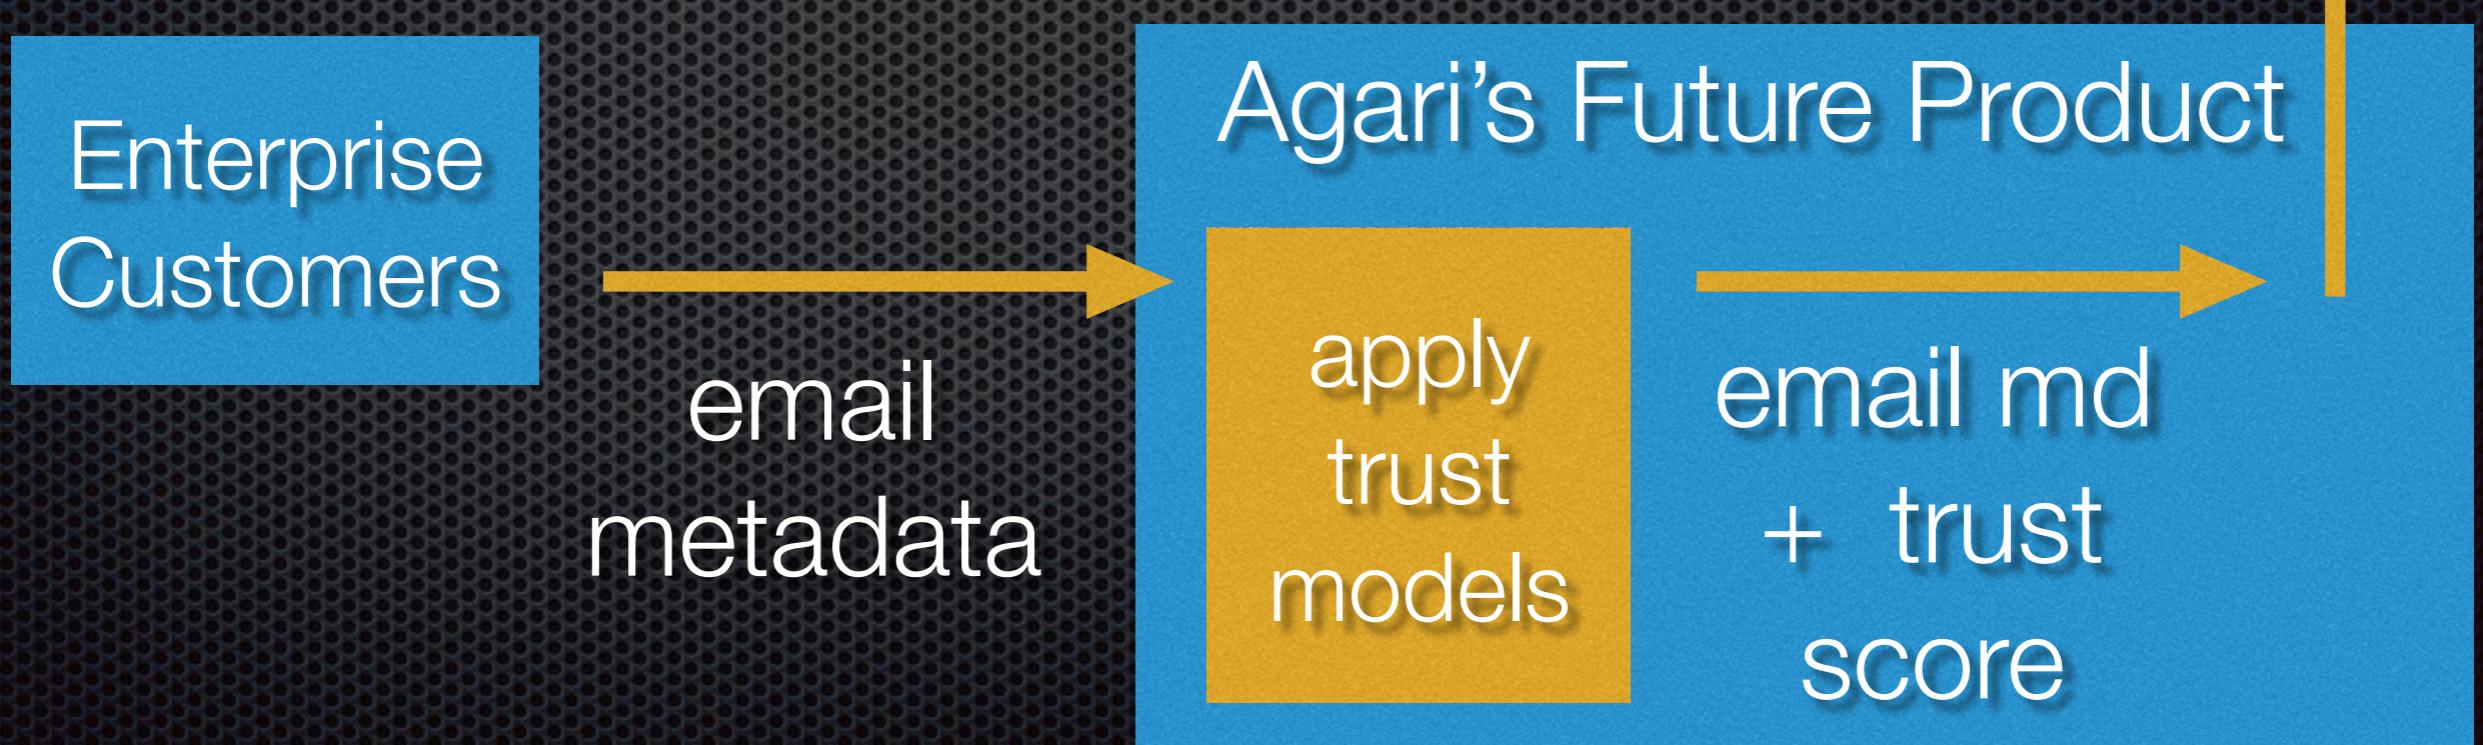
# Agari : What We Do



# Agari : What We Do



# Agari : What We Do



# Apache Airflow @ Agari

## How Do We Use It?

# 2 Classes of Orchestration

**New Product**  
(Enterprise Protect)

build trust  
models

apply trust  
models  
(message  
scoring)

**Operational Automation**

cron++  
(general  
job  
scheduler)

# 2 Classes of Orchestration

**New Product**  
(Enterprise Protect)

build trust  
models

apply trust  
models  
(message  
scoring)

**Operational Automation**

cron++  
(general  
job  
scheduler)

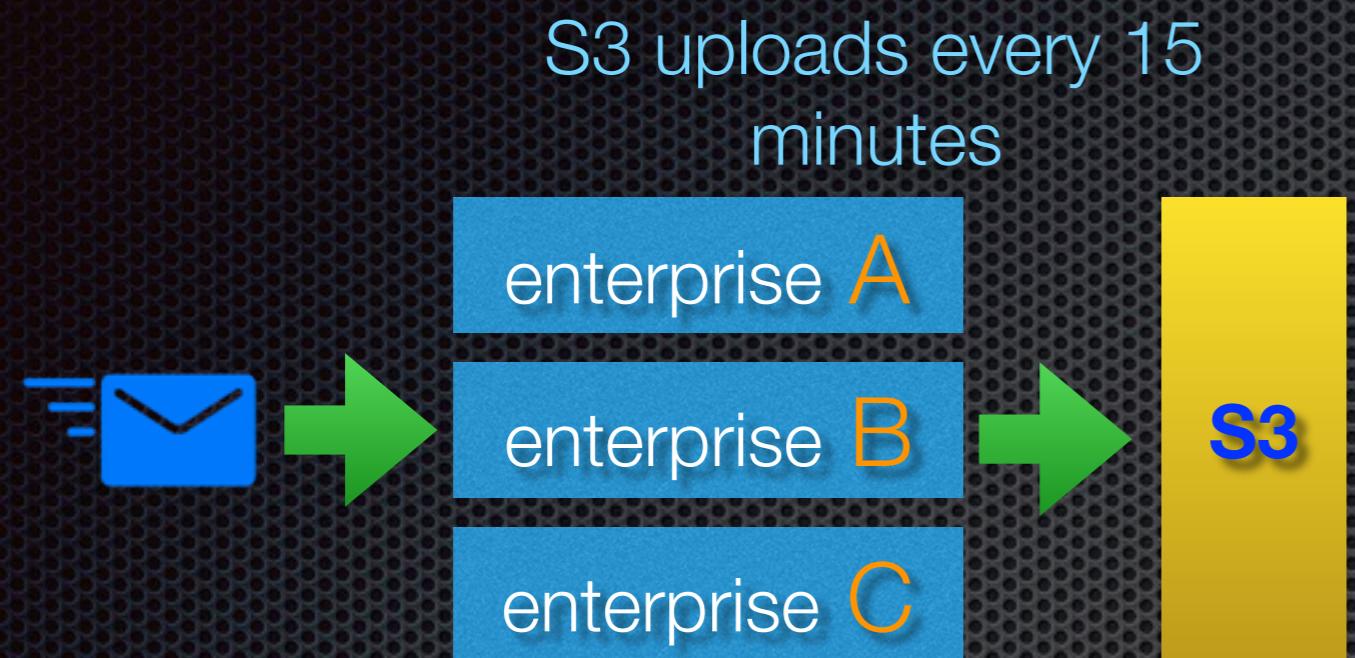
This Talk



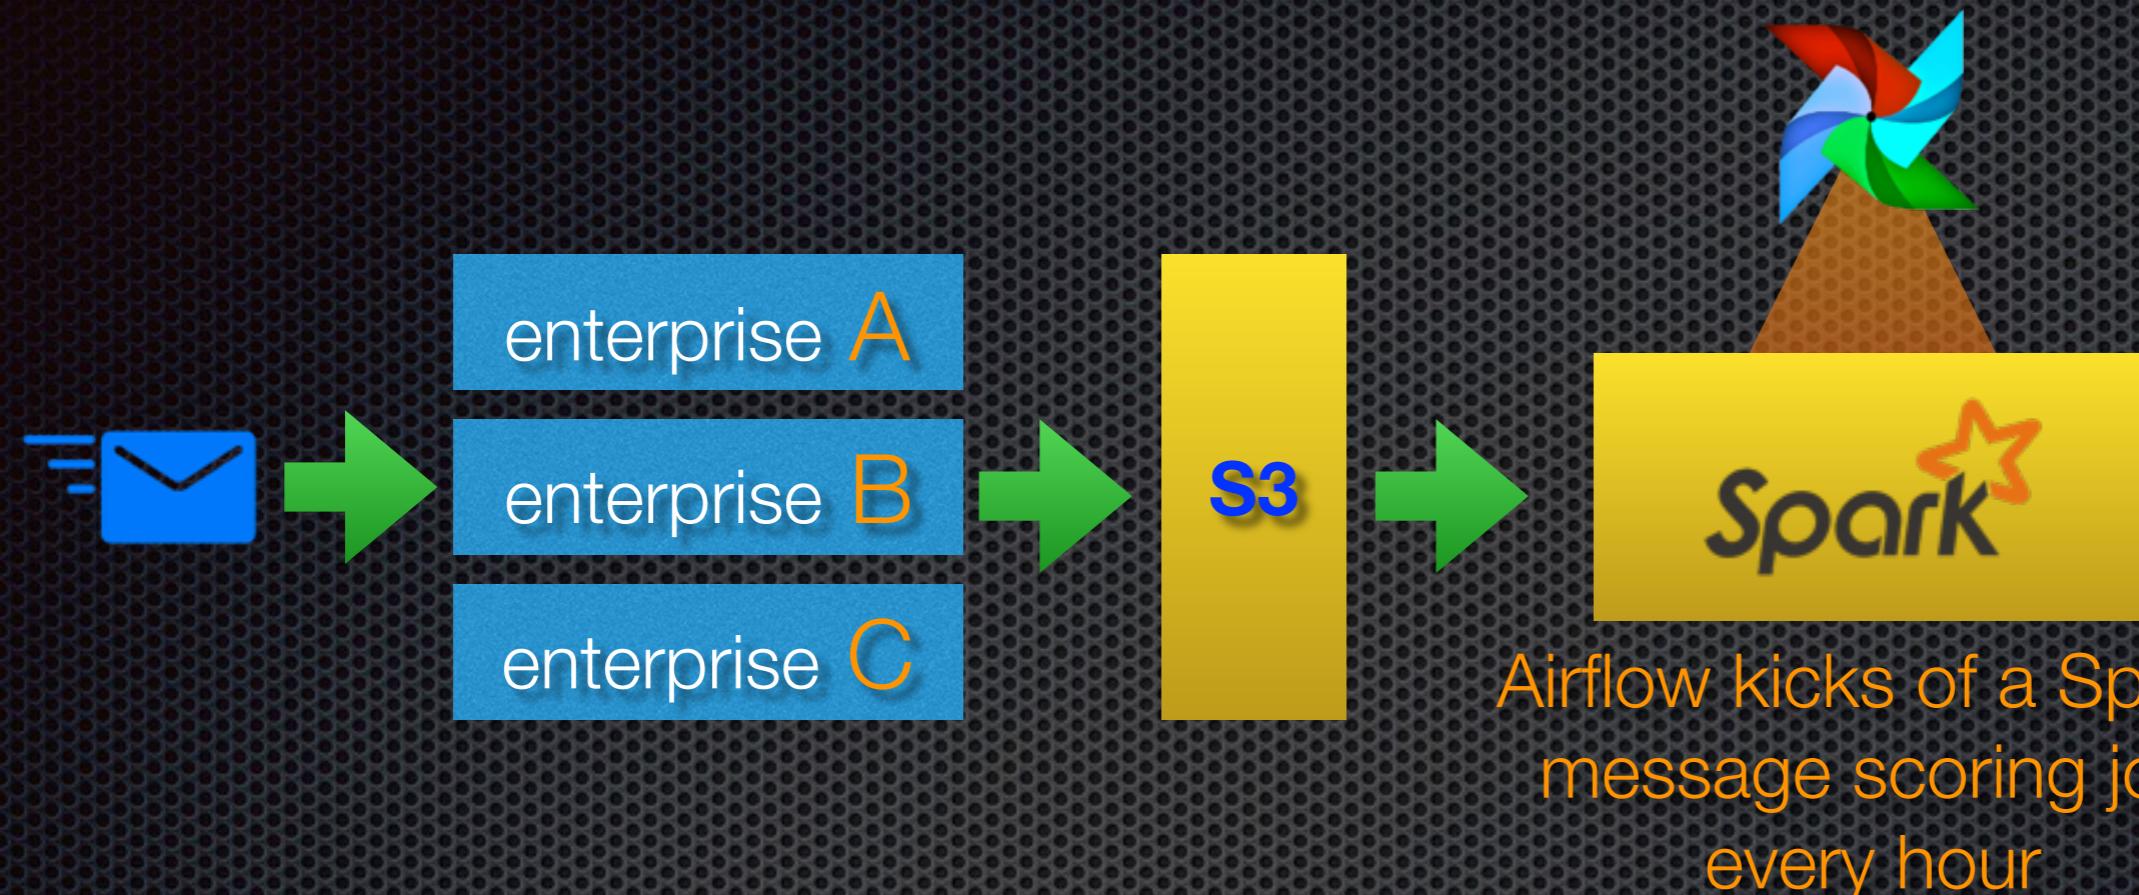
# Use-Case : Message Scoring

Batch Pipeline Architecture

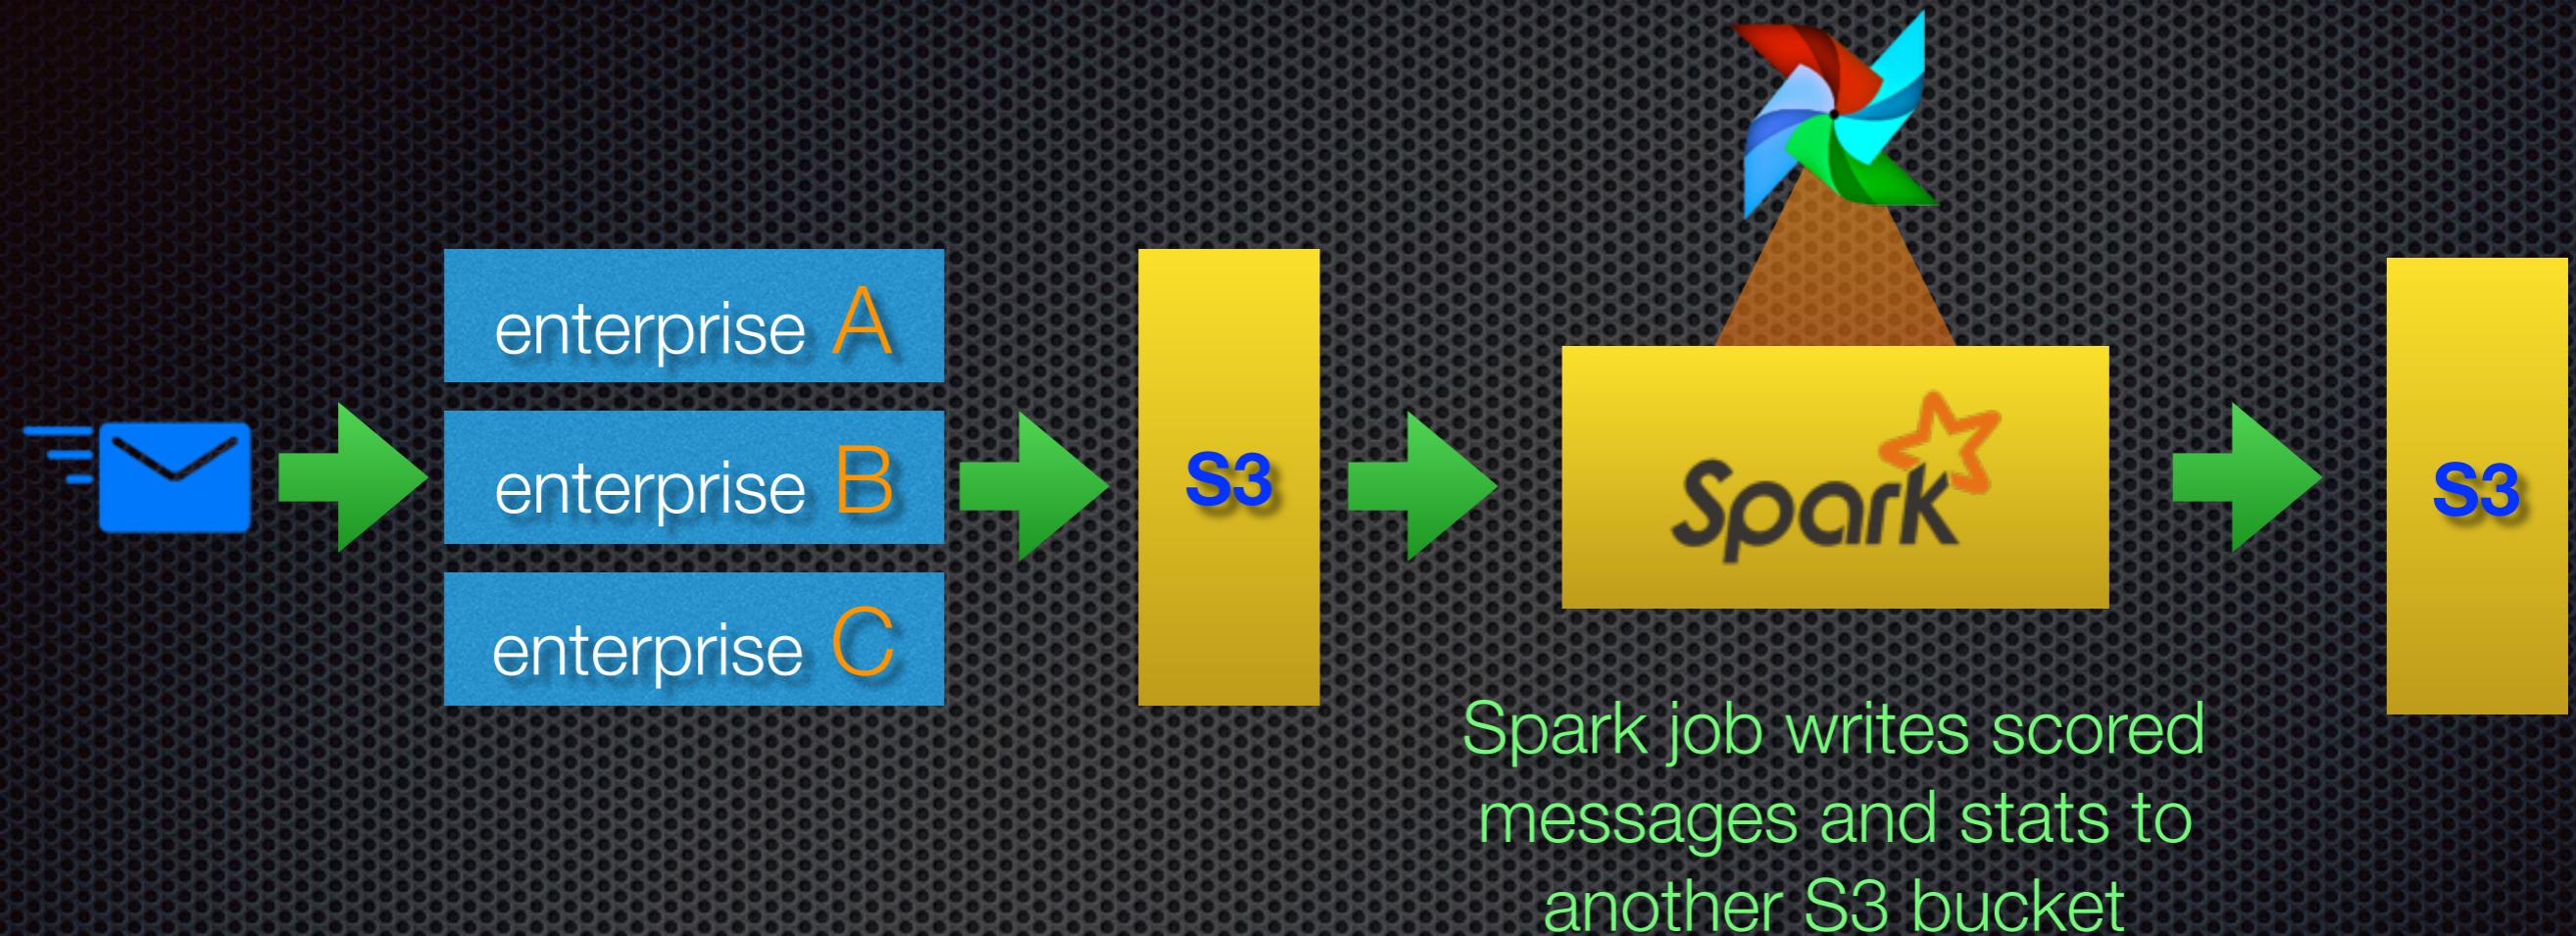
# Use-Case : Message Scoring



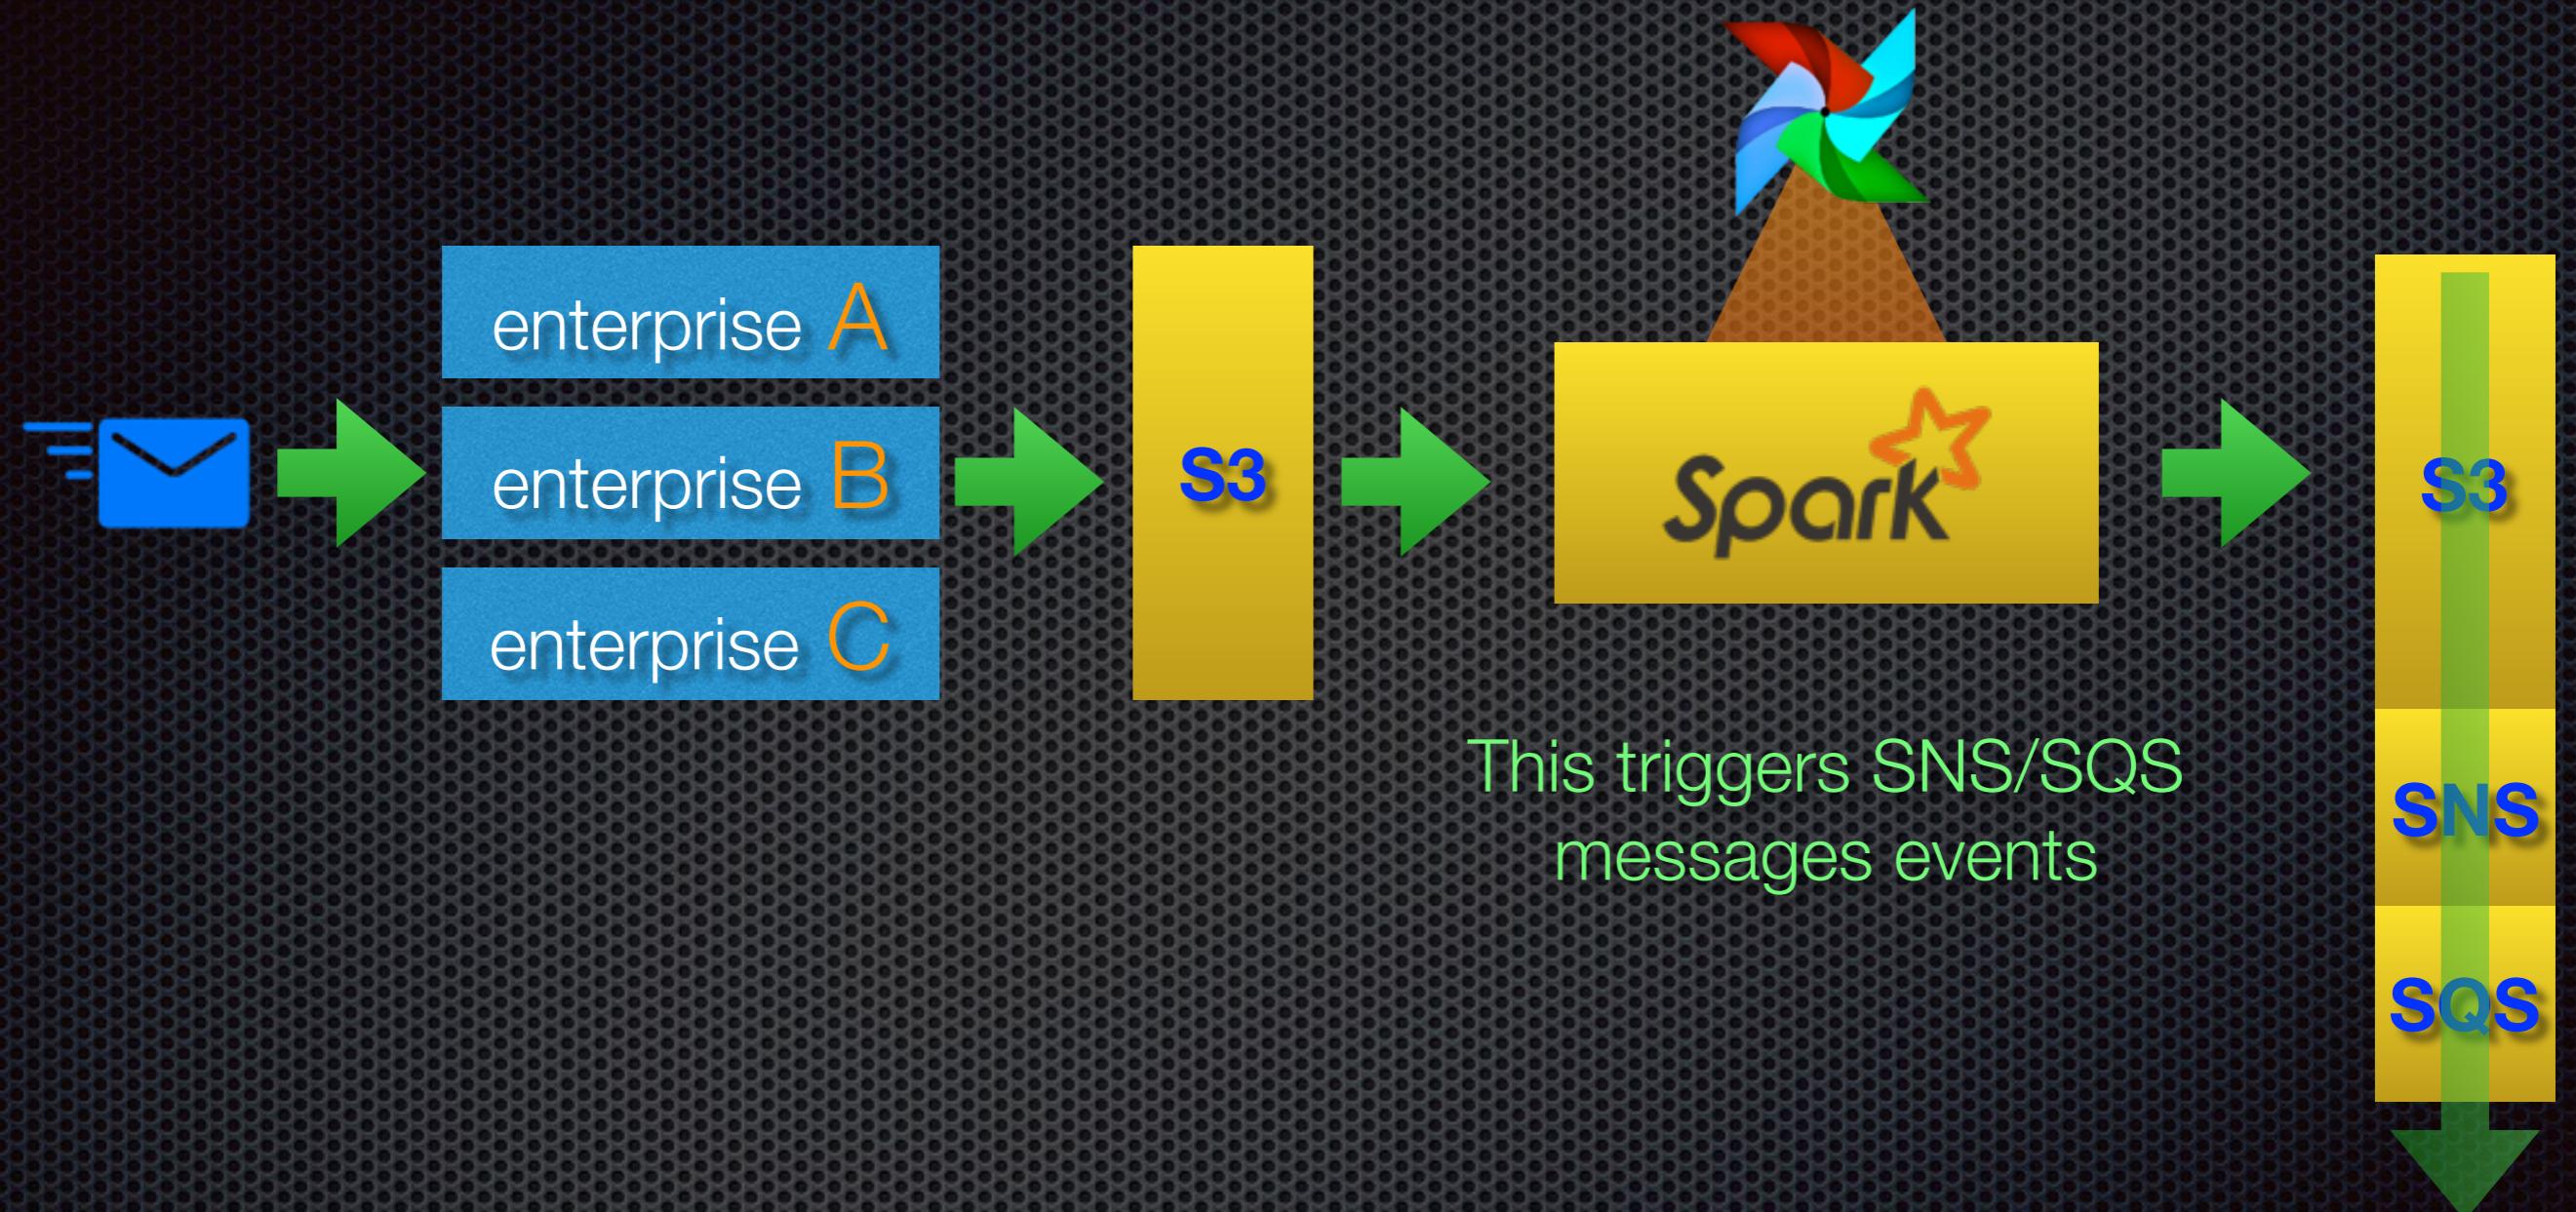
# Use-Case : Message Scoring



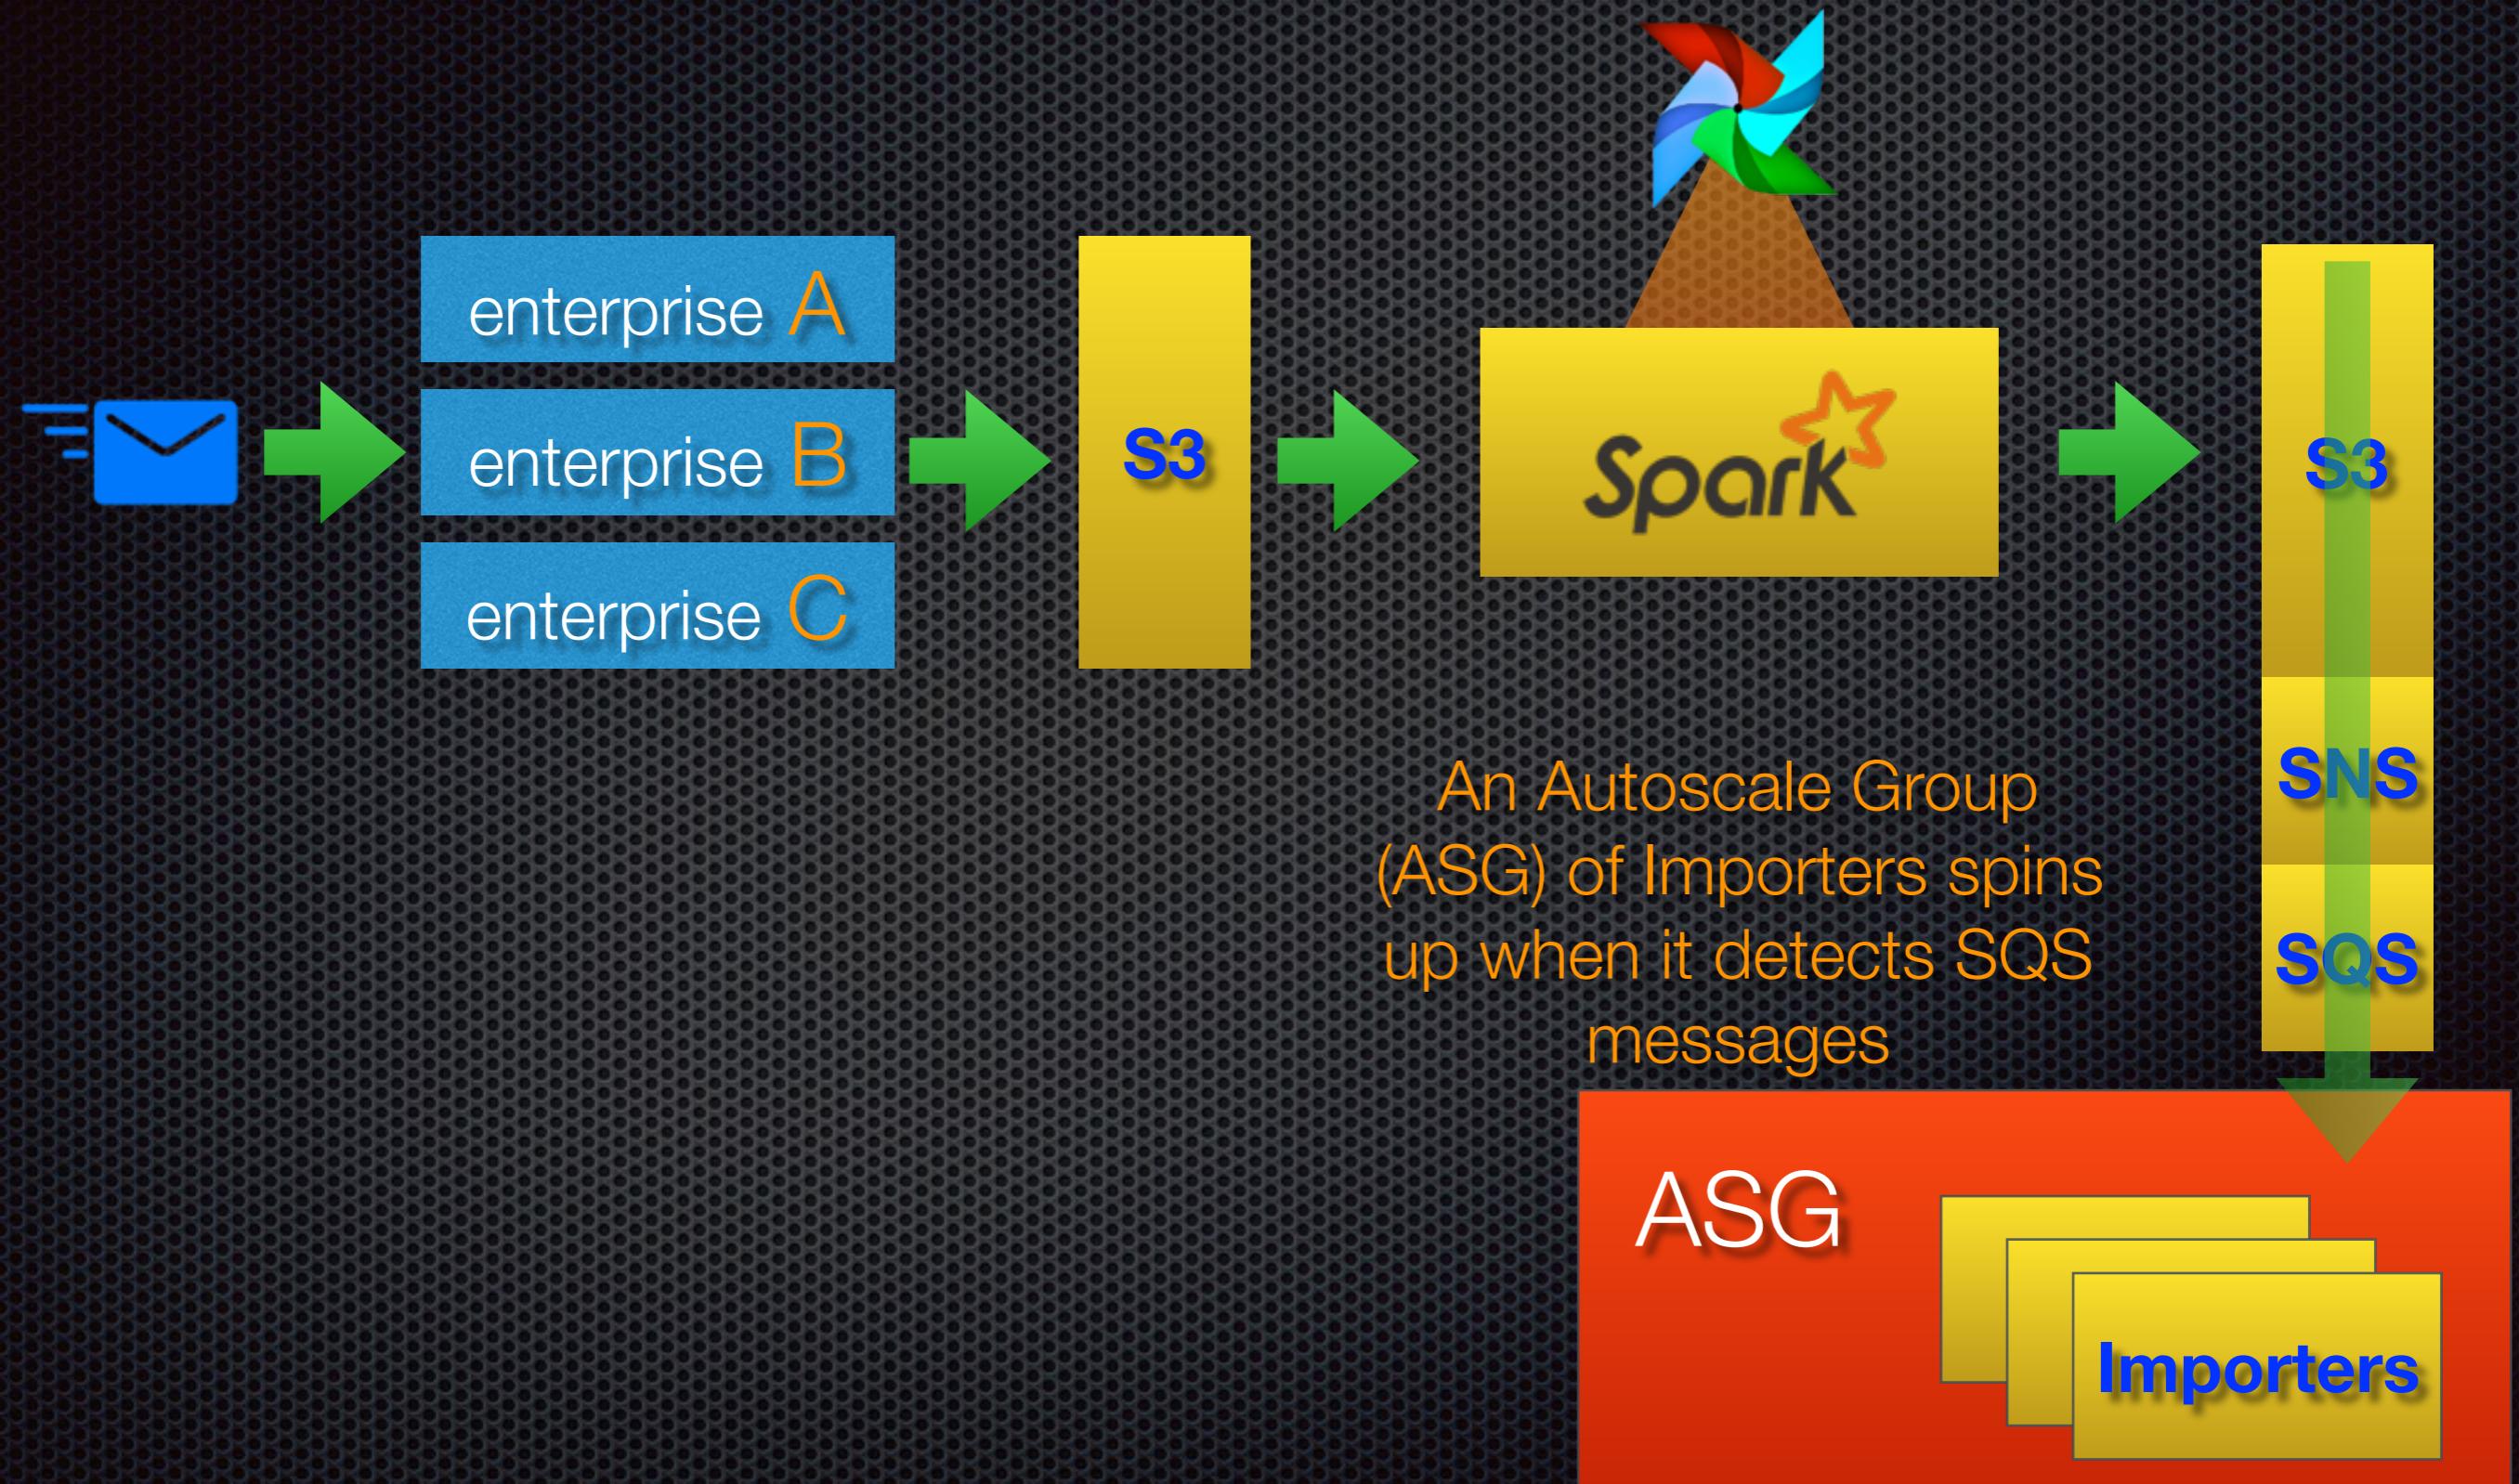
# Use-Case : Message Scoring



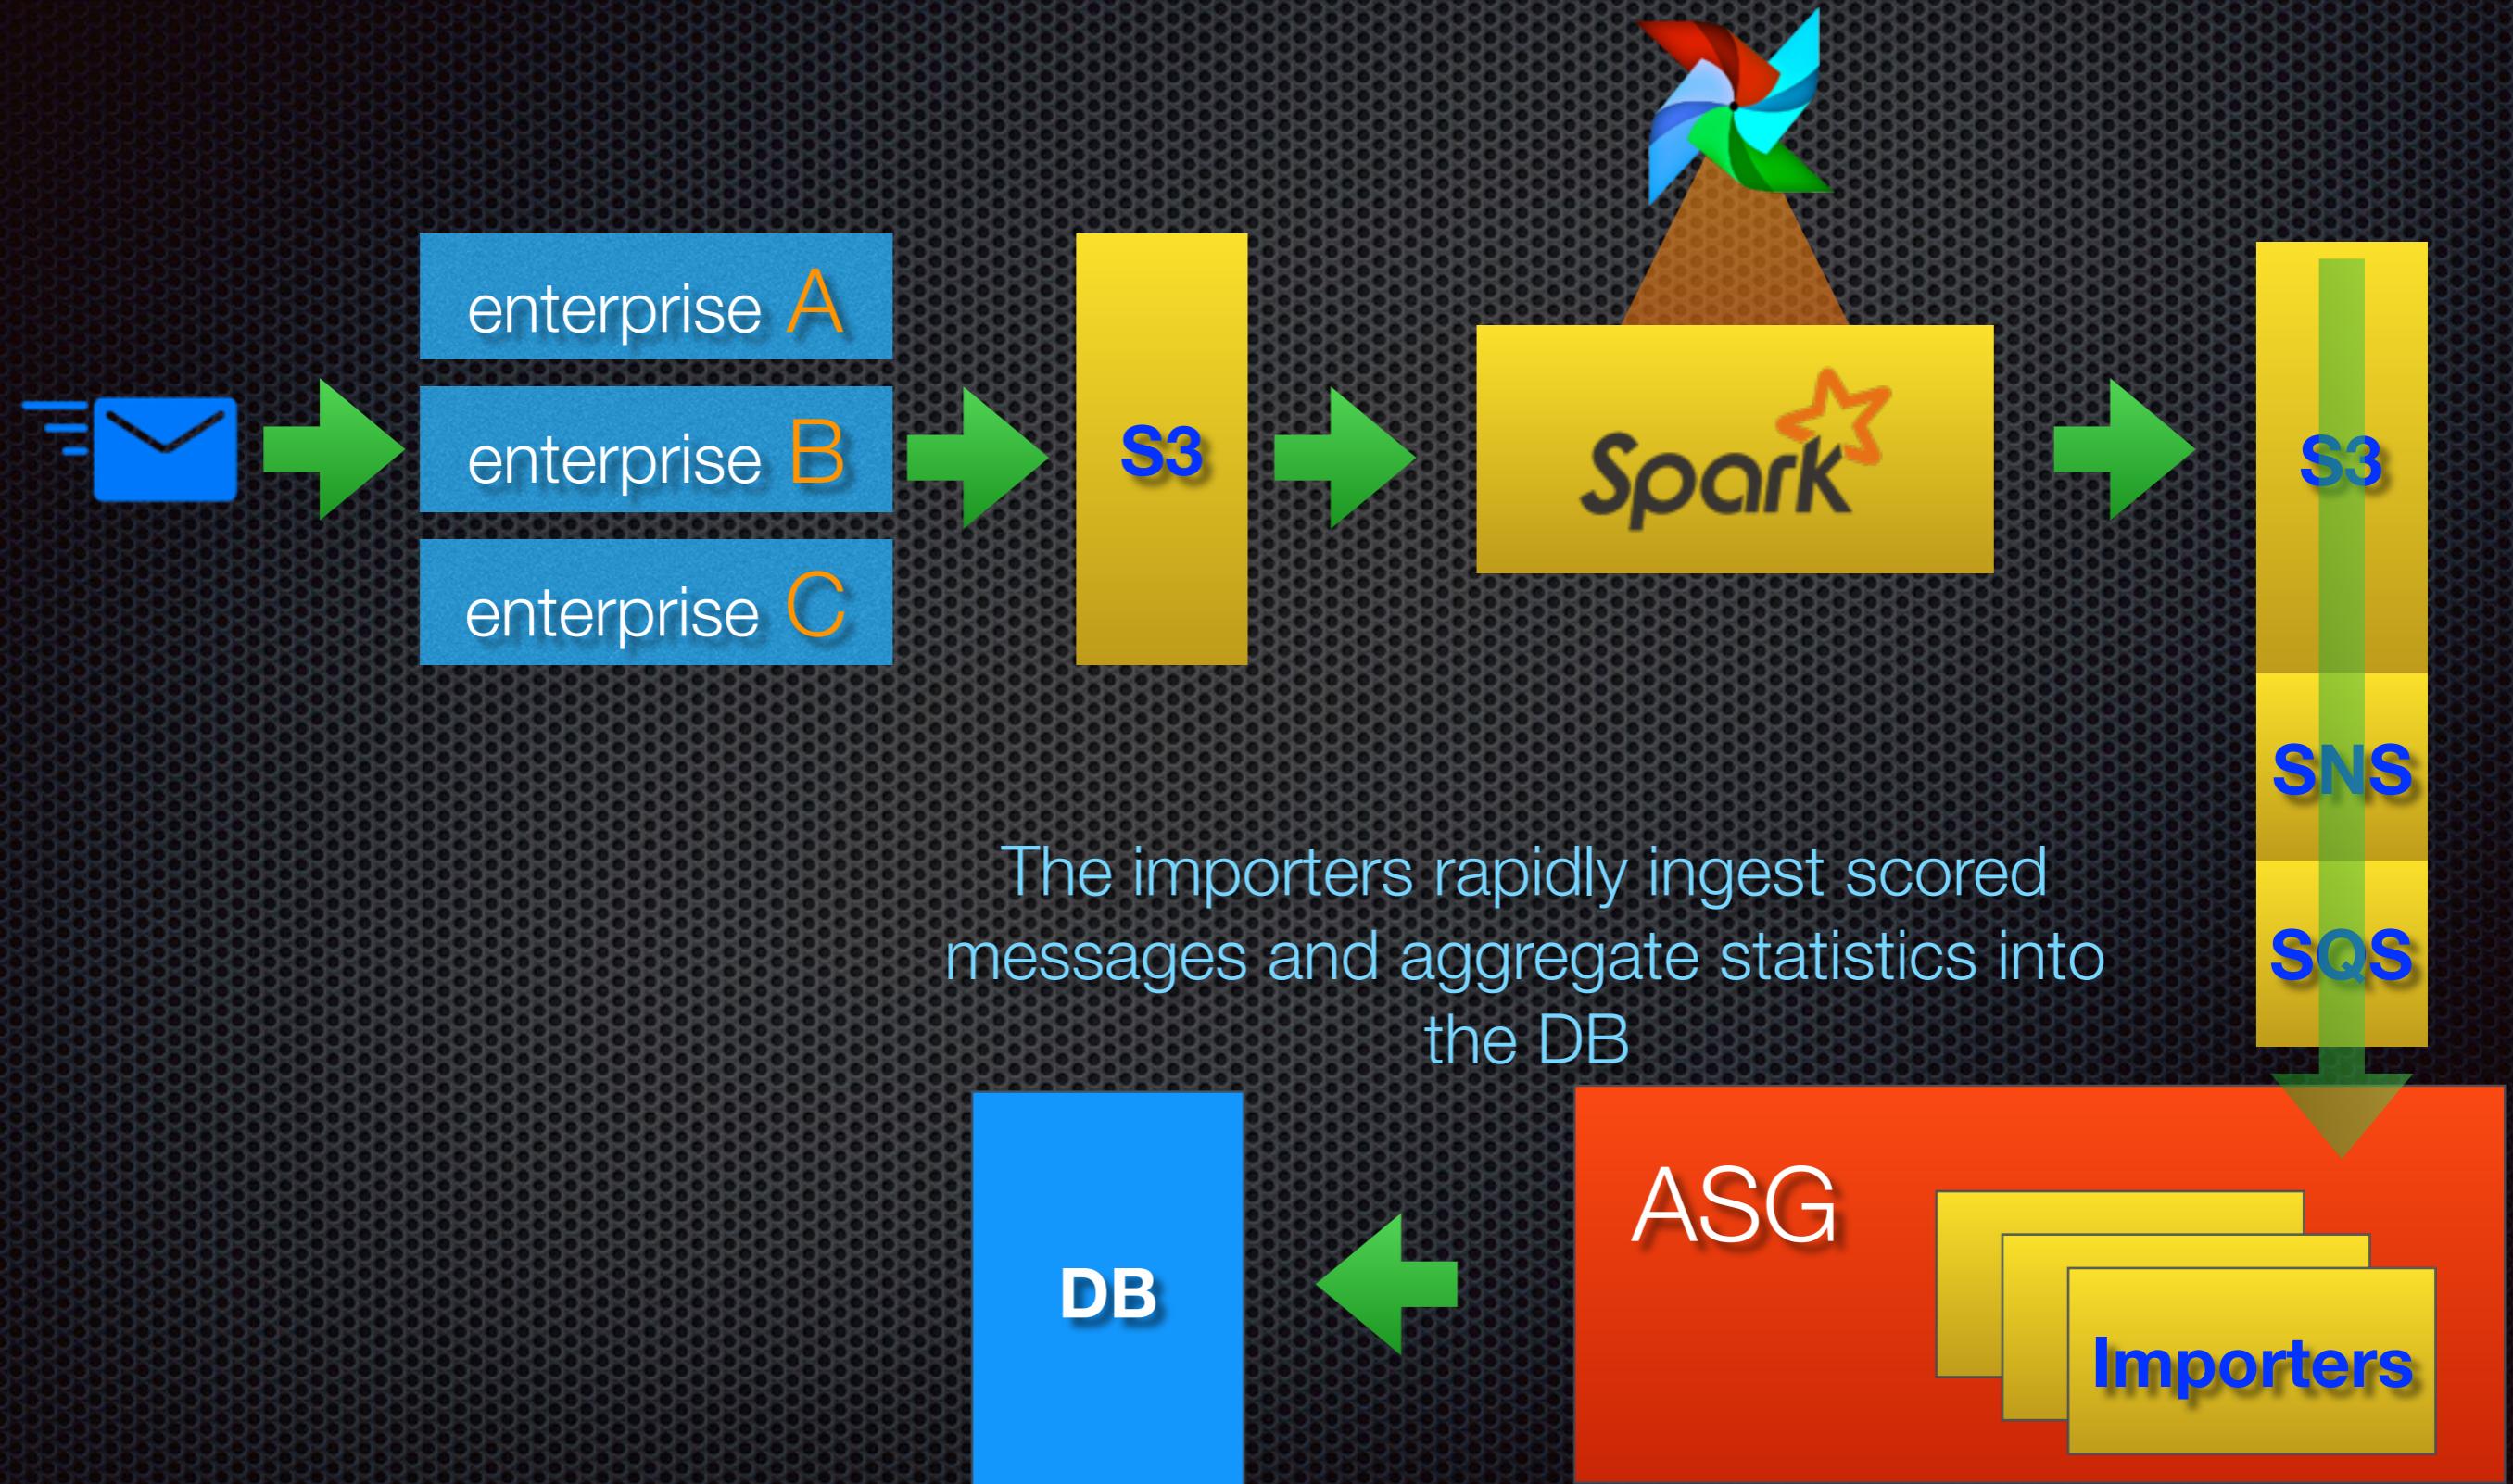
# Use-Case : Message Scoring



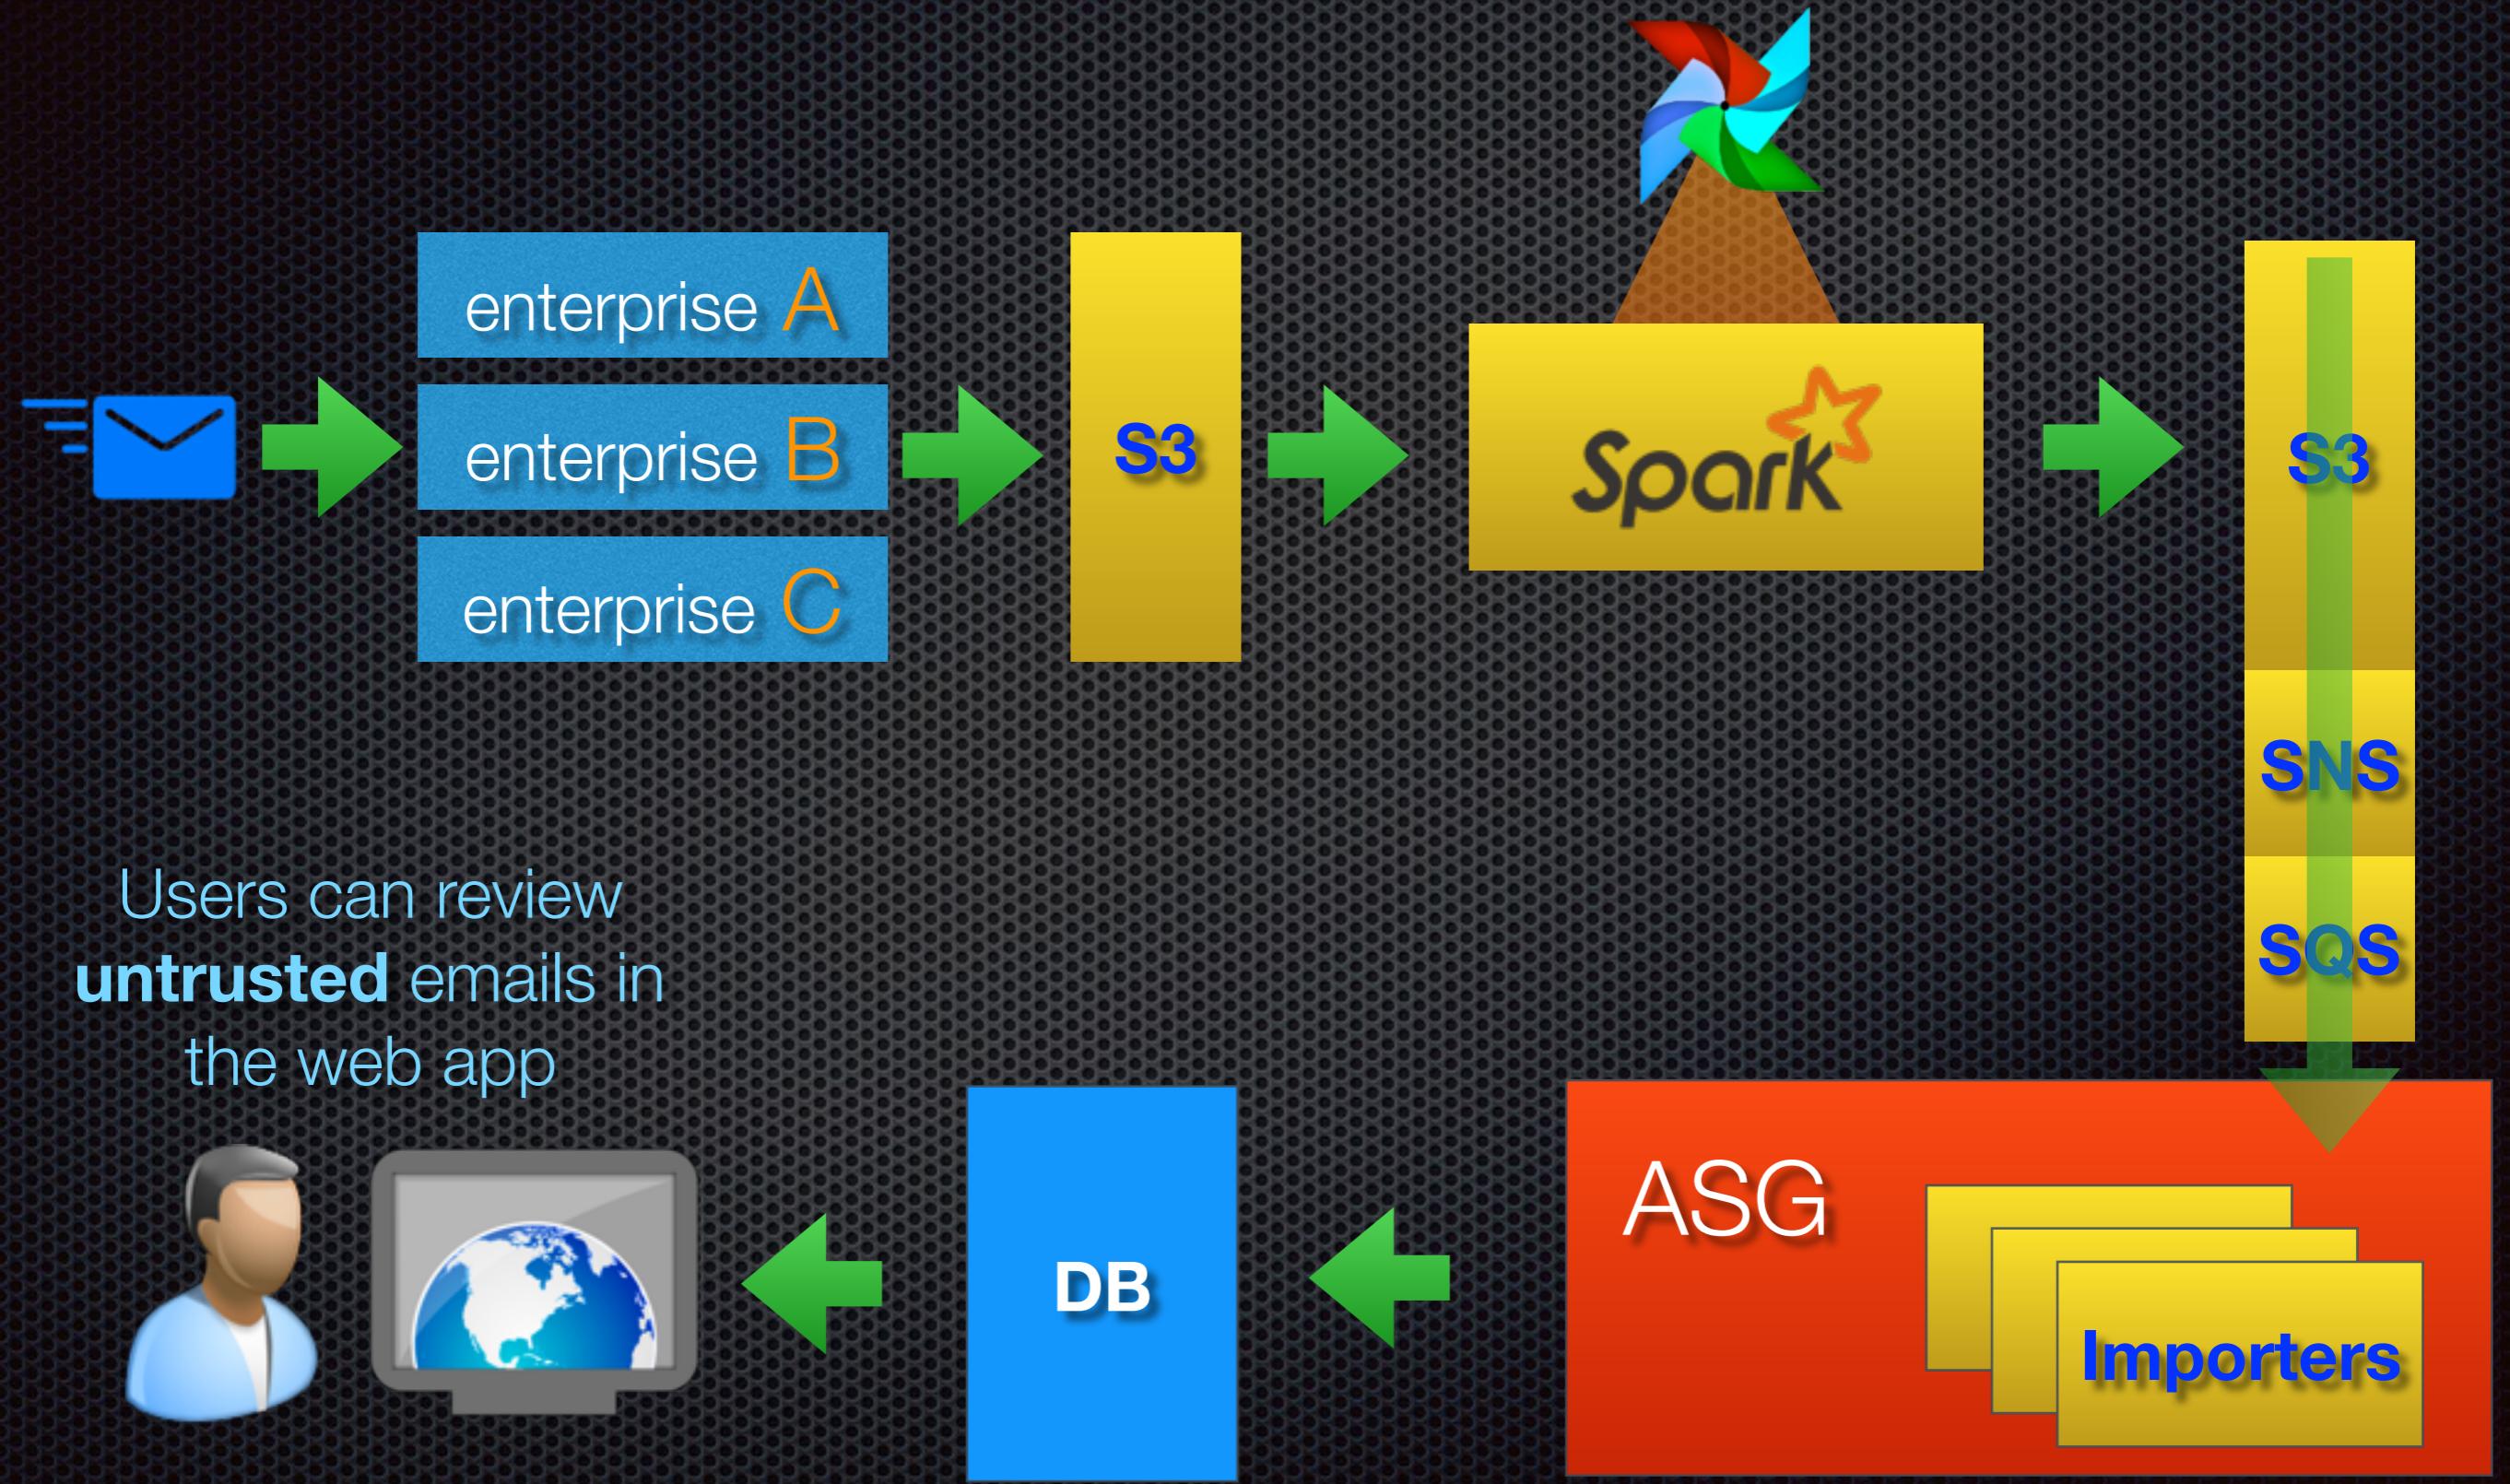
# Use-Case : Message Scoring



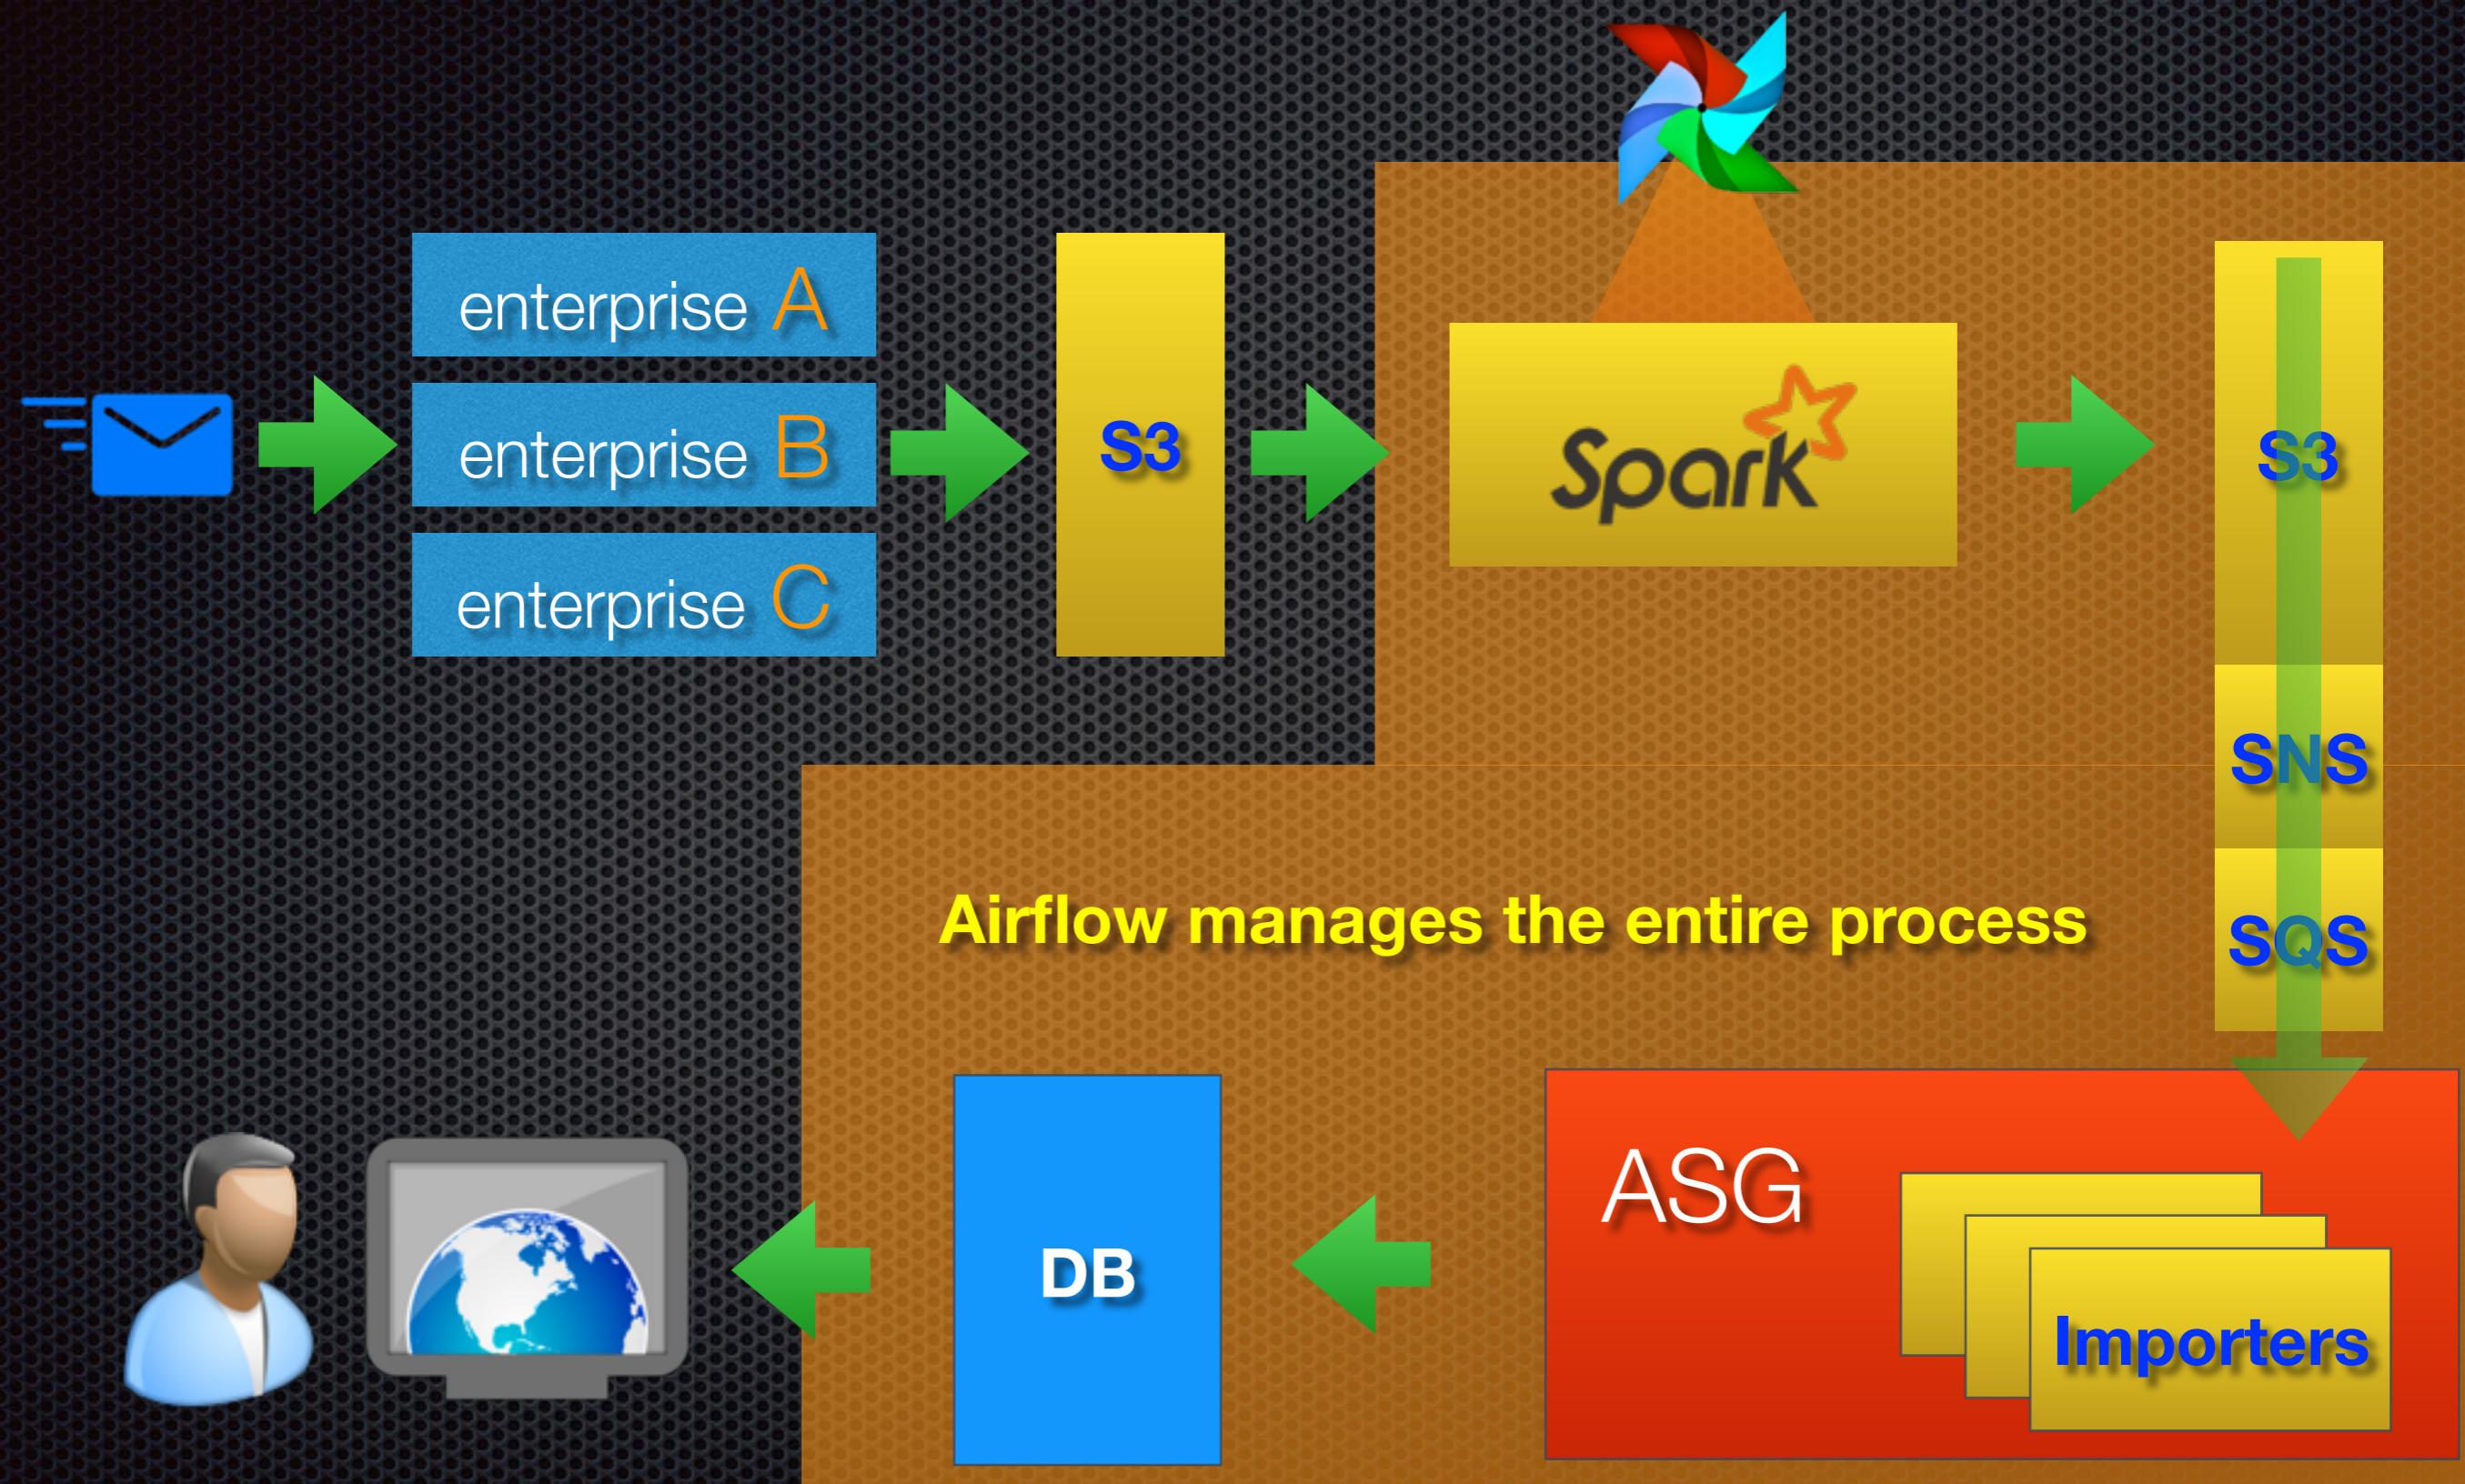
# Use-Case : Message Scoring



# Use-Case : Message Scoring



# Use-Case : Message Scoring

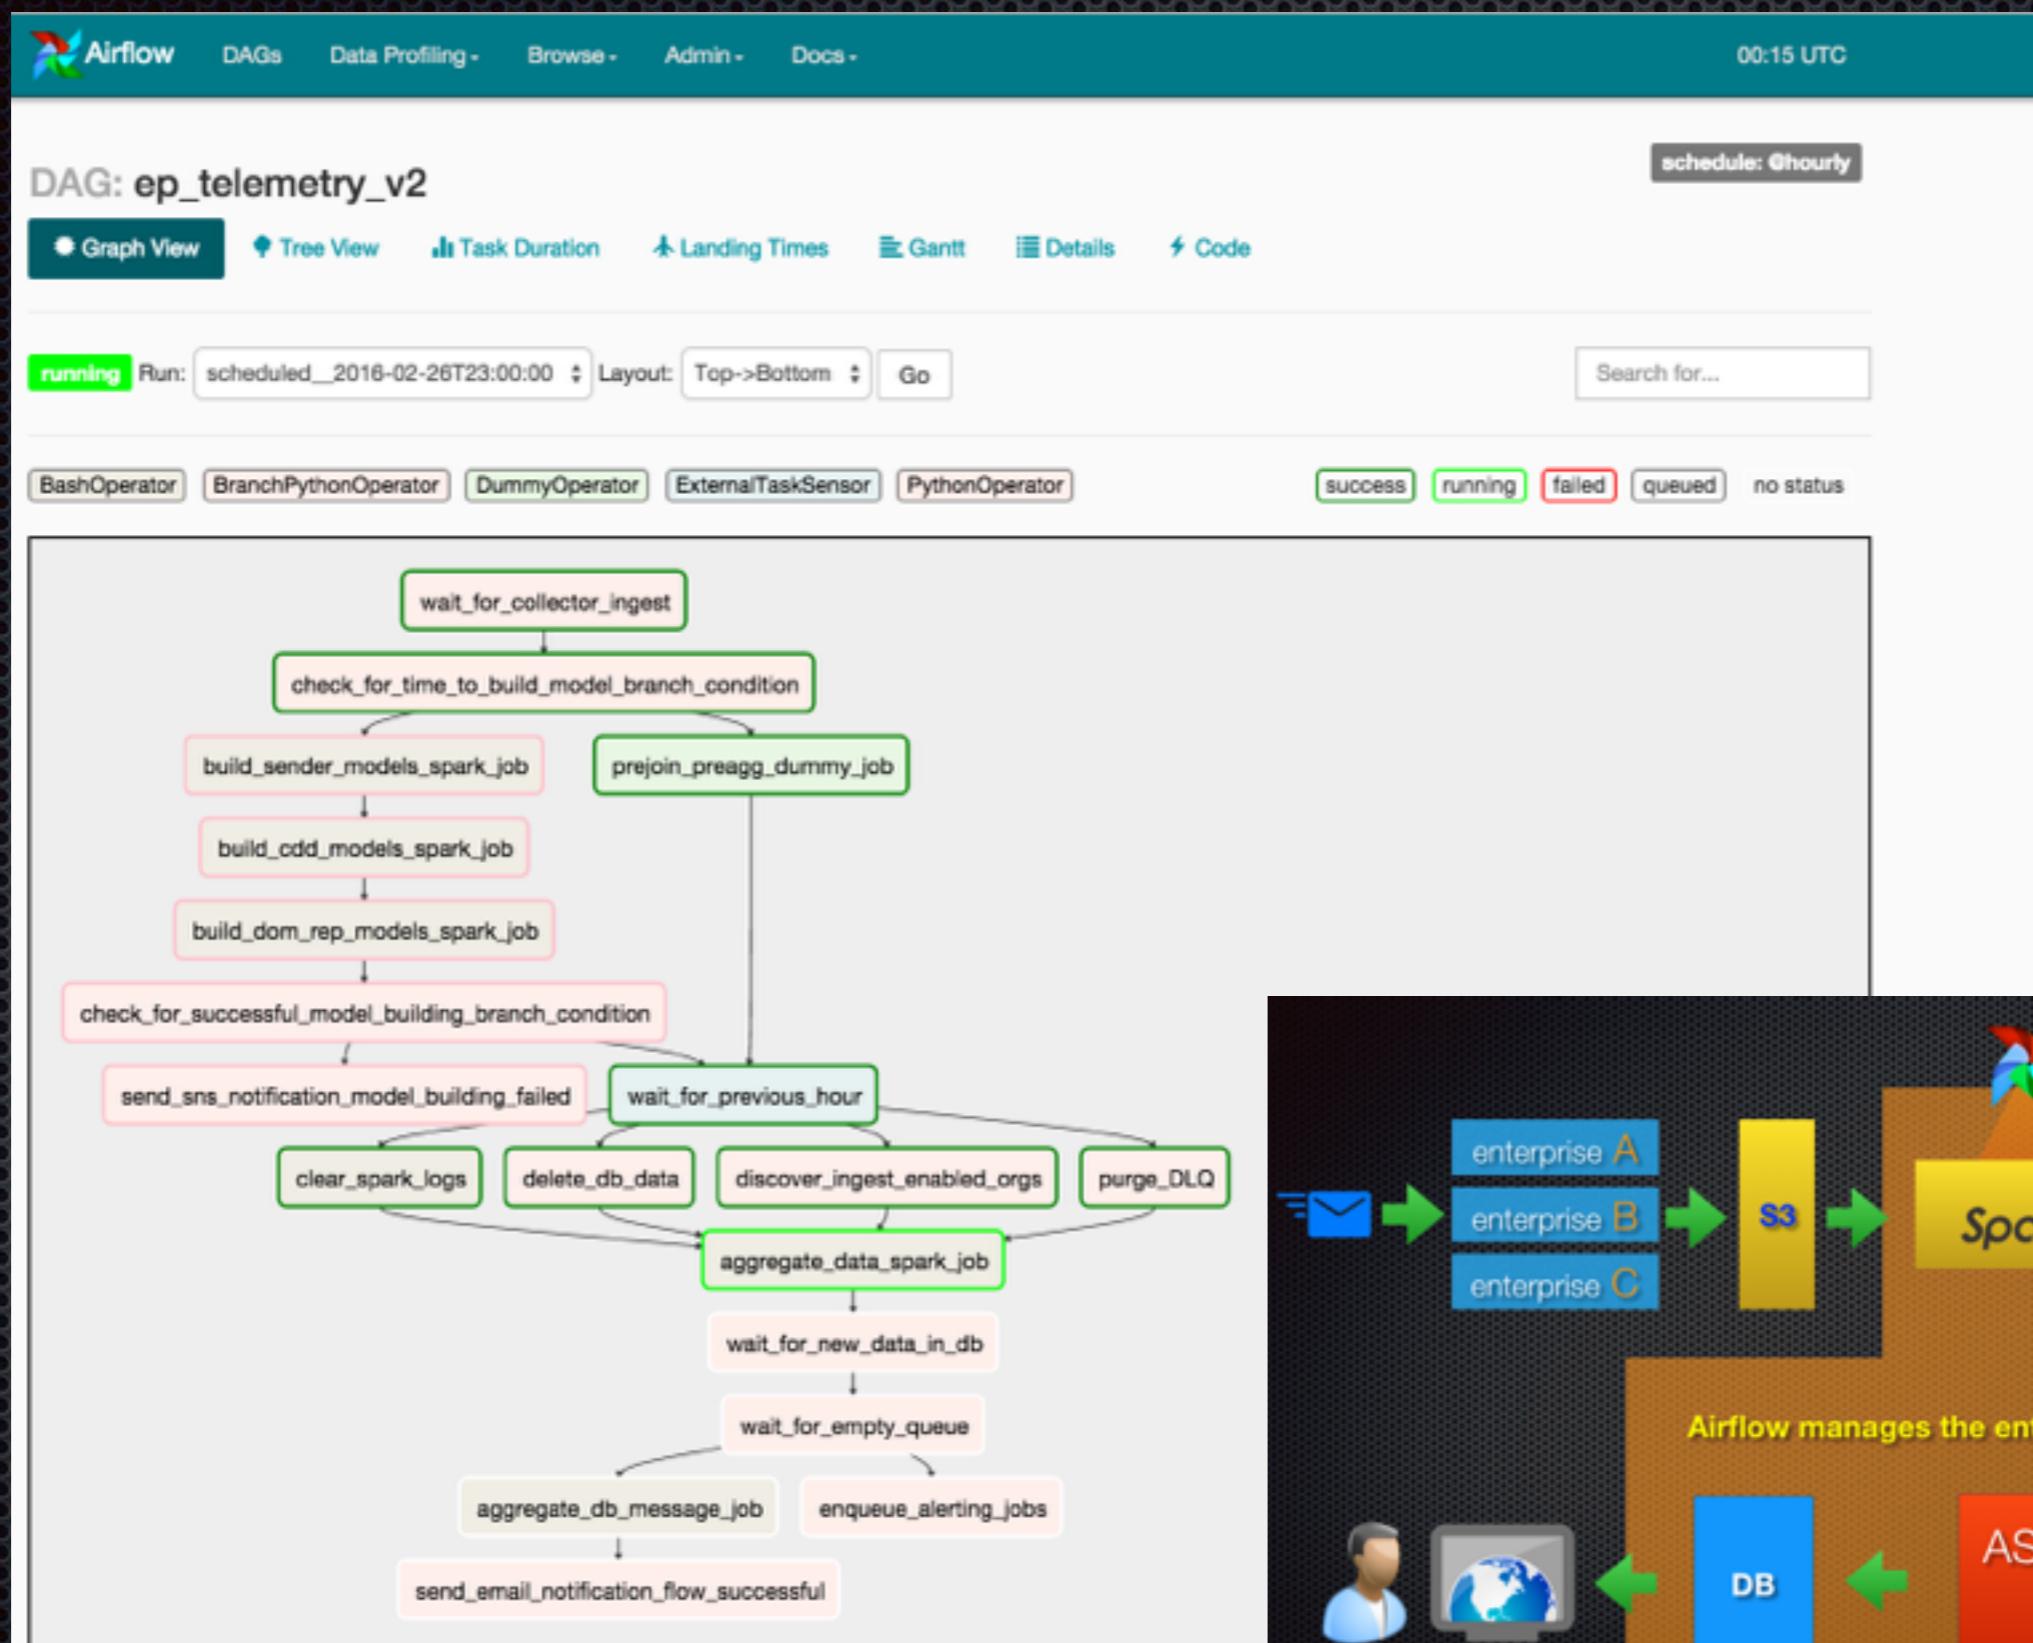


# Use-Case : Message Scoring

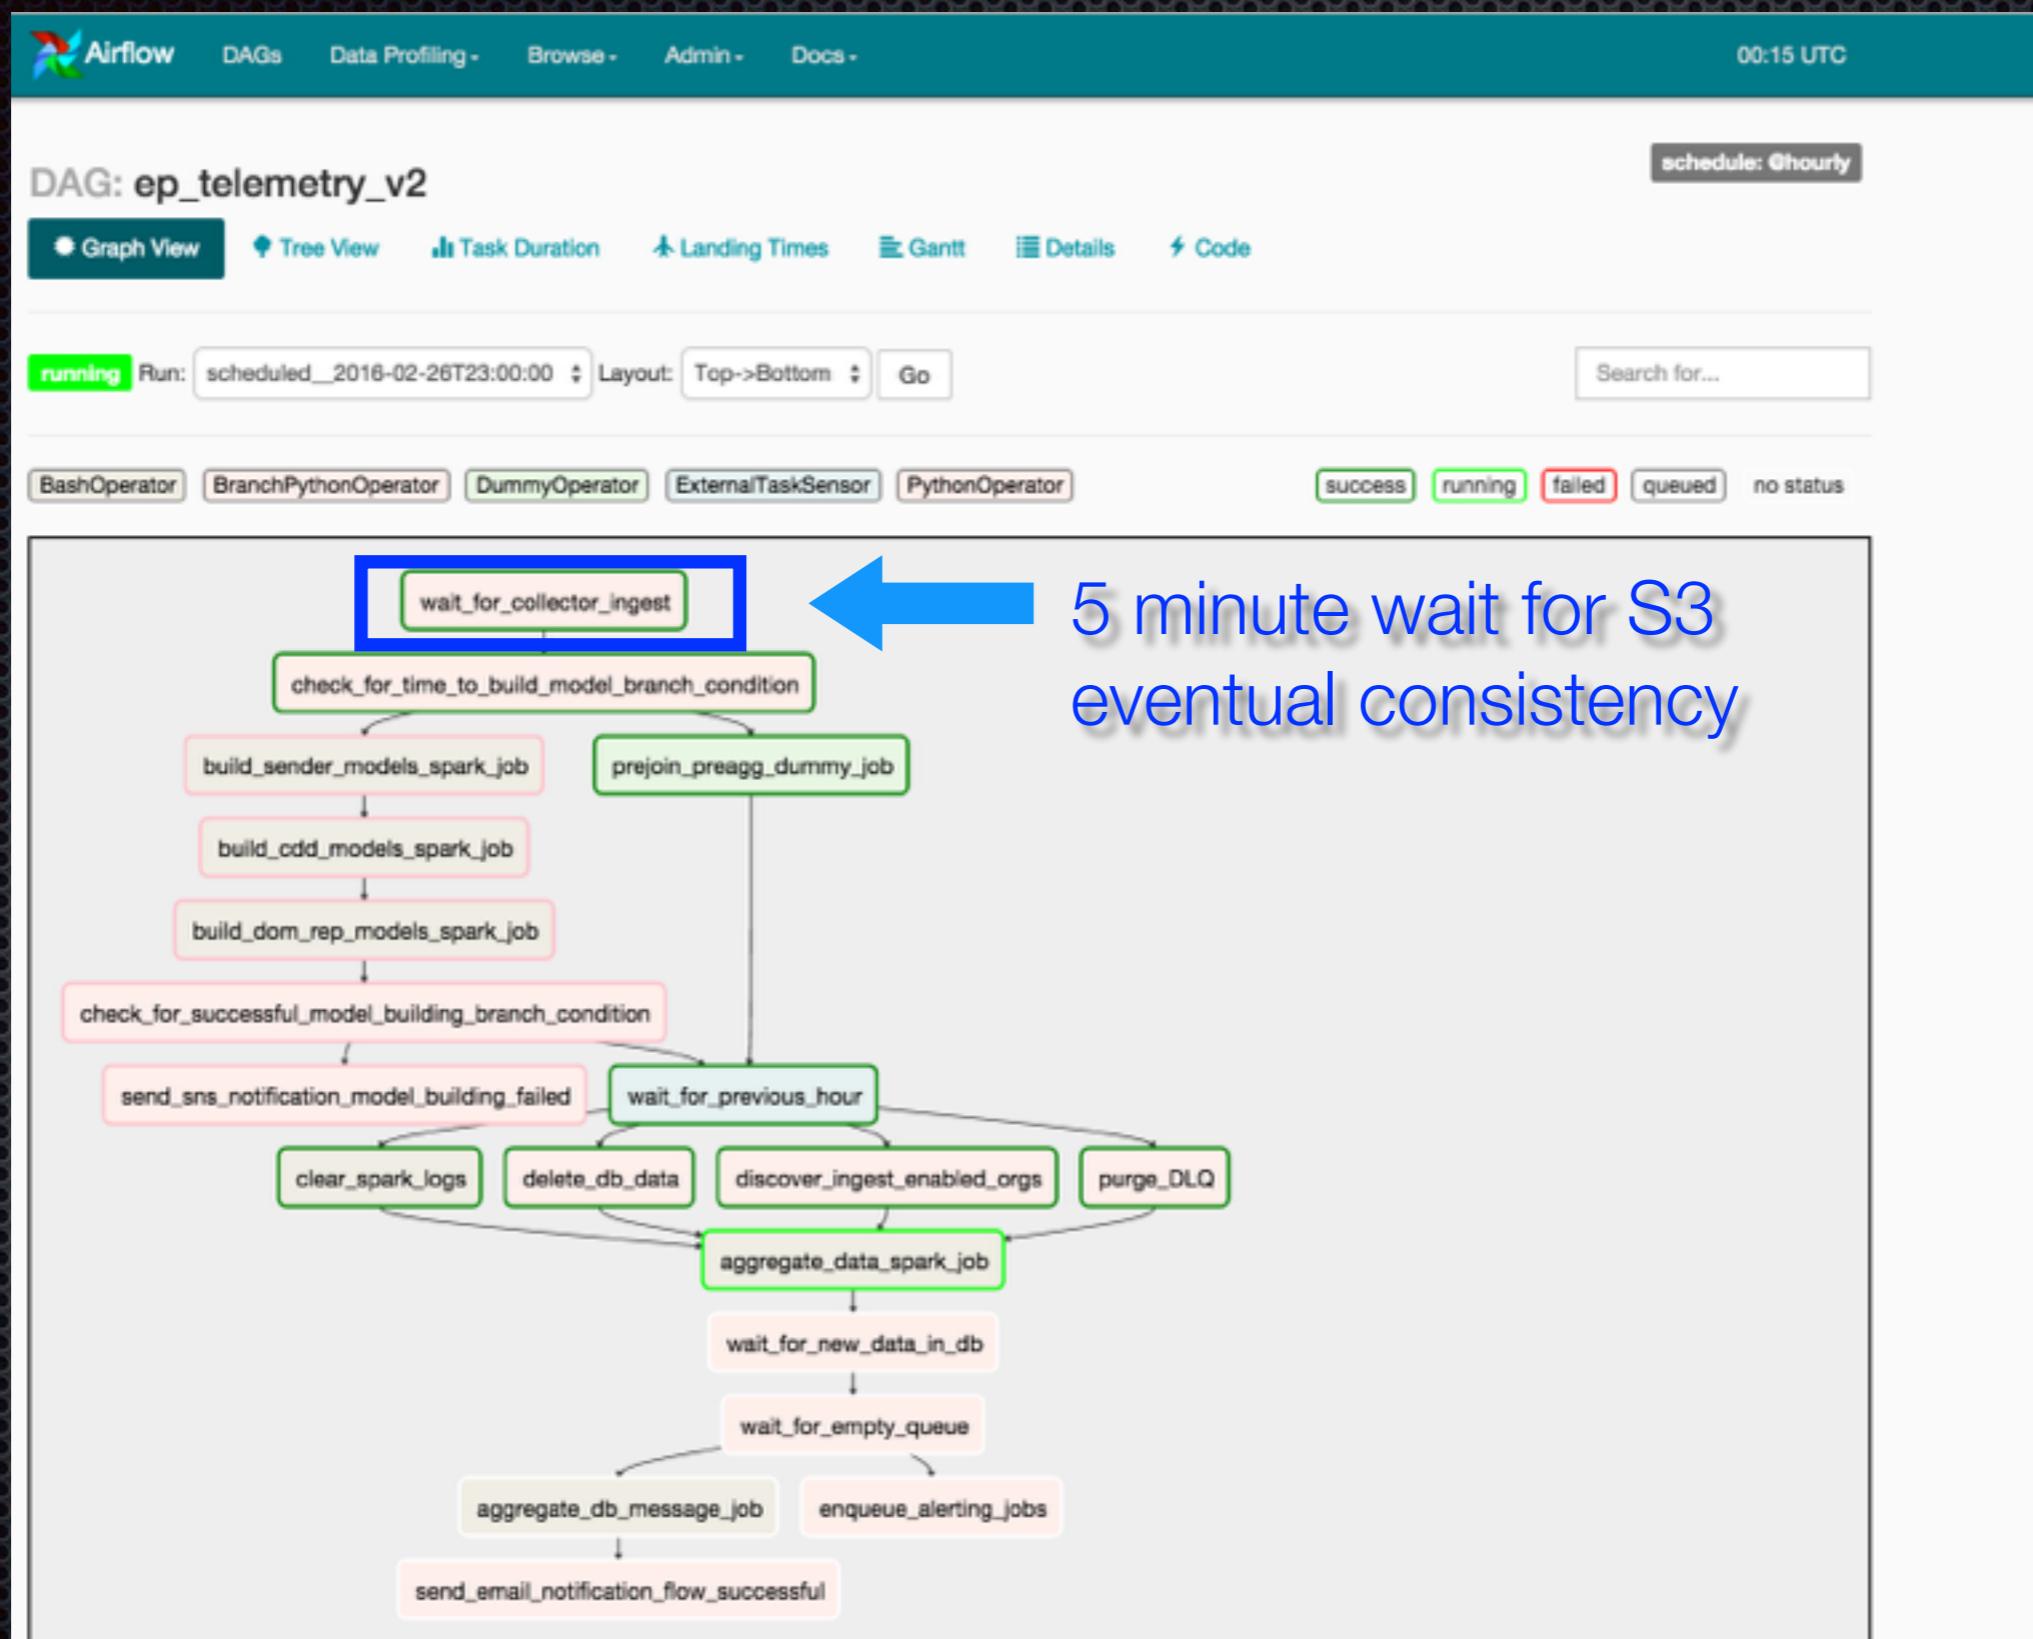
Airflow DAG



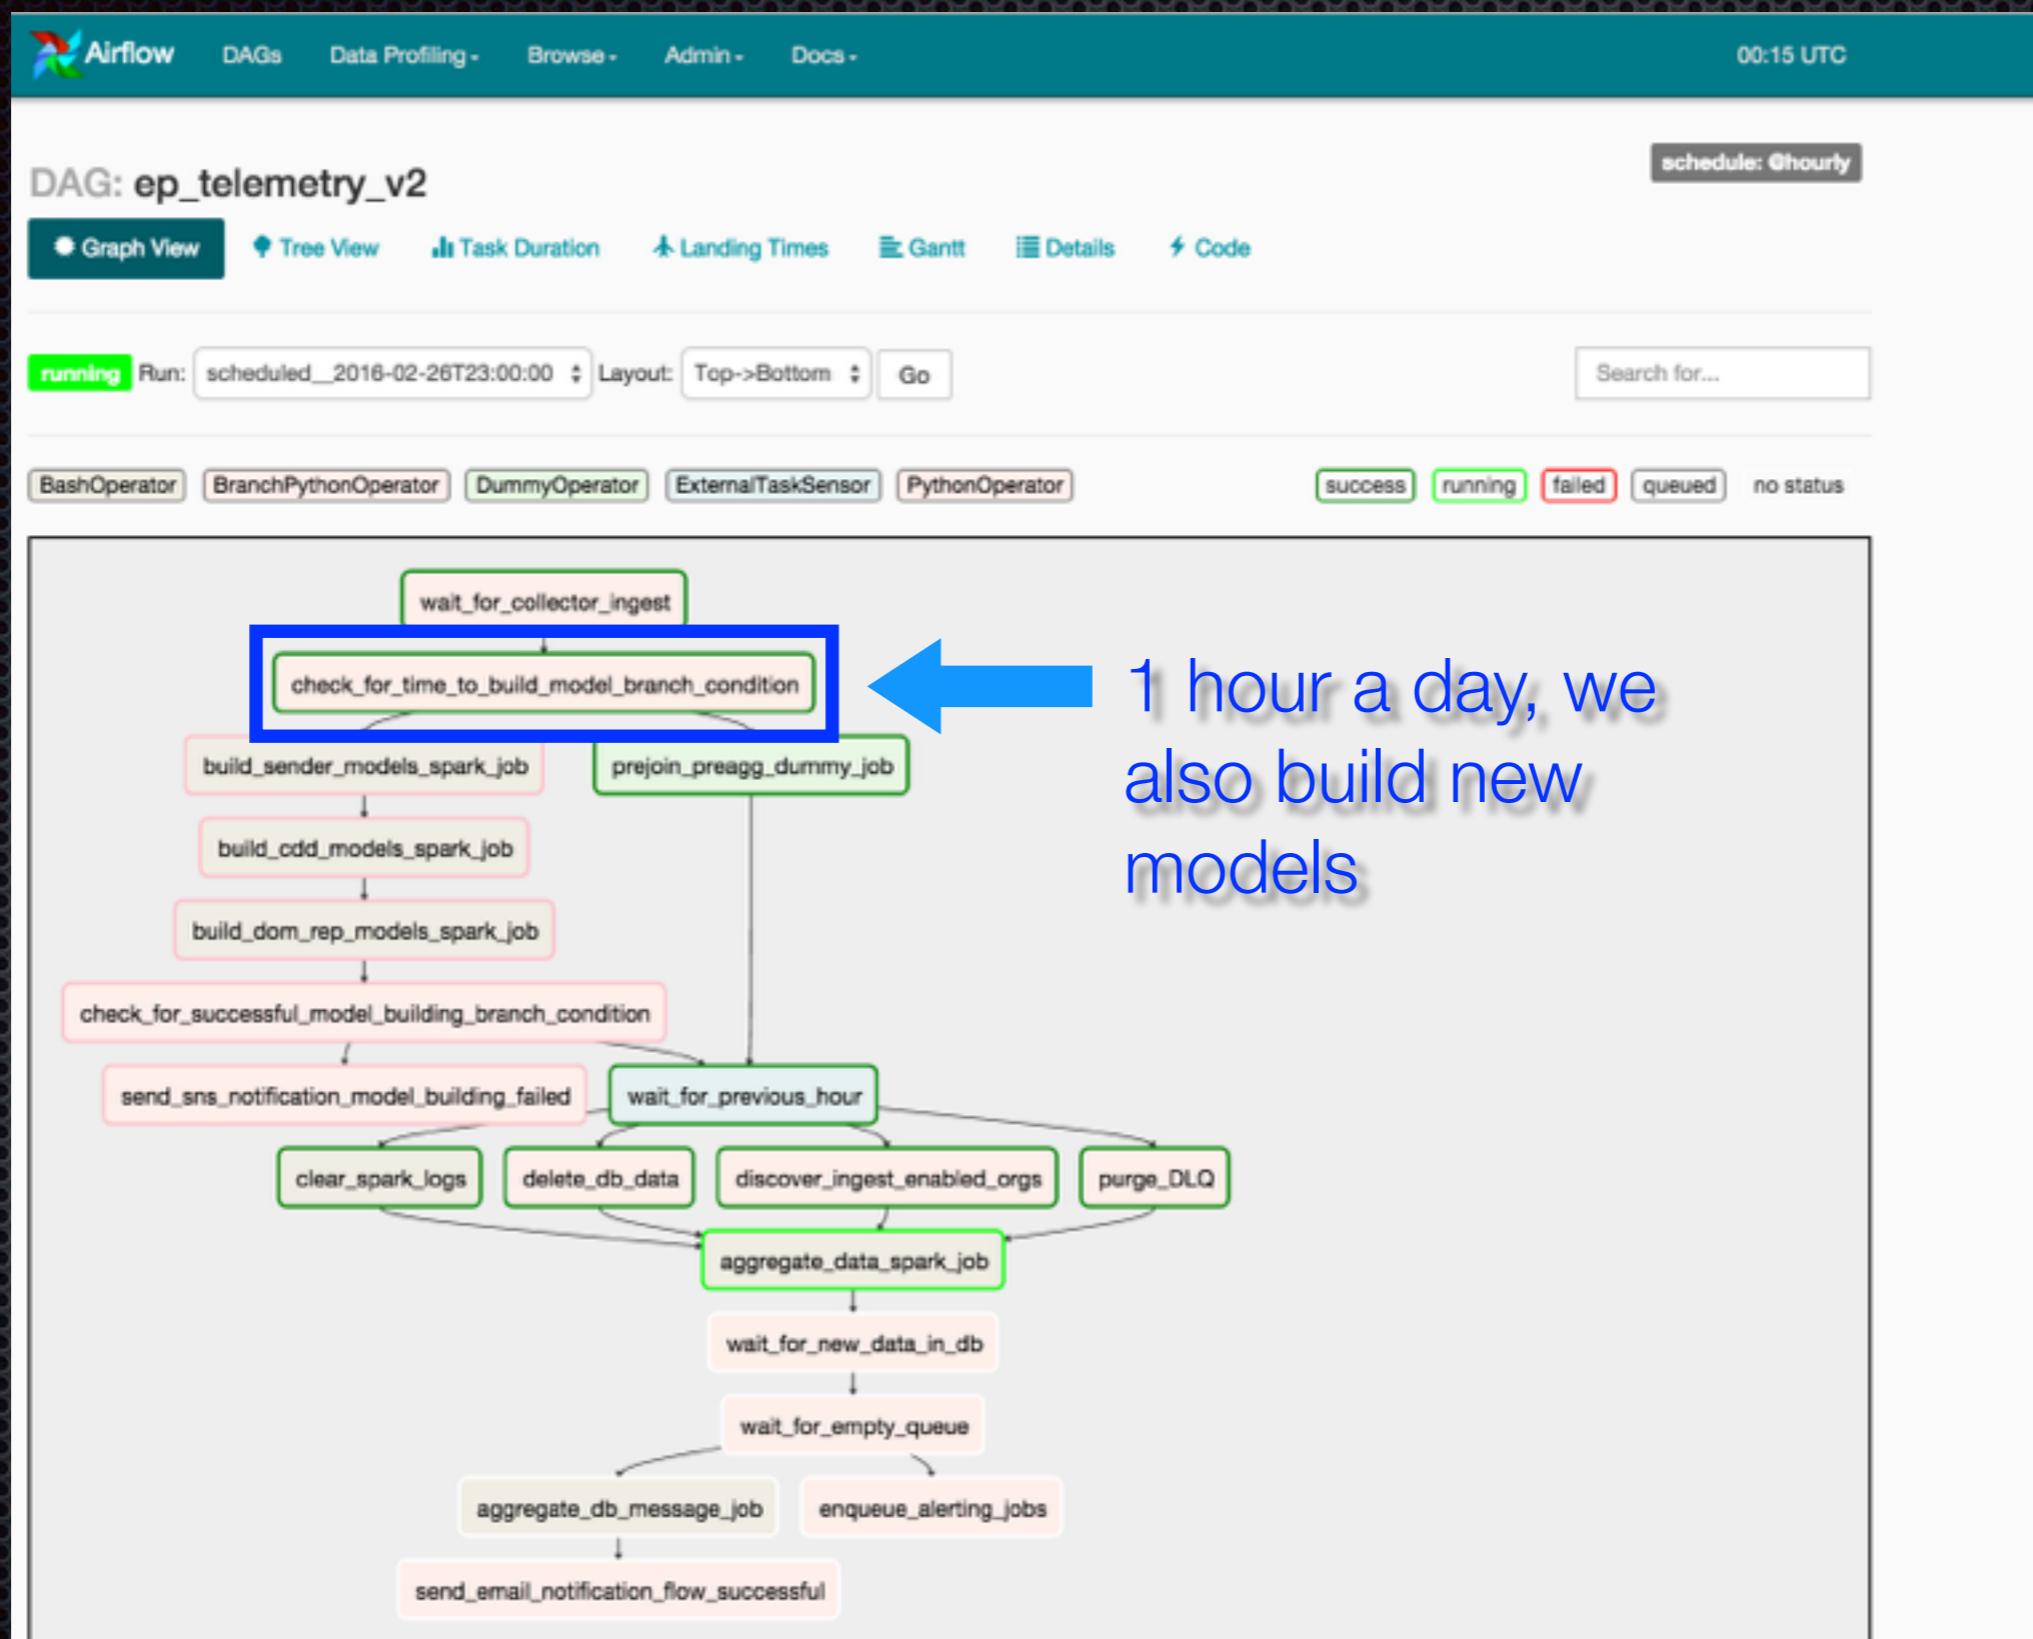
# Airflow DAG



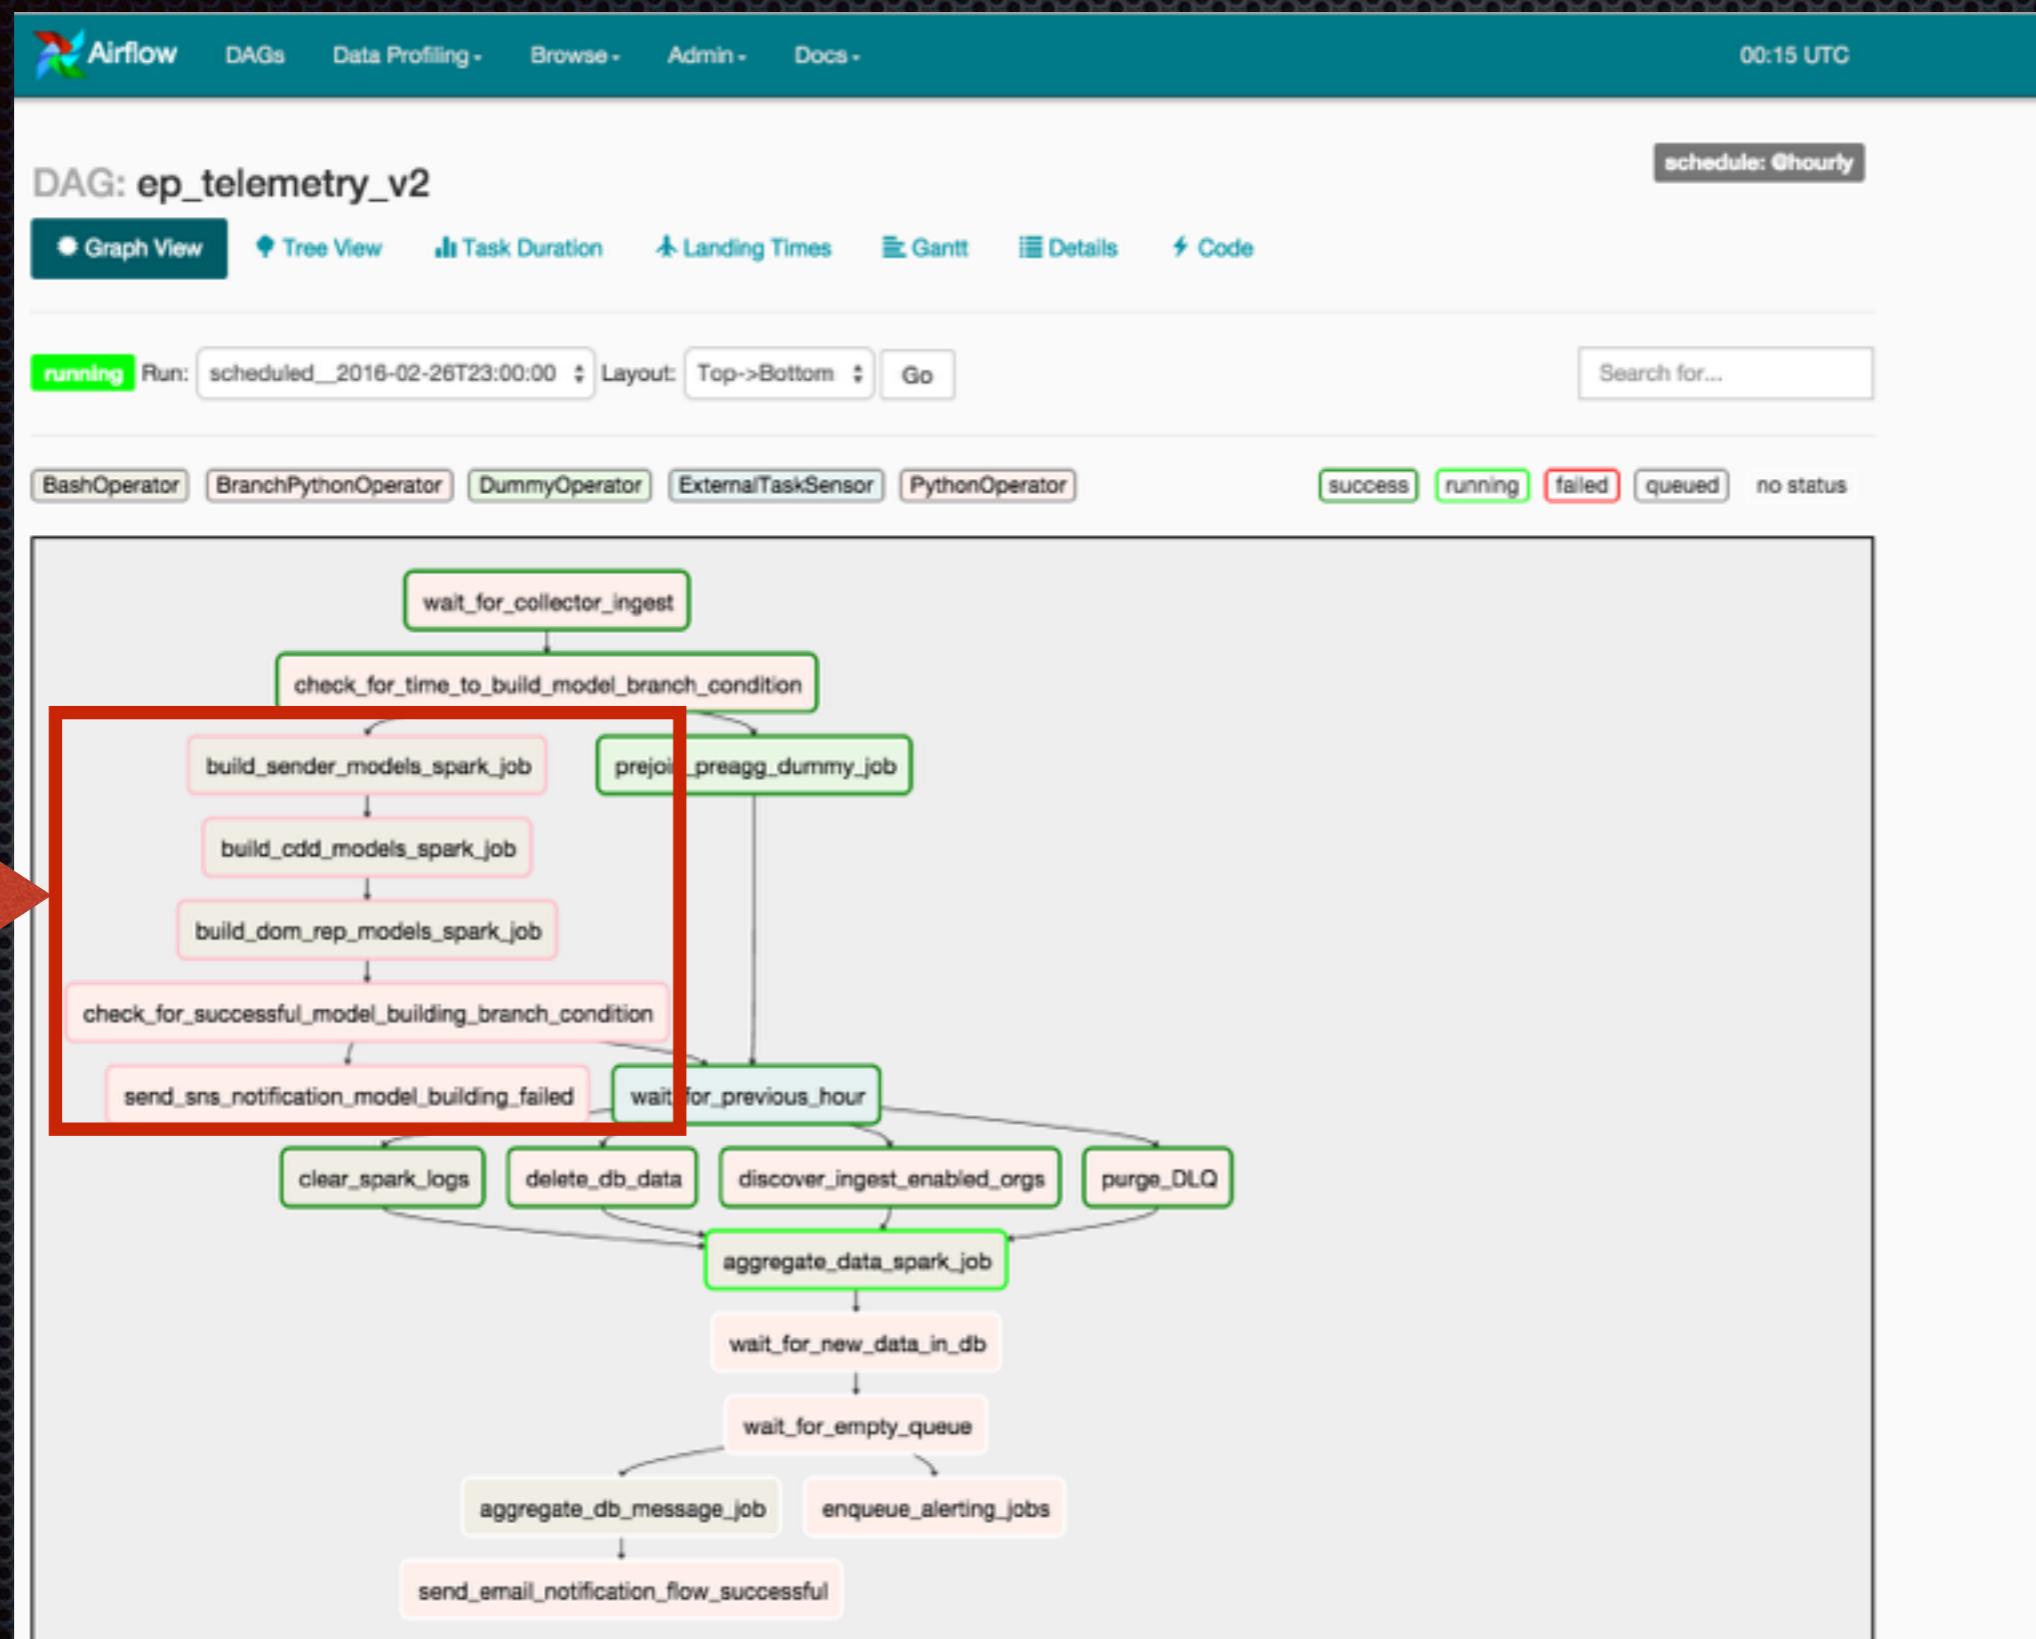
# Airflow DAG



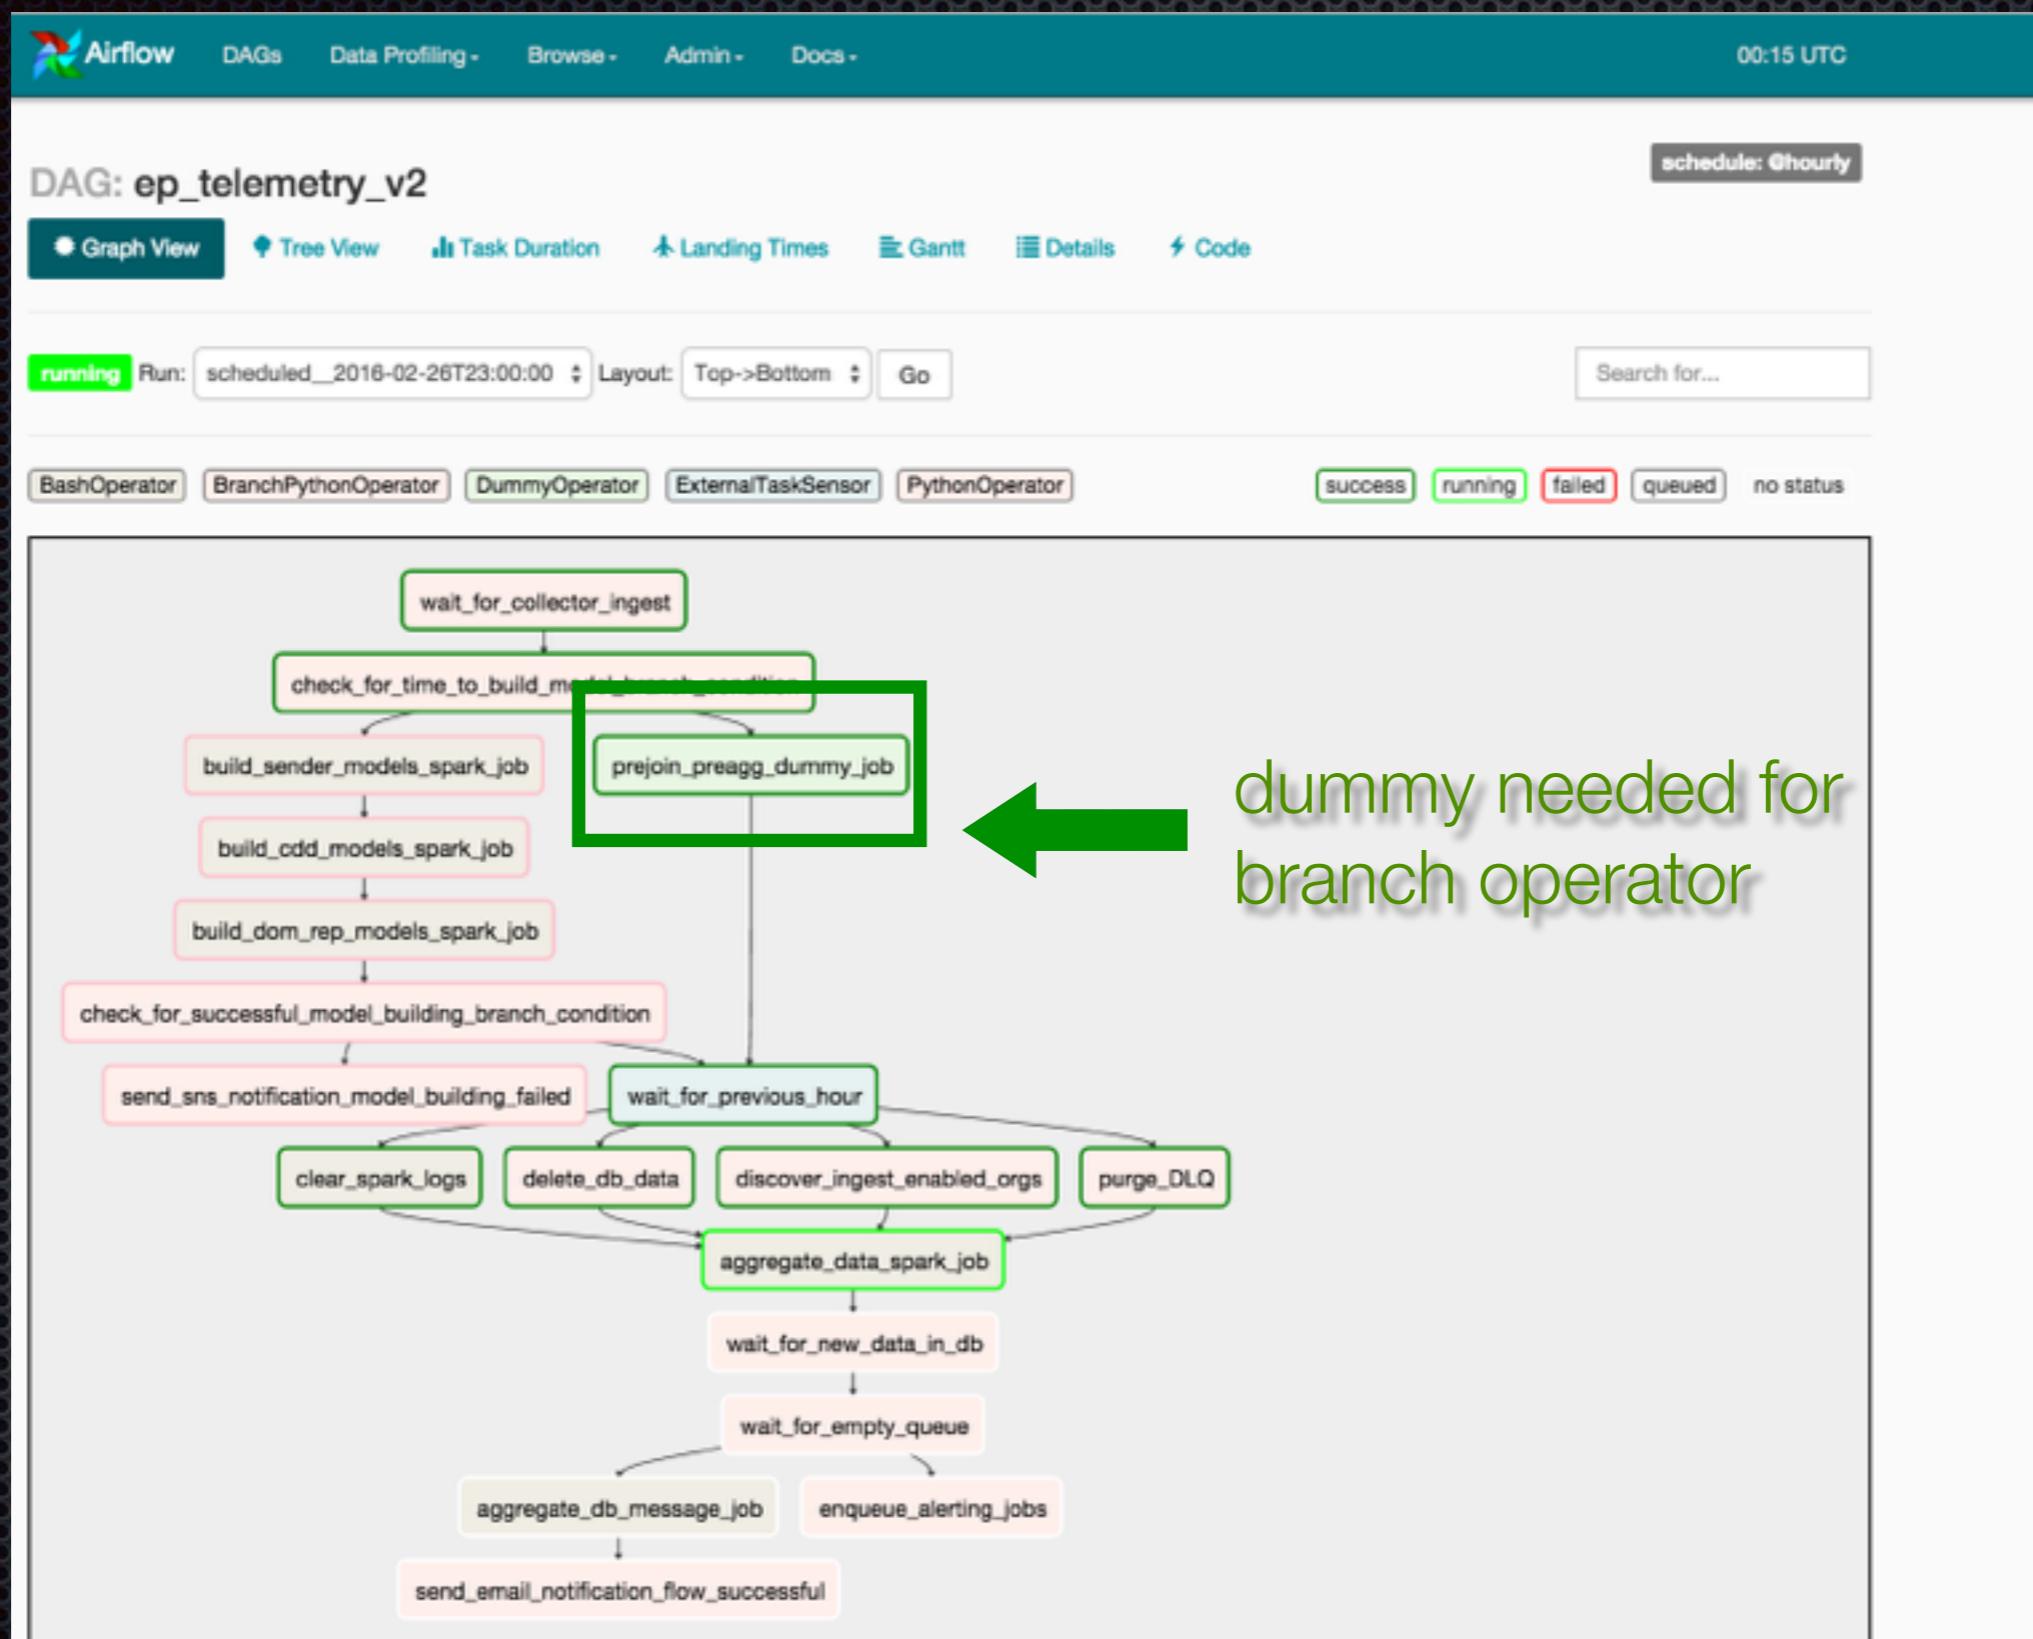
# Airflow DAG



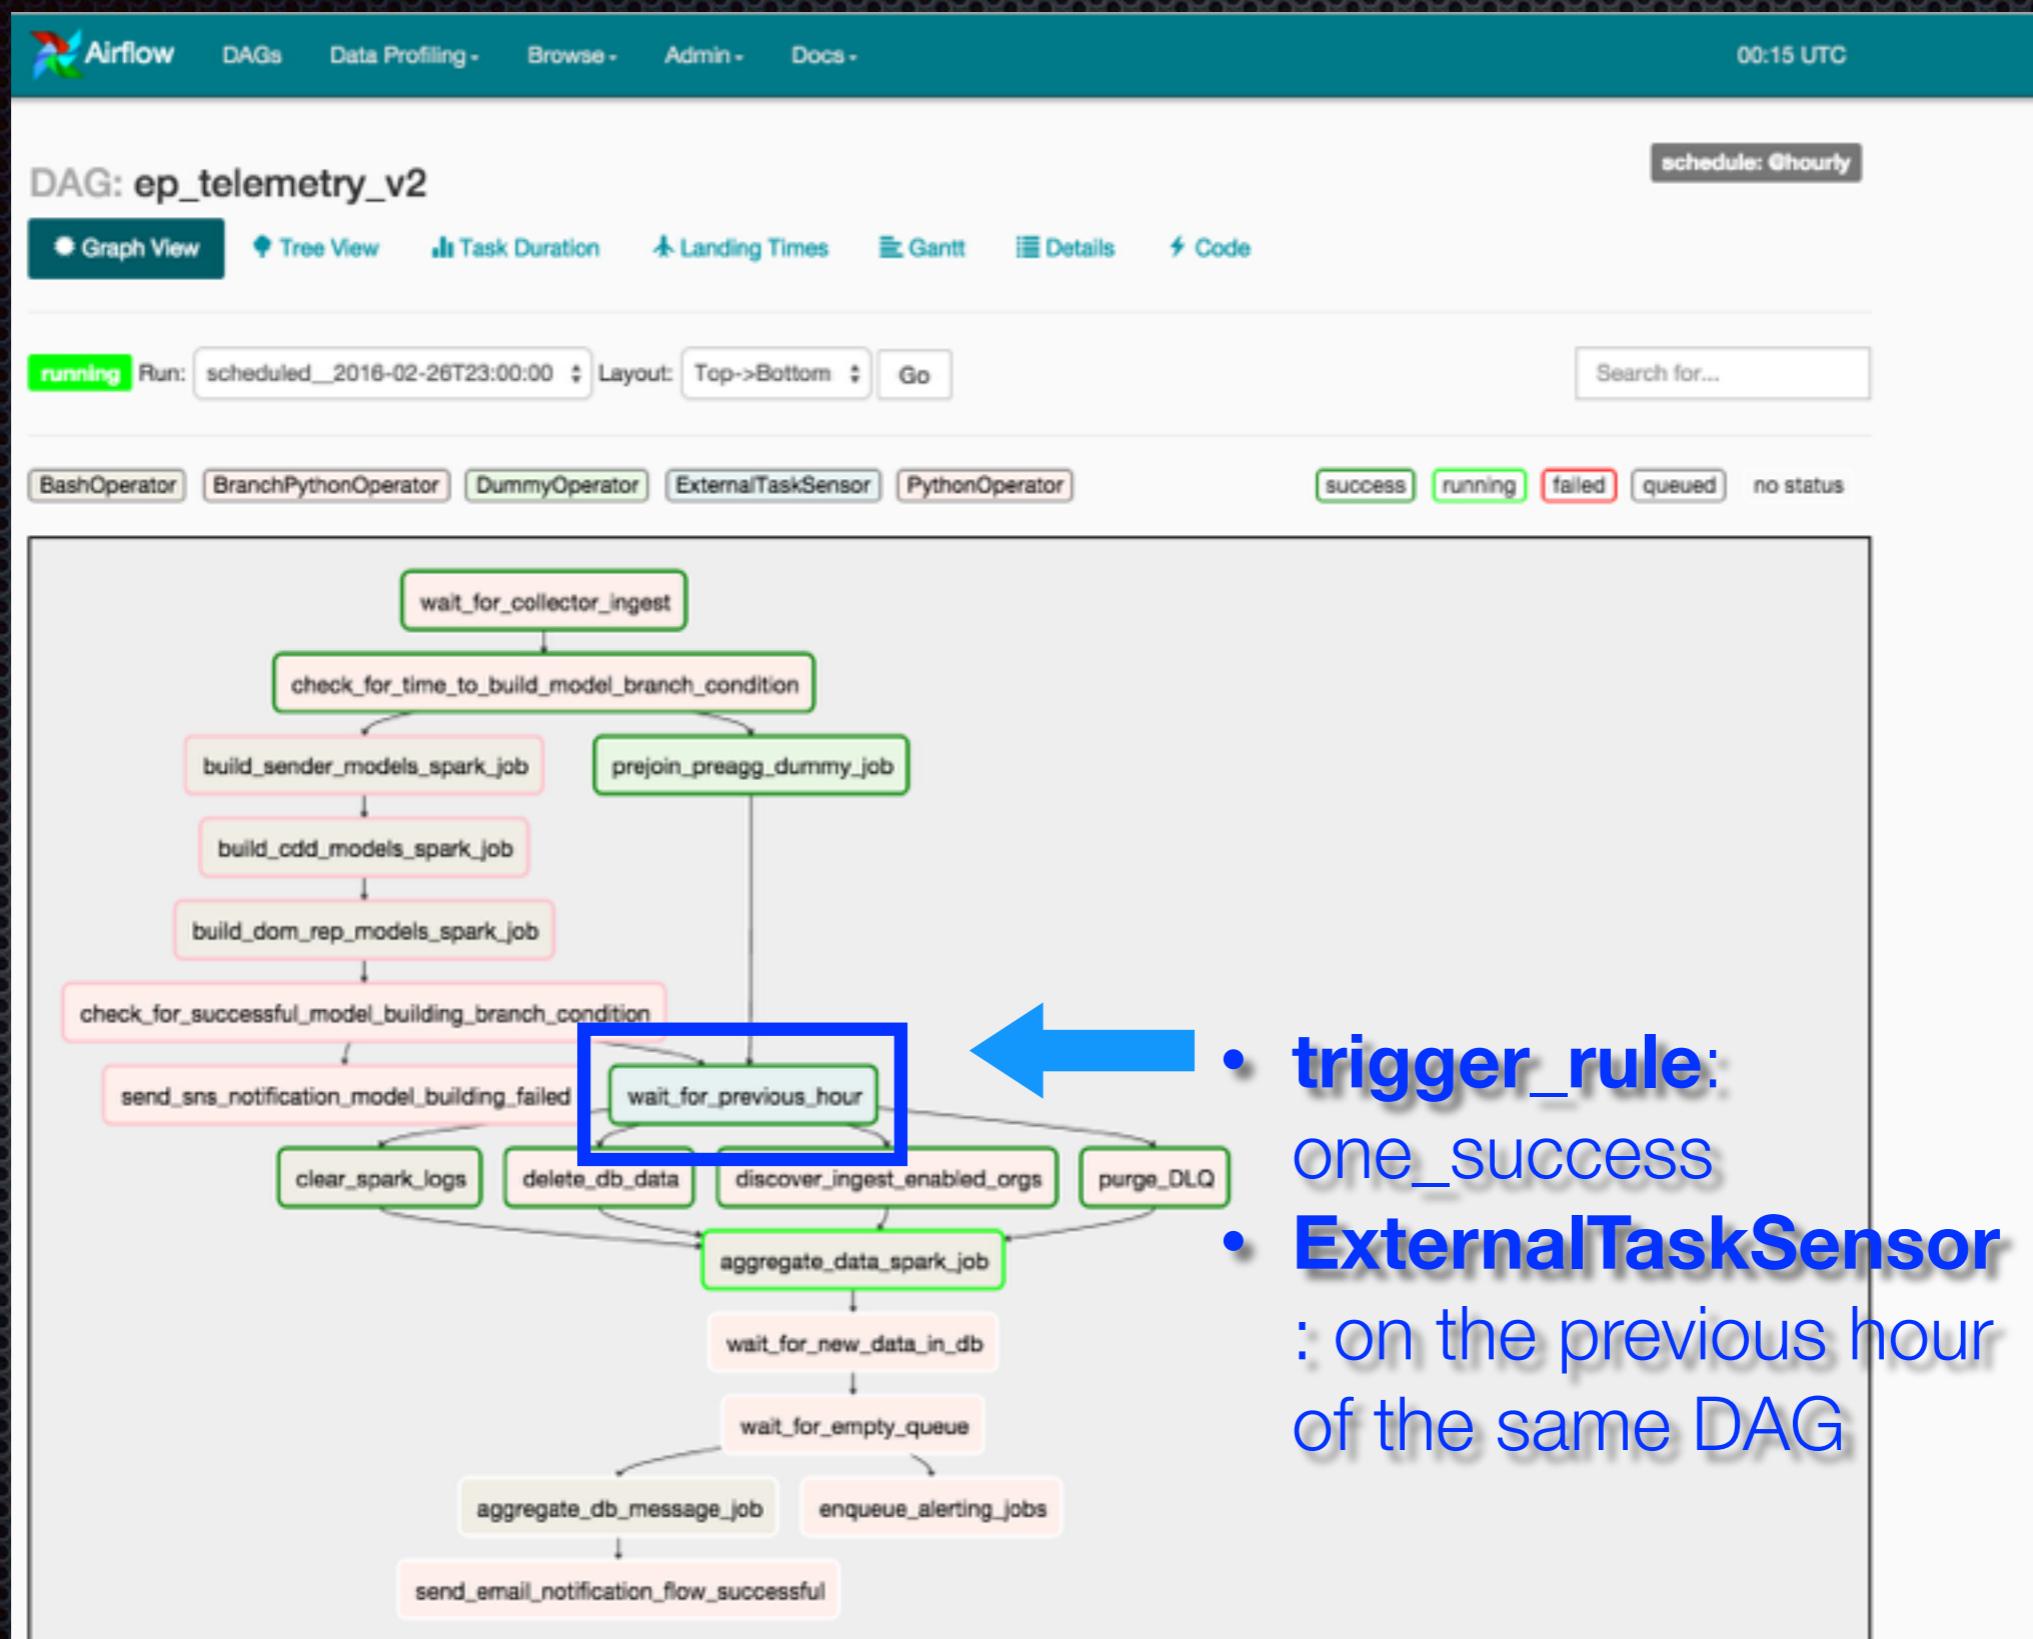
# Airflow DAG



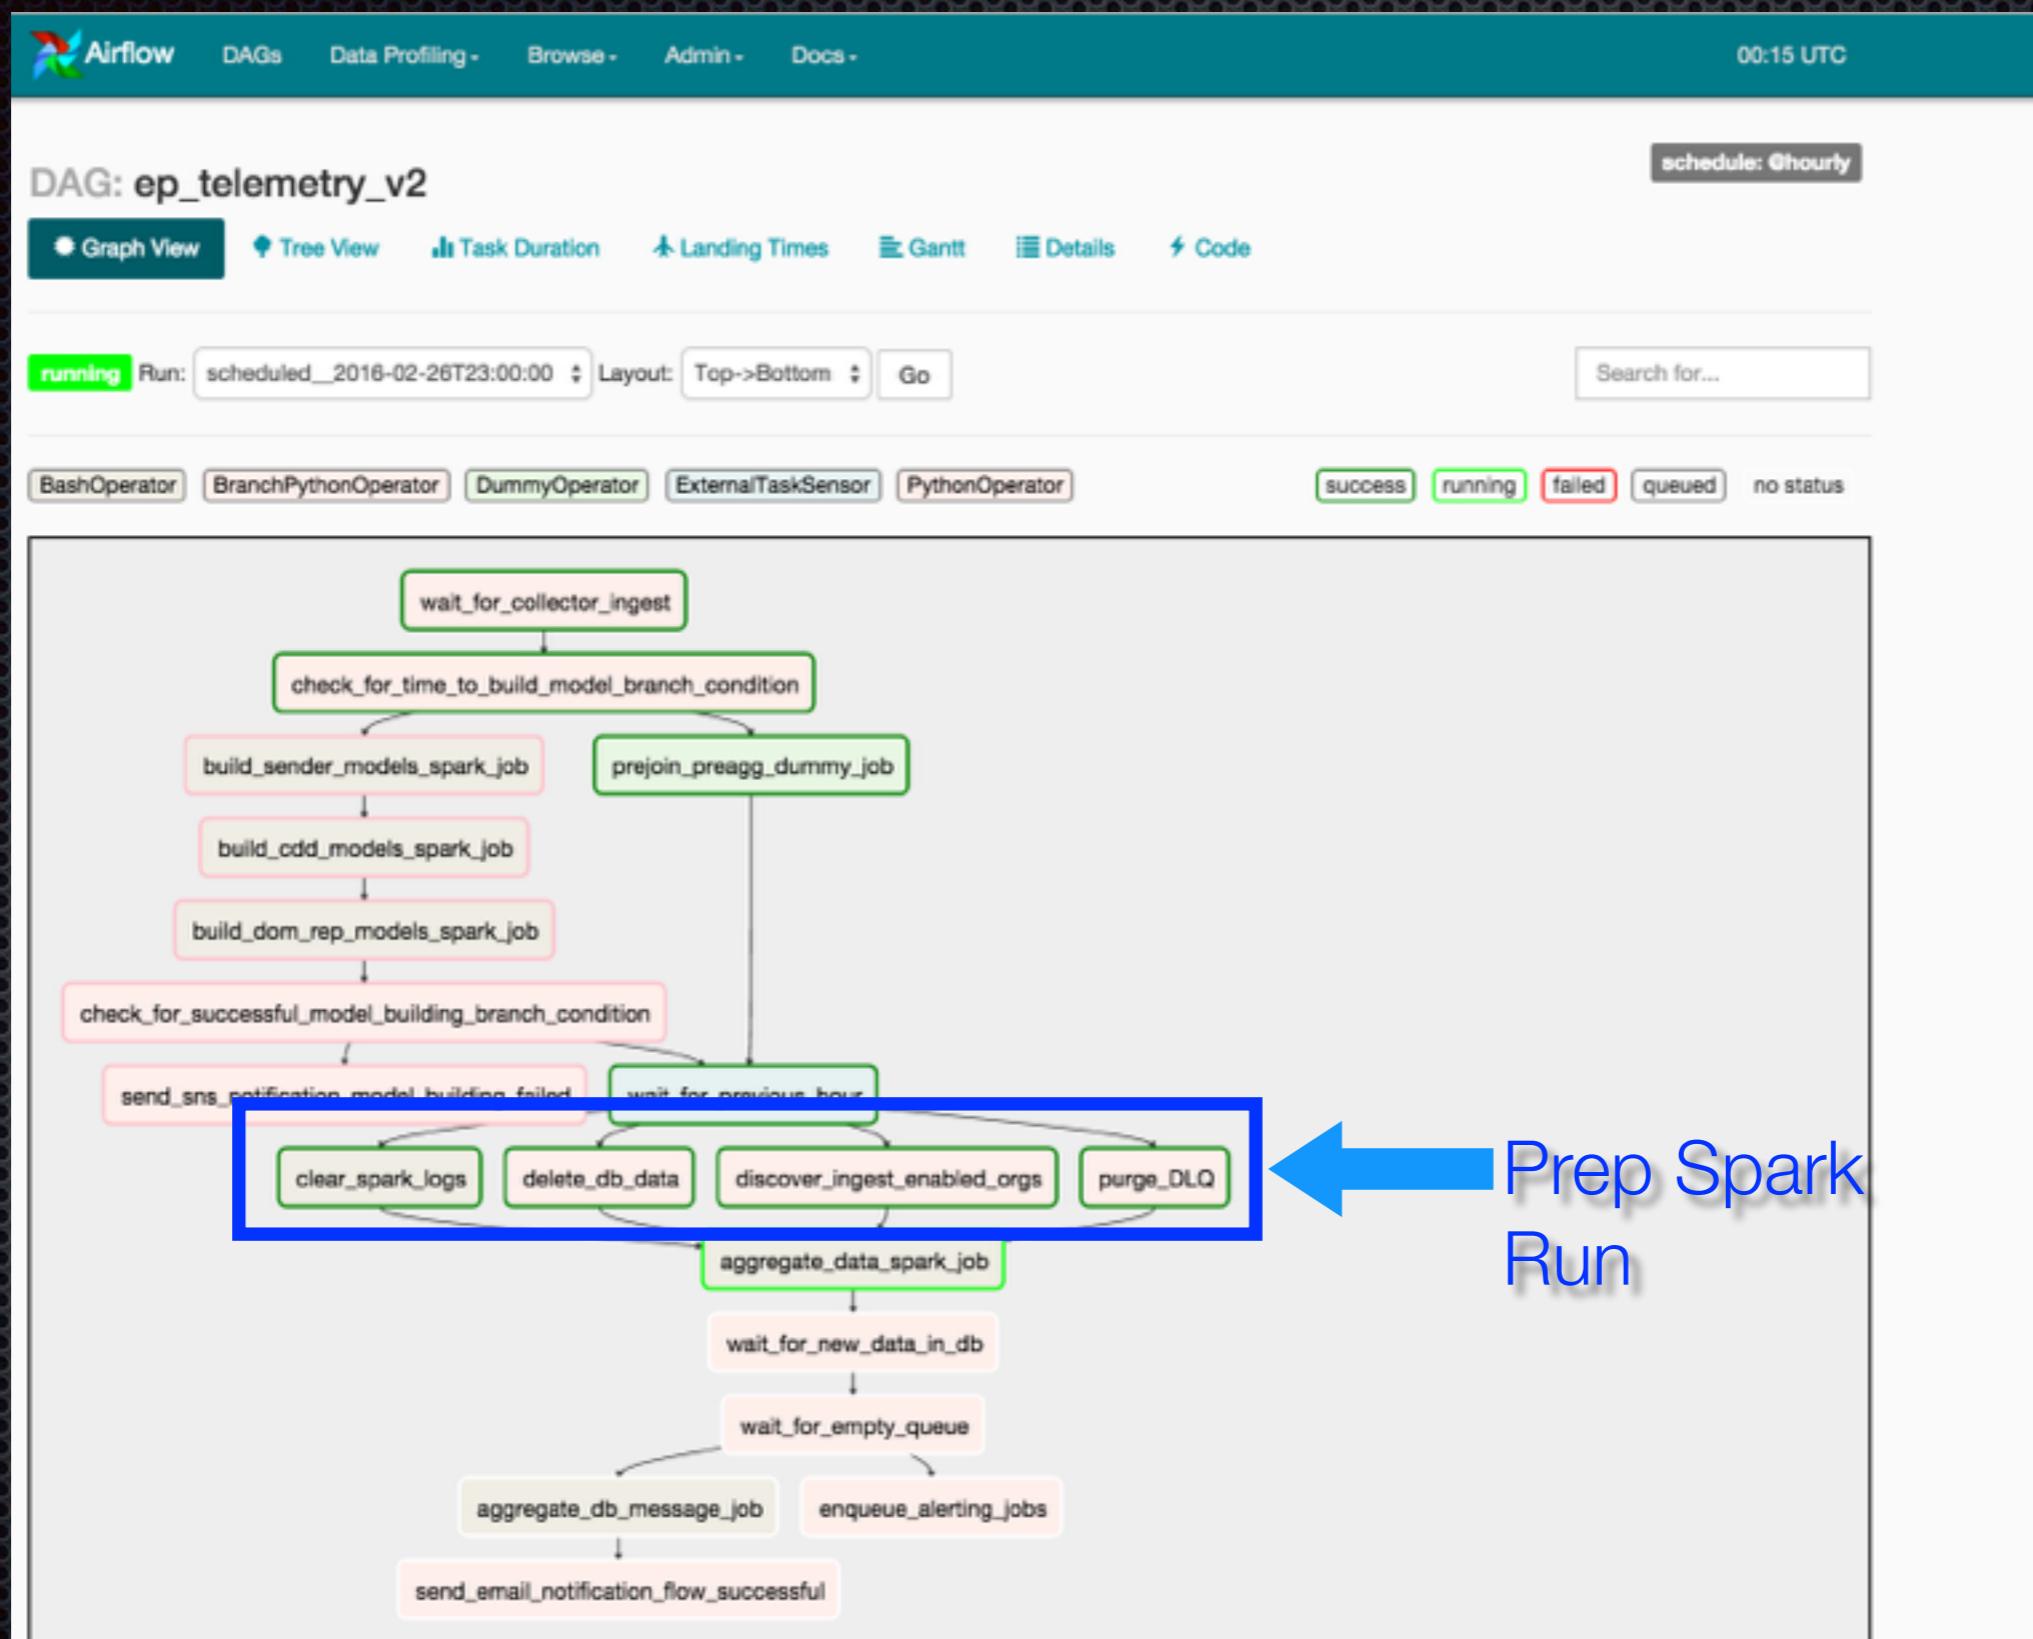
# Airflow DAG



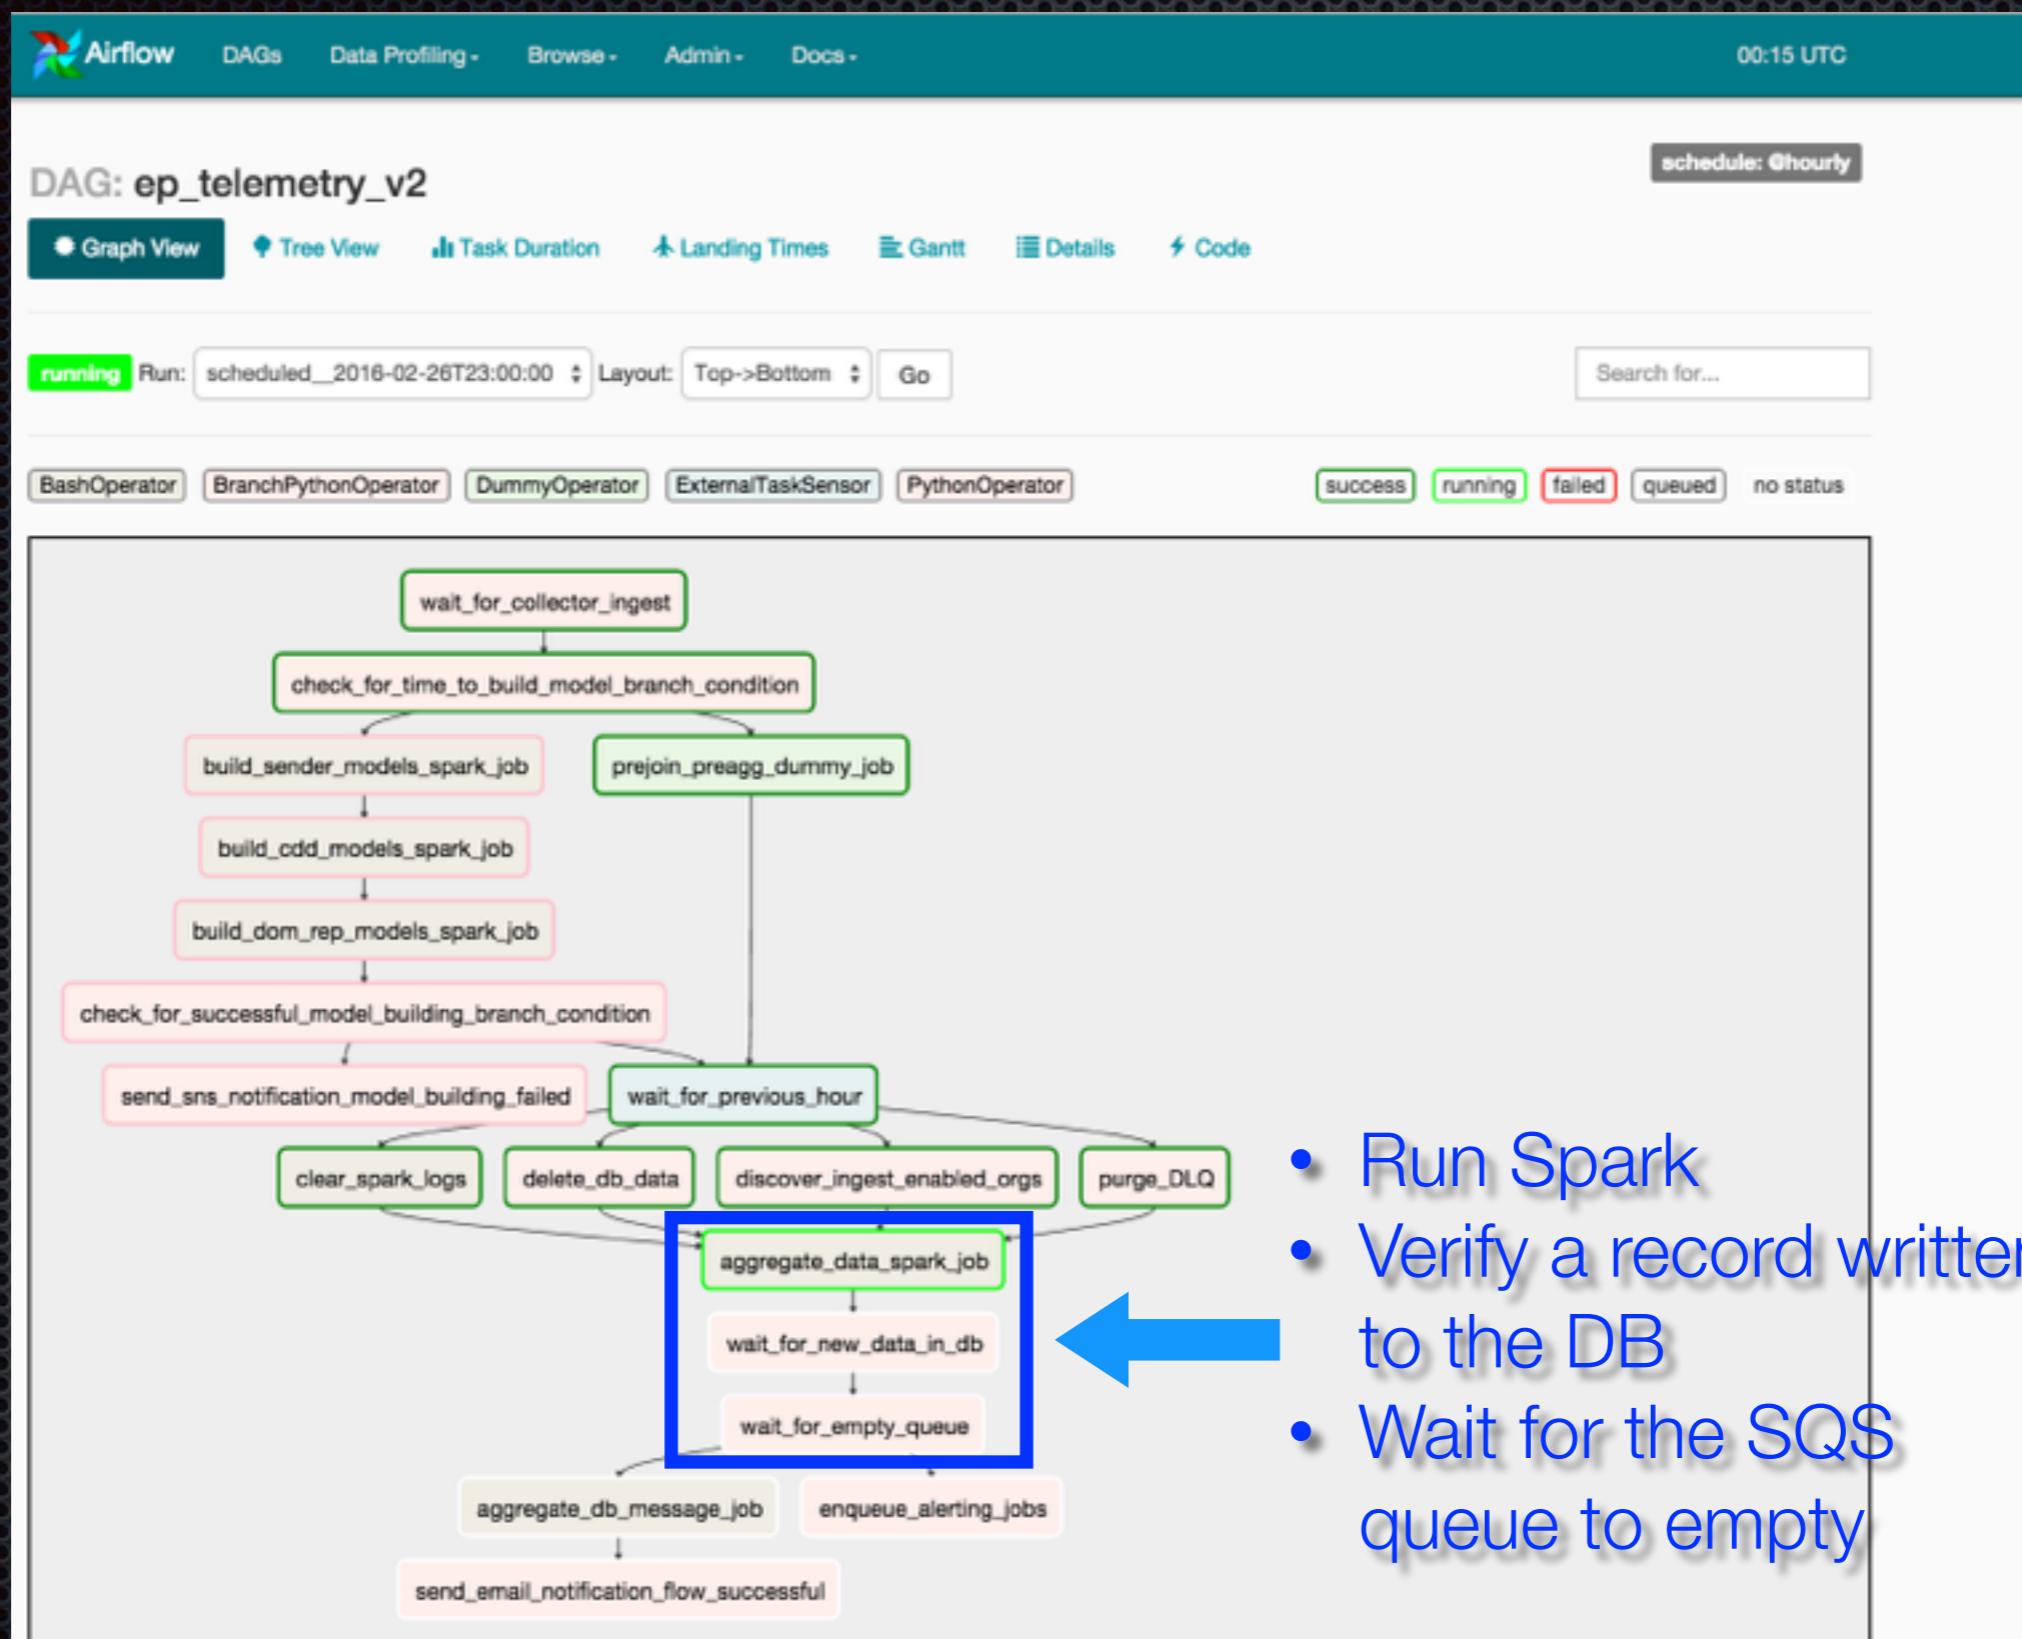
# Airflow DAG



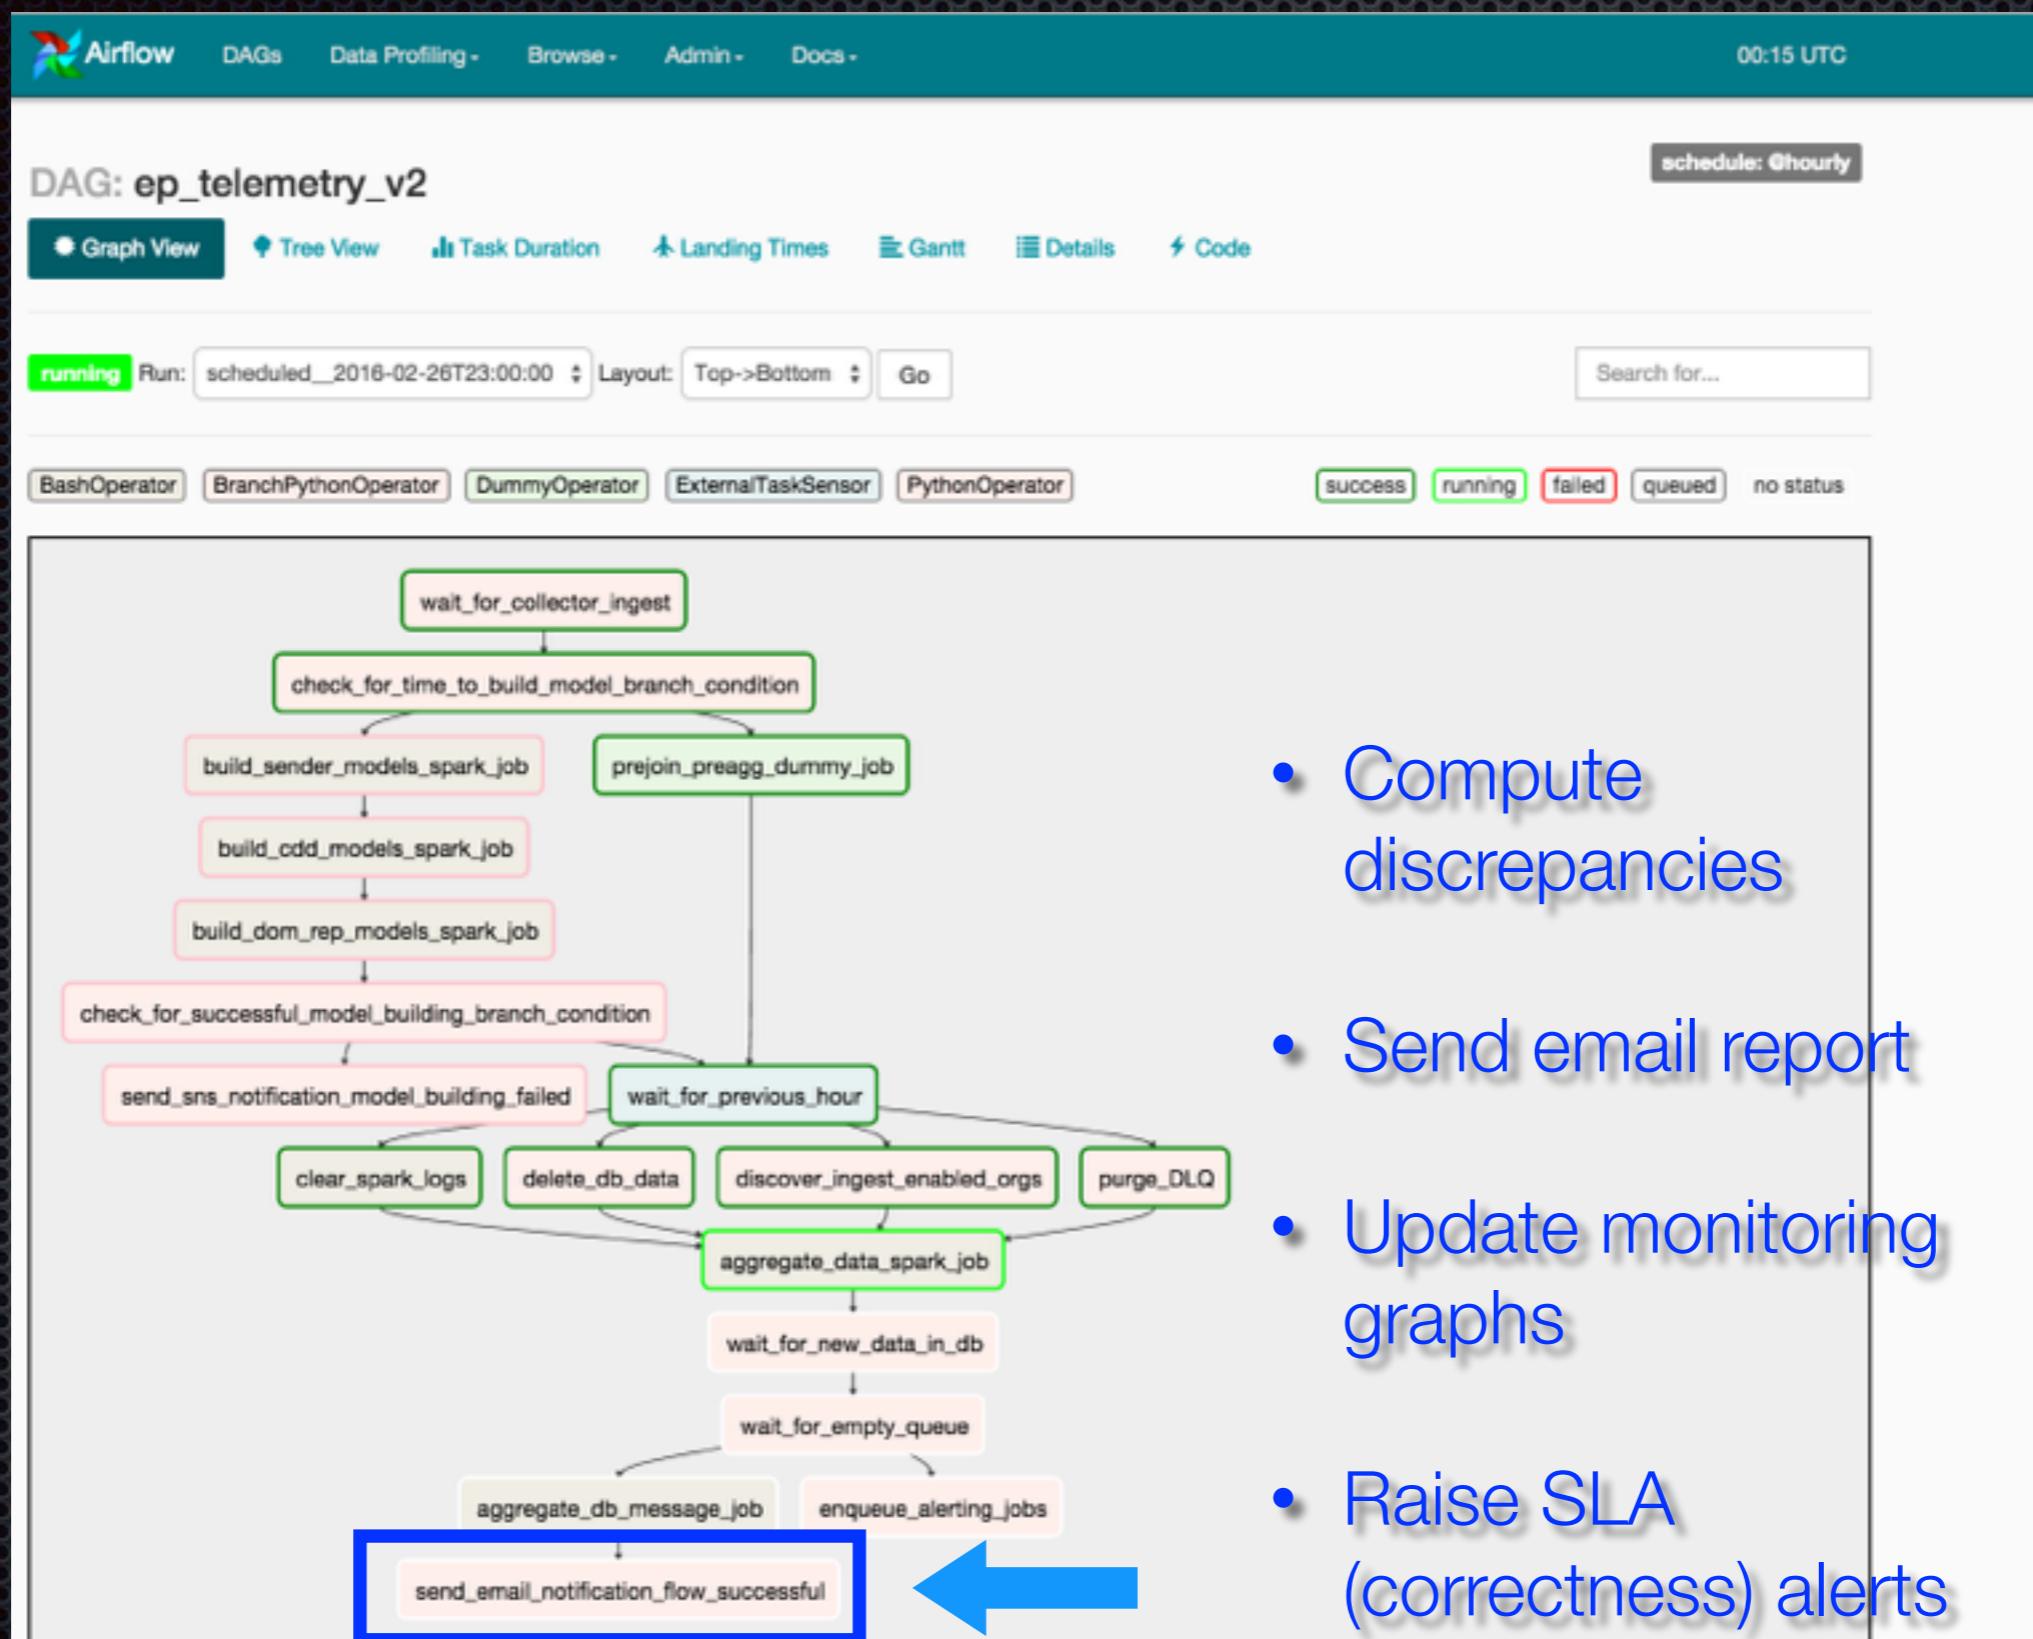
# Airflow DAG



# Airflow DAG



# Airflow DAG



# SLAs & Insights



# Desirable Qualities of a Resilient Data Pipeline

## Correctness

- Data Integrity (no loss, etc...)
- Expected data distributions

## Operability

- Fine-grained Monitoring & Alerting of Correctness & Timeliness SLAs
- Quick Recoverability

## Timeliness

- All output within time-bound SLAs (e.g. 1 hour)

## Cost

- Pay-as-you-go

# Desirable Qualities of a Resilient Data Pipeline

Correctness

- Data Integrity (no loss, etc...)
- Expected data distributions

**SLA**

Operability

- Fine-grained Monitoring & Alerting of Correctness & Timeliness SLAs
- Quick Recoverability

Timeliness

- All output within time-bound SLAs (e.g. 1 hour)

**SLA**

Cost

- Pay-as-you-go

# Correctness : Email Reporting

airflow@agari.com

to AirflowJobComp.

6:20 PM (16 hours ago)

Hi EP Folks! The EP Data Pipeline (Airflow) Loaded Data in the EP\_STAGE environment for day: 2016-04-28 12:00:00 GMT/UTC  
[Airflow](#)

Day	RDAs	Messages
2016-04-28 00:00:00	22058	261254

## Message Counts by Org and by Pipeline Stage Output

Org_ID	Collector Msgs	Agg Msgs	DB Msgs	Discrepancy	Discr %	NRT DB Msgs	NRT Discrepancy	NRT Discr %	Skipped	Notes
1	340	266	266	0	0.0	340	-74	-21.8	74	
5	43115	43115	43115	0	0.0	358	42757	99.2	0	
6	156	153	153	0	0.0	0	153	98.1	3	
7	73	68	68	0	0.0	0	68	93.2	5	
12	315	111	111	0	0.0	0	111	35.2	204	
14	35648	28731	28731	0	0.0	0	28731	80.6	6917	
16	35383	34010	34010	0	0.0	0	34010	96.1	1373	
25	20530	20530	20530	0	0.0	0	20530	100.0	0	
54	30003	26190	26190	0	0.0	0	26190	87.3	3813	
56	108084	108080	108080	0	0.0	0	108080	100.0	4	

orgs



For each org, we check for duplicate or missing data as a count & percentage

# Correctness : Email Reporting

airflow@agari.com

to AirflowJobComp.

6:20 PM (16 hours ago)

Hi EP Folks! The EP Data Pipeline (Airflow) Loaded Data in the EP\_STAGE environment for day: 2016-04-28 12:00:00 GMT/UTC  
[Airflow](#)

Day	RDAs	Messages
2016-04-28 00:00:00	22058	261254

## Message Counts by Org and by Pipeline Stage Output

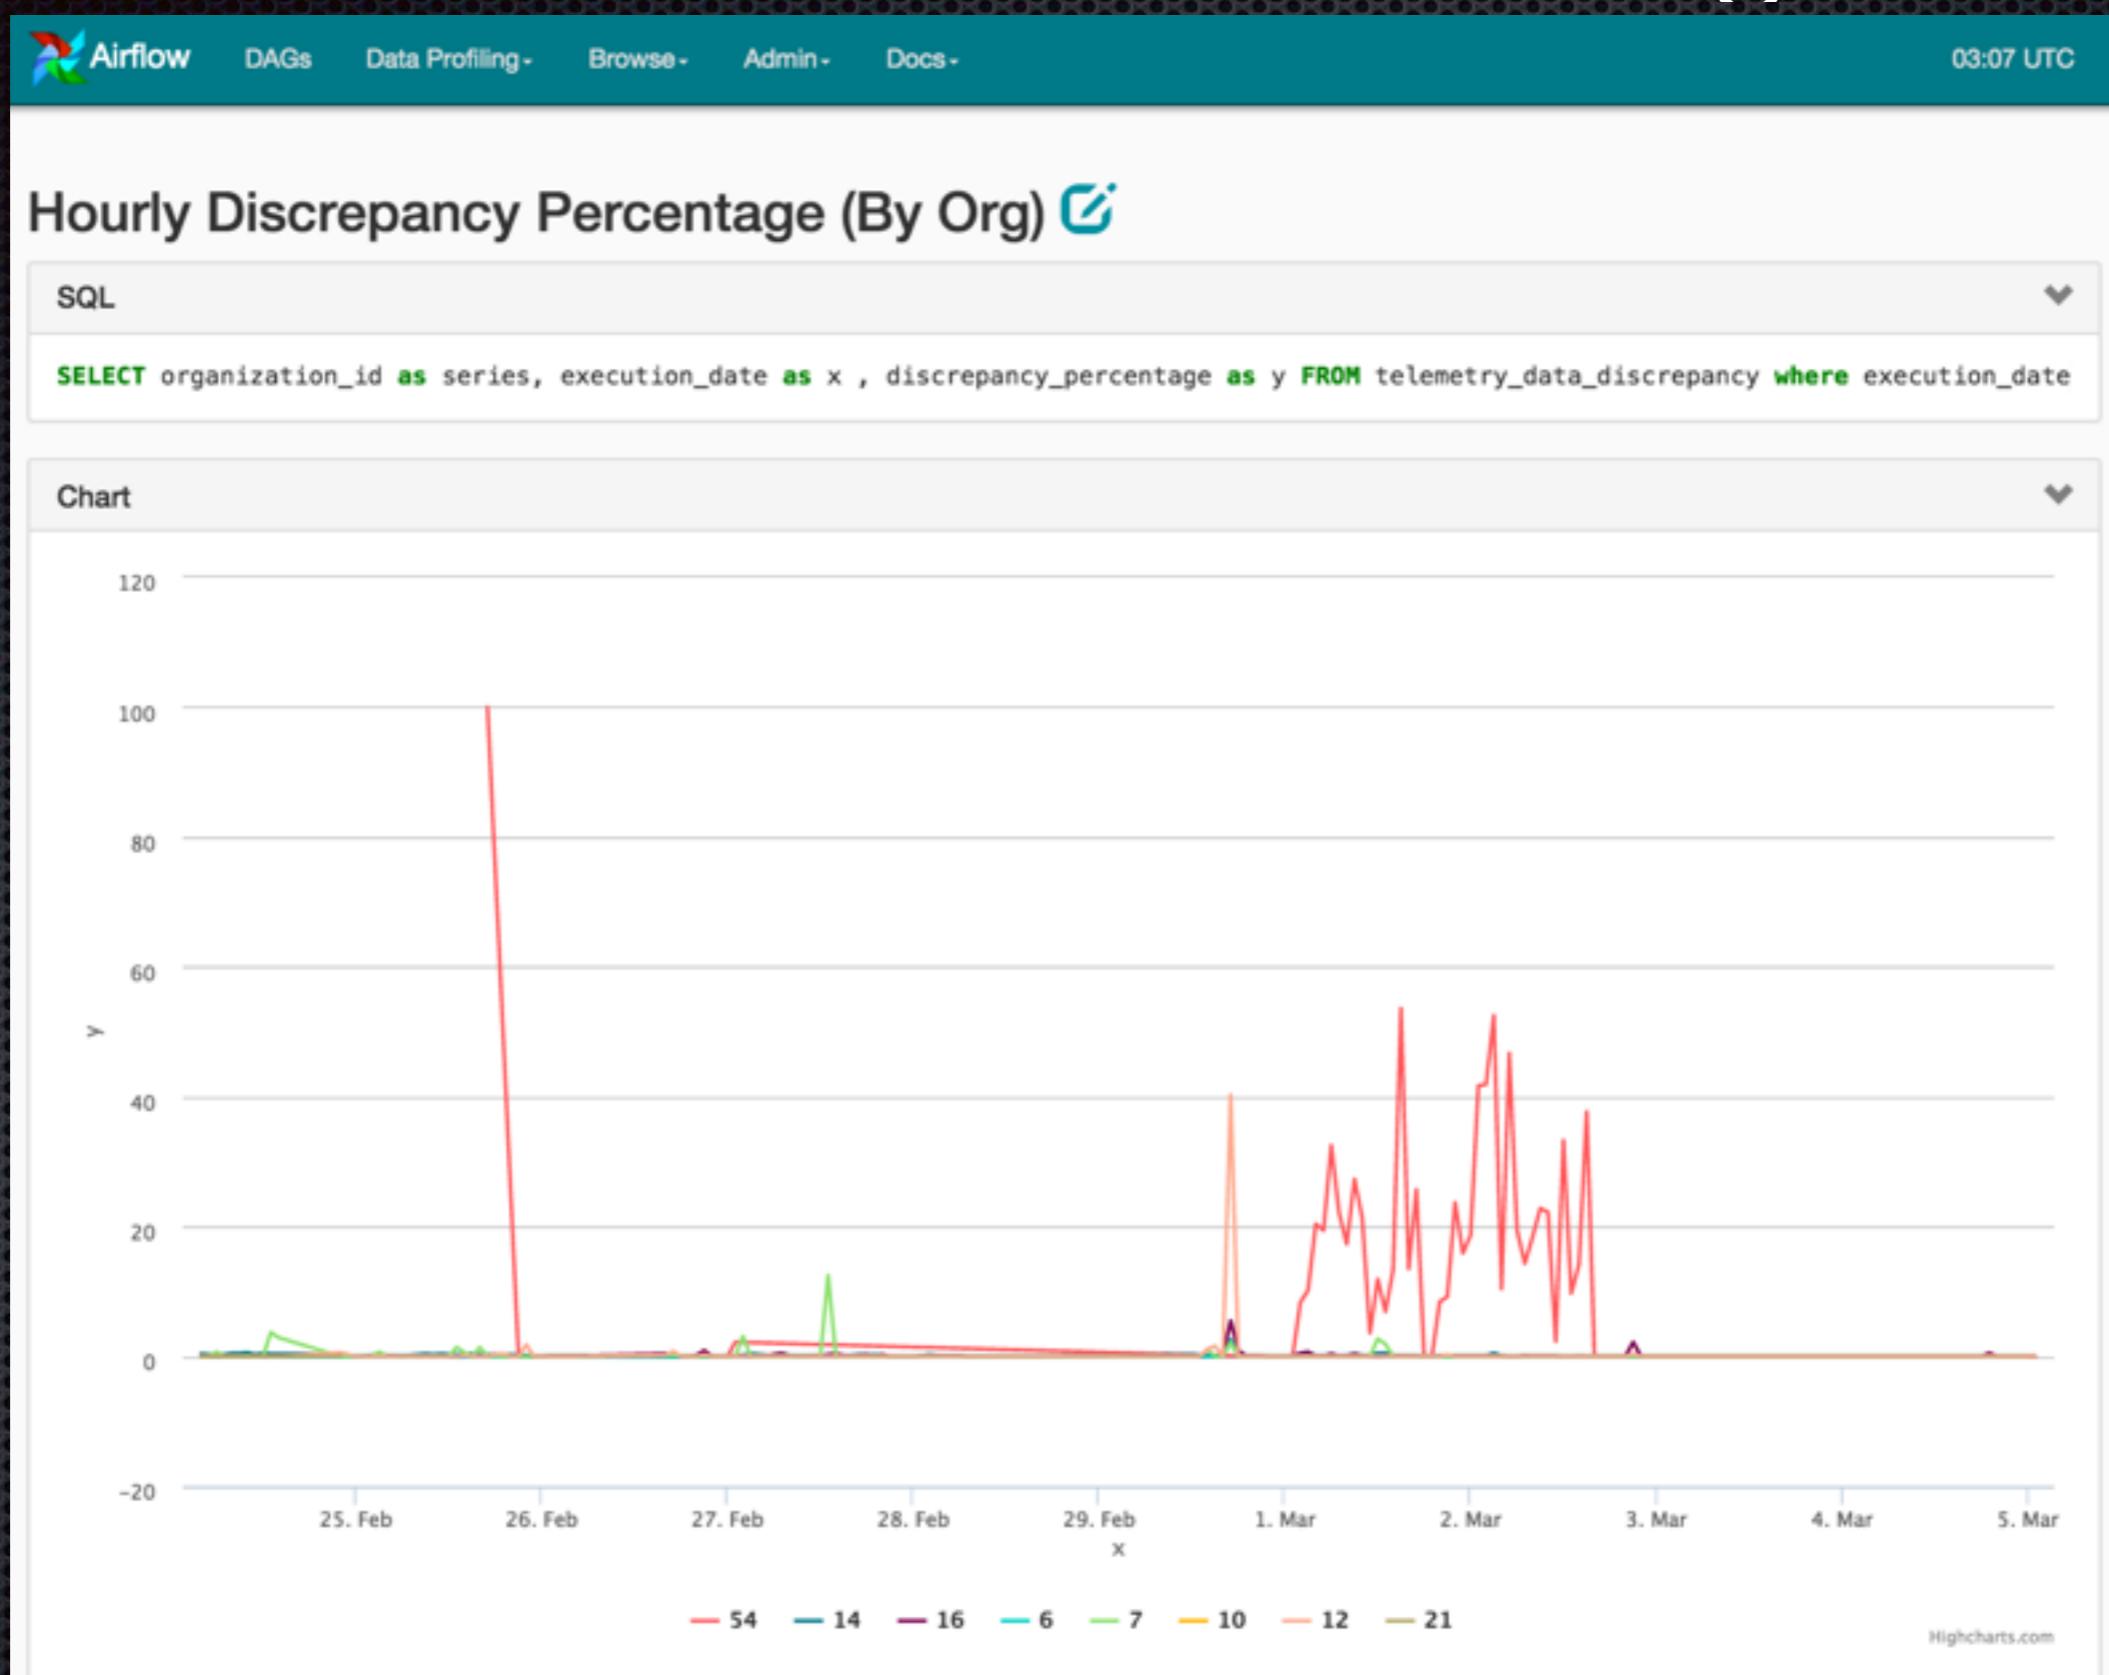
Org_ID	Collector Msgs	Agg Msgs	DB Msgs	Discrepancy	Discr %	NRT DB Msgs	NRT Discrepancy	NRT Discr %	Skipped	Notes
1	340	266	266	0	0.0	340	-74	-21.8	74	
5	43115	43115	43115	0	0.0	358	42757	99.2	0	
6	156	153	153	0	0.0	0	153	98.1	3	
7	73	68	68	0	0.0	0	68	93.2	5	
12	315	111	111	0	0.0	0	111	35.2	204	
14	35648	28731	28731	0	0.0	0	28731	80.6	6917	
16	35383	34010	34010	0	0.0	0	34010	96.1	1373	
25	20530	20530	20530	0	0.0	0	20530	100.0	0	
54	30003	26190	26190	0	0.0	0	26190	87.3	3813	
56	108084	108080	108080	0	0.0	0	108080	100.0	4	

↑  
↑  
↑  
orgs



These are the 3 stages of the pipeline. We can detect where a discrepancy is coming from - often related to a code push!

# Correctness : Monitoring



# Correctness & Timeliness : Alerting

The screenshot shows a Slack interface with the following details:

- Left Sidebar:** Shows the team "Agari" and channels: sid, STARRED, # airflow, # analysis, # eng, # eng-ep, **# ep-ops** (highlighted in green), # ep-real-time-alerting, and @ scotfree, kevin, wforr...
- Header:** Slack logo, user profile, channel name "#ep-ops", channel description "A channel with info that can help you resolve VictorOps alerts", member count "13", and search bar.
- Today's Messages:**
  - Airflow BOT** 10:00 AM: EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful](#) on 2016-02-26T18:00:00
  - sid** 10:02 AM: So, we need to catch up here.. any reason we don't want to just mark success for the many hours it is behind?
  - Airflow BOT** 11:00 AM: EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful](#) on 2016-02-26T19:00:00
- February 2nd:**
  - Airflow BOT** 5:27 PM: ep\_telemetry\_v2 on [etl-00.ep-old.prod.agari.com](#) completed 2016-02-03 00:00:00 with High Discrepancies
- February 3rd:**
  - Airflow BOT** 8:48 PM: ep\_telemetry\_v2 on [workflow-00.ep.stage.agari.com](#) completed 2016-02-26 03:00:00 with 1 DLQs : Sample Exception

# Correctness & Timeliness : Alerting

Timeliness SLA miss

The screenshot shows a Slack interface for the '#ep-ops' channel. A blue double-headed arrow points from the text 'Timeliness SLA miss' to the channel sidebar, which lists several channels including '# ep-ops' (highlighted in green) and '# airflow'. The main message list shows the following:

- Airflow BOT 10:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful on 2016-02-26T18:00:00](#)
- sid 10:02 AM So, we need to catch up here.. any reason we don't want to just mark success for the many hours it is behind?
- Airflow BOT 11:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful on 2016-02-26T19:00:00](#)
- Airflow BOT 5:27 PM ep\_telemetry\_v2 on [etl-00.ep-old.prod.agari.com](#) completed [2016-02-03 00:00:00](#) with High Discrepancies
- Airflow BOT 8:48 PM ep\_telemetry\_v2 on [workflow-00.ep.stage.agari.com](#) completed [2016-02-26 03:00:00](#) with 1 DLQs : Sample Exception

Today

# Correctness & Timeliness : Alerting

Timeliness SLA miss

Agari sid

STARRED # airflow # analysis # eng # eng-ep # ep-ops # ep-real-time-testing @ scotfree, kiran, wforr

Slack

#ep-ops A channel with info that can help you resolve VictorOps alerts 13 Search

Today

Airflow BOT 10:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful on 2016-02-26T18:00:00](#)

sid 10:02 AM So, we need to catch up here.. any reason we don't want to just mark success for the many hours it is behind?

Airflow BOT 11:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful on 2016-02-26T19:00:00](#)

Airflow BOT 5:27 PM ep\_telemetry\_v2 on [etl-00.ep-old.prod.agari.com](#) completed [2016-02-03 00:00:00](#) with High Discrepancies

Airflow BOT 8:48 PM ep\_telemetry\_v2 on [workflow-00.ep.stage.agari.com](#) completed [2016-02-26 03:00:00](#) with 1 DLQs : Sample Exception

February 3rd

```
dag = DAG(DAG_NAME,  
          schedule_interval='@hourly',  
          default_args=default_args,  
          sla_miss_callback=sla_alert_func)
```

# Correctness & Timeliness : Alerting

Timeliness SLA miss

Correctness SLA miss

Slack

#ep-ops A channel with info that can help you resolve VictorOps alerts

Today

Airflow BOT 10:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful on 2016-02-26T18:00:00](#)

sid 10:02 AM So, we need to catch up here.. any reason we don't want to just mark success for the many hours it is behind?

Airflow BOT 11:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful on 2016-02-26T19:00:00](#)

Airflow BOT 5:27 PM ep\_telemetry\_v2 on [etl-00.ep-old.prod.agari.com](#) completed [2016-02-03 00:00:00](#) with High Discrepancies

Airflow BOT 8:48 PM ep\_telemetry\_v2 on [workflow-00.ep.stage.agari.com](#) completed [2016-02-26 03:00:00](#) with 1 DLQs : Sample Exception

# Correctness & Timeliness : Alerting

Slack

#ep-ops A channel with info that can help you resolve VictorOps alerts

Today

Airflow BOT 10:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful](#) on 2016-02-26T18:00:00

sid 10:02 AM So, we need to catch up here.. any reason we don't want to just mark success for the many hours it is behind?

Airflow BOT 11:00 AM EP\_STAGE - ep\_telemetry\_v2 SLA Miss for task [send\\_email\\_notification\\_flow\\_successful](#) on 2016-02-26T19:00:00

Airflow BOT 5:27 PM ep\_telemetry\_v2 on [etl-00.ep-old.prod.agari.com](#) completed 2016-02-03 00:00:00 with High Discrepancies

February 3rd

Airflow BOT 8:48 PM ep\_telemetry\_v2 on [workflow-00.ep.stage.agari.com](#) completed 2016-02-26 03:00:00 with 1 DLQs : Sample Exception

Agari Timeline Settings Reports

Teams Users

On-call Engaged

Teammates All Users

Andrew Flury @affury Enterprise Protect everyone

Christopher Haag @chaaq Enterprise Protect everyone

Spencer Sun @Hsun everyone healthchecks.io

Triggered: 0

Acked: 0

Resolved

#102 ALERT - Prod EP Pipeline Discrepancy workflow-00.ep.prod.agari.com/agari.log/Ahari Data.Inc

Time	Feb 26, 2016 13:00:00
Email	ALERT - Prod EP ...
State	OK
Resolved By	affury

#101 airflow/telemetry/SLA miss

Time	Feb 26, 2016 13:00:00
API	airflow/telemetry...
State	INFO
Resolved By	SYSTEM

Timeliness &  
Correctness SLA  
misses sent to  
PagerDuty/VictorOps

# Operations

## Managing Airflow

# Agari in AWS

- Agari runs in the AWS US-West-2 Region & leverages multiple Availability Zones (AZs) for Fault Tolerance
- We have 3 environments (Dev, Stage, Prod), each with its own AWS account

Dev

- A Developer's playground
  - developers prototype new services here

# Agari in AWS

- Agari runs in the AWS US-West-2 Region & leverages multiple Availability Zones (AZs) for Fault Tolerance
- We have 3 environments (Dev, Stage, Prod), each with its own AWS account



- A Developer's playground
  - developers prototype new services here



- A Pre-Production Env
  - All new code & system upgrades are baked here prior to Prod deployment
  - It gets a copy of prod data & stage data : helps us test for scale!

# Agari in AWS

- Agari runs in the AWS US-West-2 Region & leverages multiple Availability Zones (AZs) for Fault Tolerance
- We have 3 environments (Dev, Stage, Prod), each with its own AWS account

Dev

- A Developer's playground
  - developers prototype new services here

Stage

- A Pre-Production Env
  - All new code & system upgrades are baked here prior to Prod deployment
  - It gets a copy of prod data & stage data : helps us test for scale!

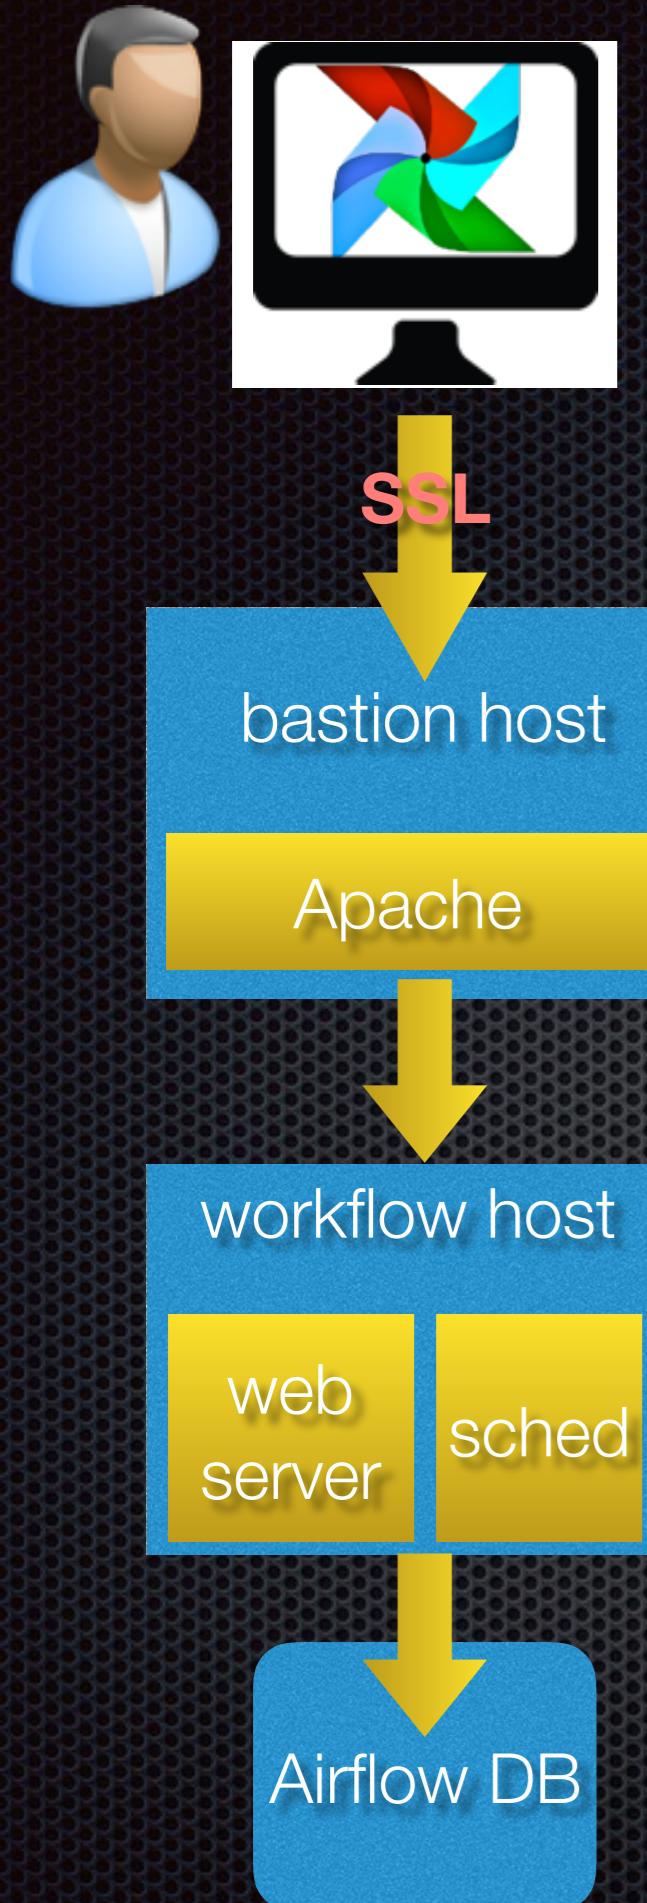
Prod

- Prod Env

# Agari in AWS

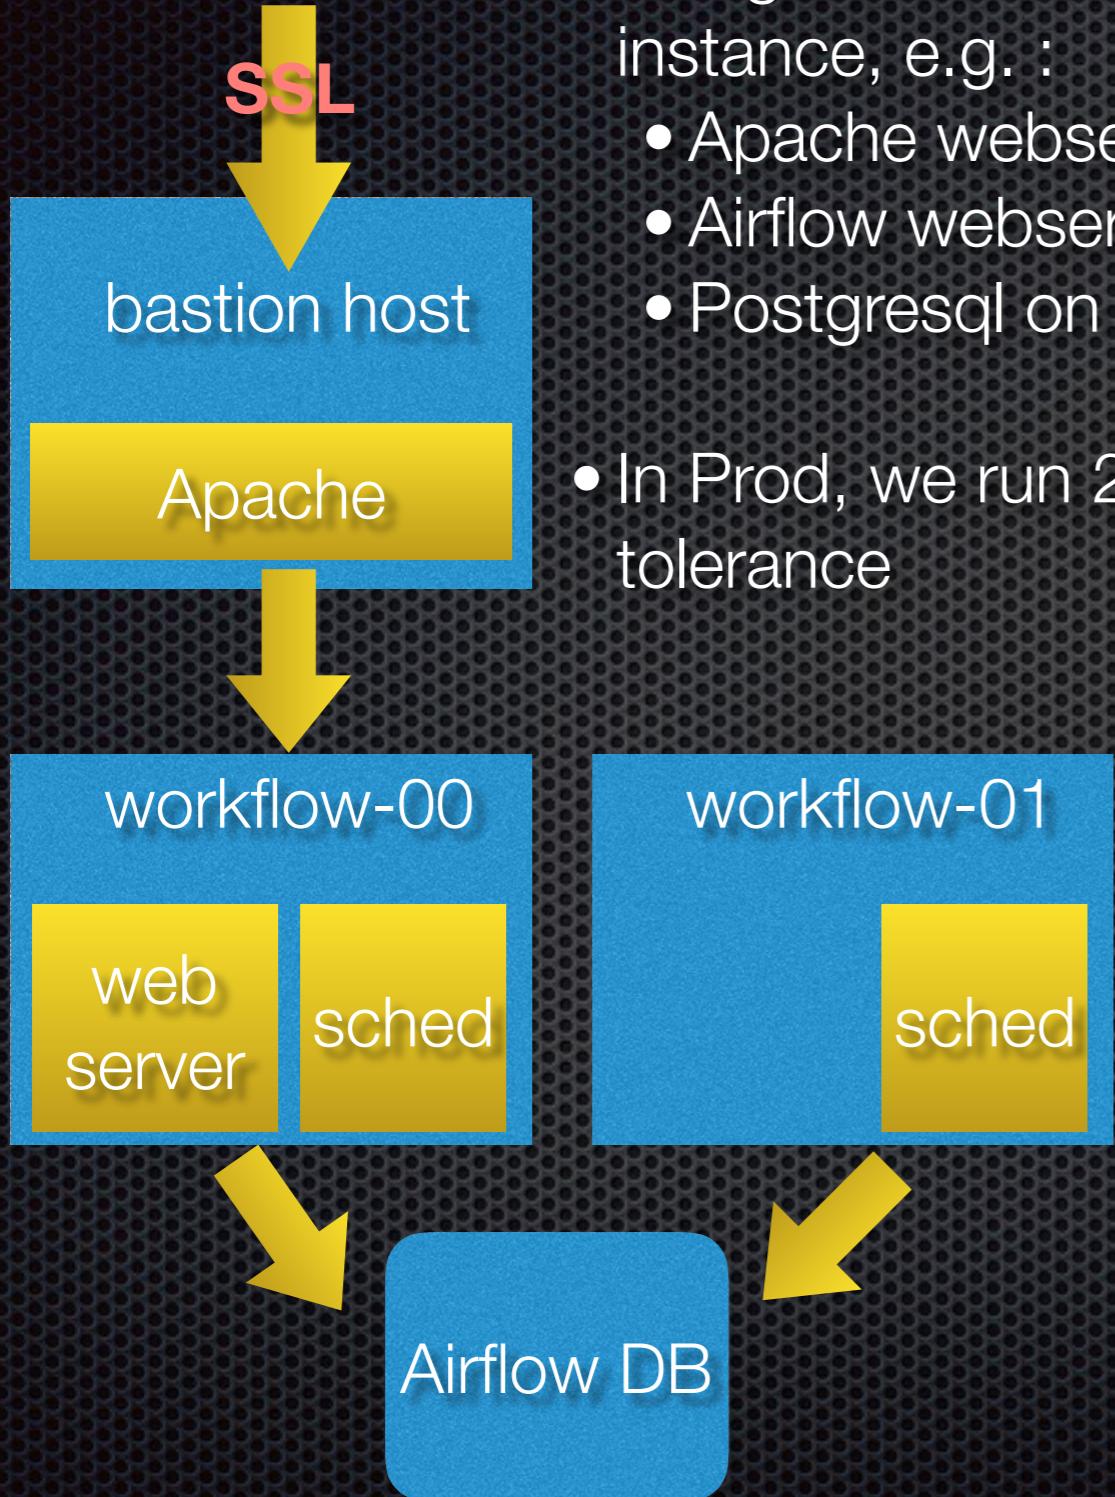
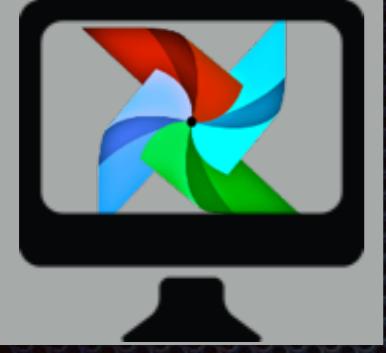
## Security

- Each environment has its own AWS account for isolation
- We use **Security Groups** & **VPCs** to lock every EC2 instance down
  - In other words, EC2 nodes only accept inbound traffic from other hosts in the same or peered VPCs
- Exception : Each env has one **bastion** host
  - The **bastion** host provides indirect access to the nodes in the VPC



# Agari's Airflow in AWS

- An Apache Reverse-proxy, running on the bastion host, does:
  - SSH termination &
  - Authentication via mod\_auth
- The workflow machine hosts
  - airflow webserver
  - airflow scheduler (running LocalExecutor)
- A Postgresql DB instance hosts the Airflow database



# Fault Tolerance

- At Agari, we use **Monit** to keep processes running on any EC2 instance, e.g. :
  - Apache webserver on the **bastion** host
  - Airflow webserver & scheduler on the **workflow** host
  - Postgresql on the **DB** host
- In Prod, we run 2 schedulers for increased scalability and fault-tolerance

# Airflow Github Repos

Apache/  
Airflow

r39132/  
Airflow  
(fork)

Agari/  
Airflow  
(fork)

# Airflow Github Repos

Apache/  
Airflow

r39132/  
Airflow  
(fork)

Agari/  
Airflow  
(fork)



Deployed to  
stage & prod

# Airflow Github Repos

Apache/  
Airflow

r39132/  
Airflow  
(fork)

Agari/  
Airflow  
(fork)

Agari / Airflow pinned  
to a release version  
(e.g. 1.6.2)



# Airflow Github Repos

Apache/  
Airflow

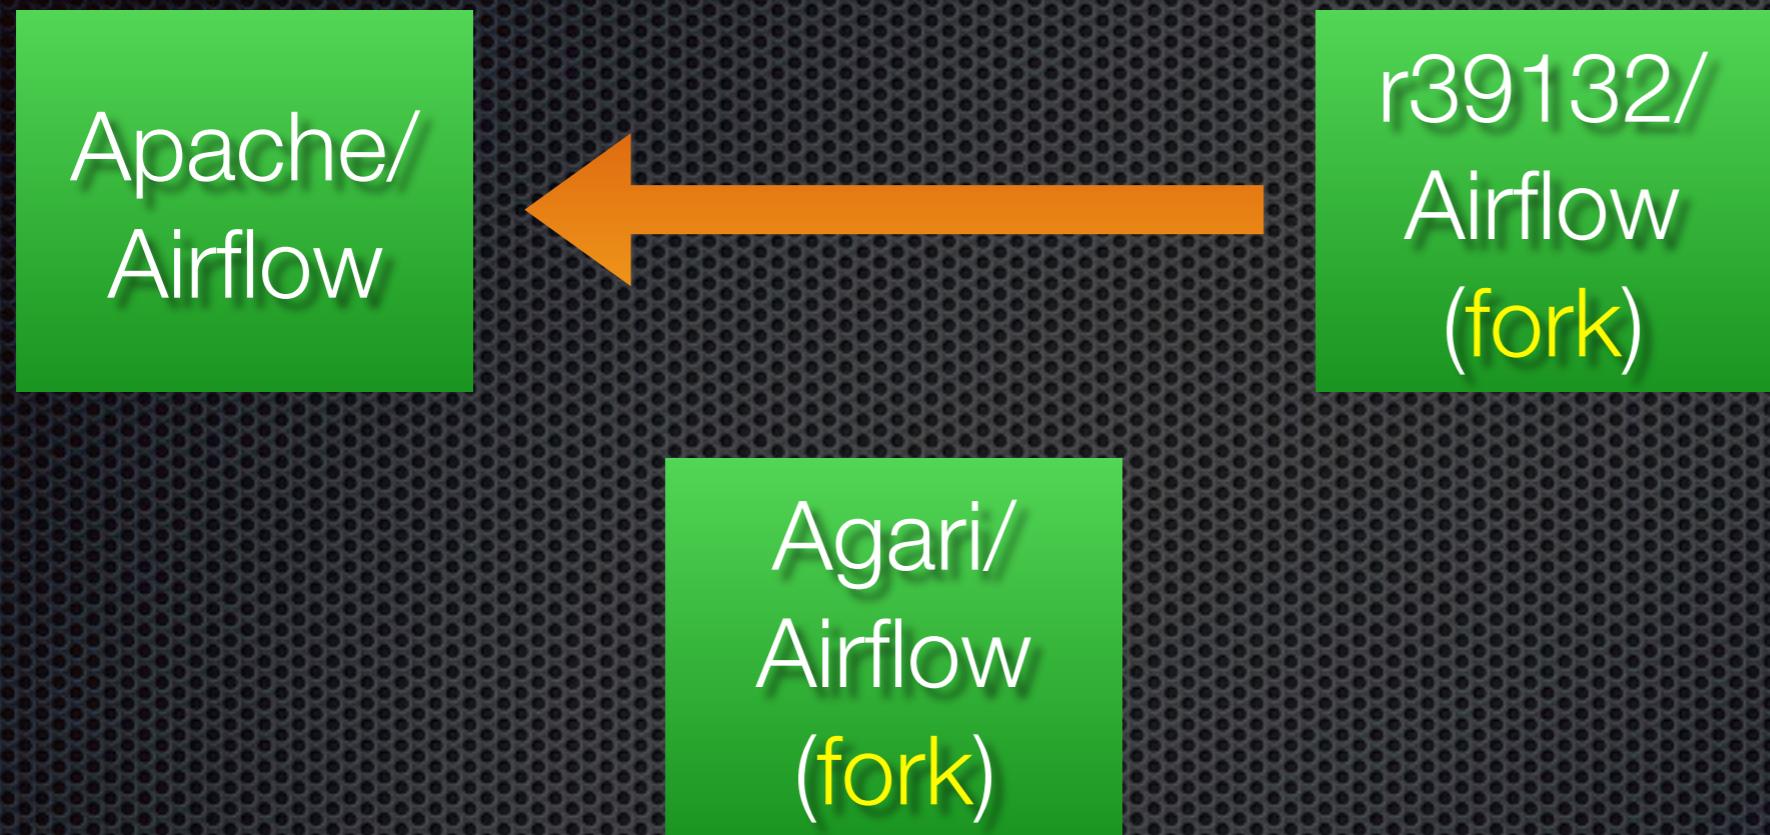


r39132 / Airflow kept in sync  
upstream master

r39132/  
Airflow  
(fork)

Agari/  
Airflow  
(fork)

# Airflow Github Repos



# Airflow Github Repos

Apache/  
Airflow

r39132/  
Airflow  
(fork)

Agari/  
Airflow  
(fork)



PR 191 is cherry-  
picked into Agari /  
Airflow

# Airflow Github Repos

Apache/  
Airflow

r39132/  
Airflow  
(fork)

Agari/  
Airflow  
(fork)



Deployed to  
stage & prod

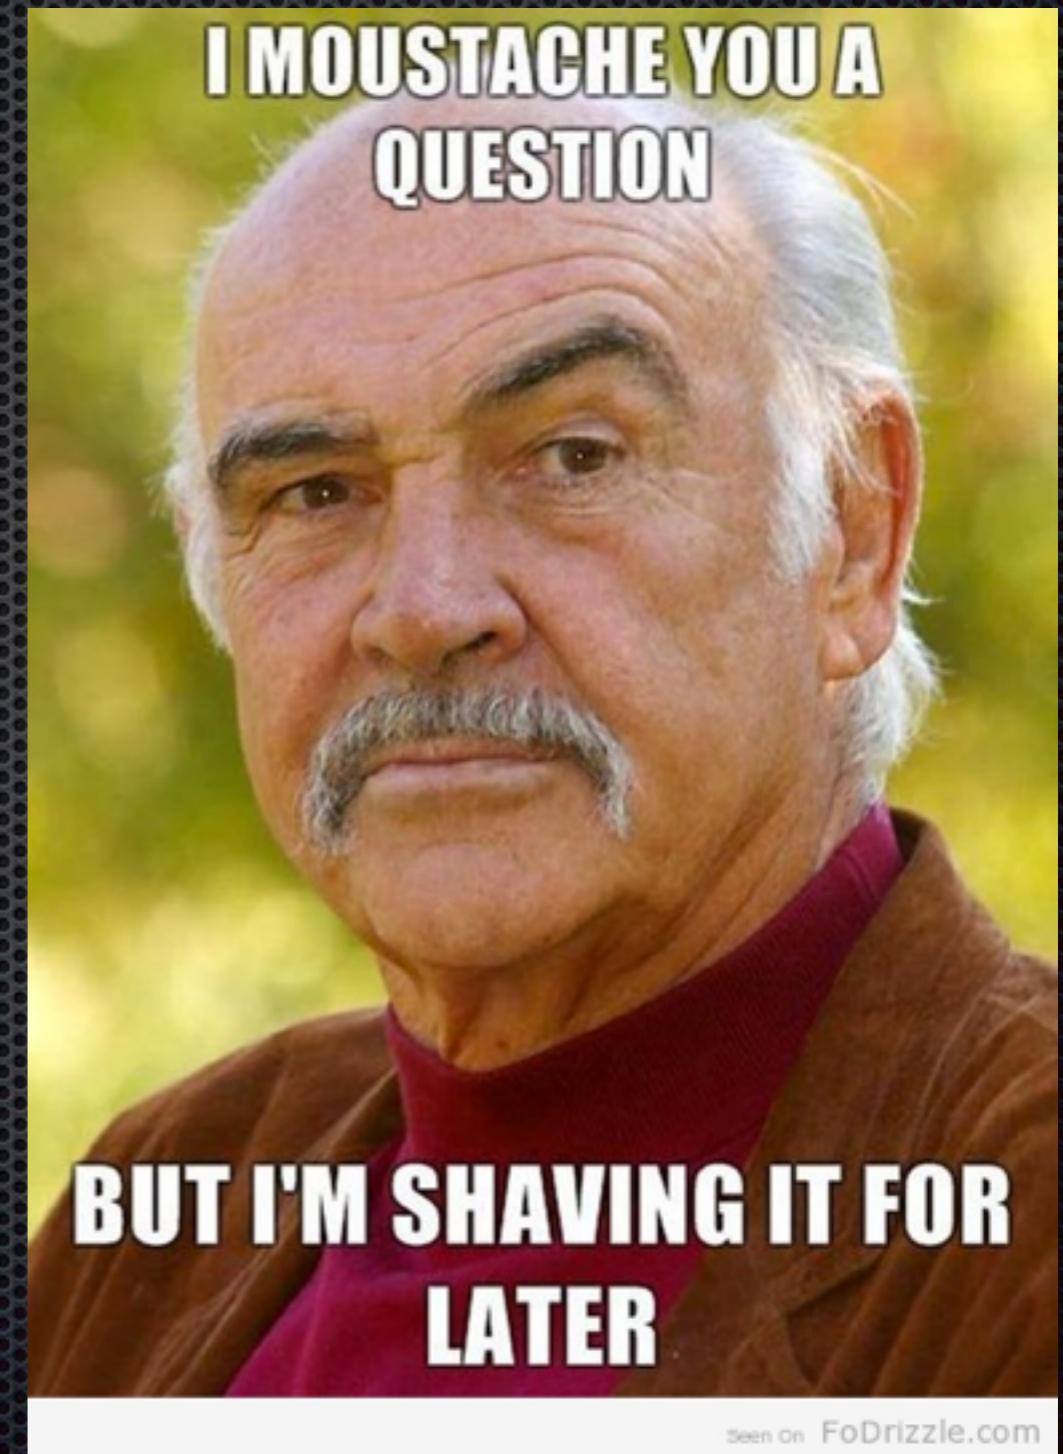
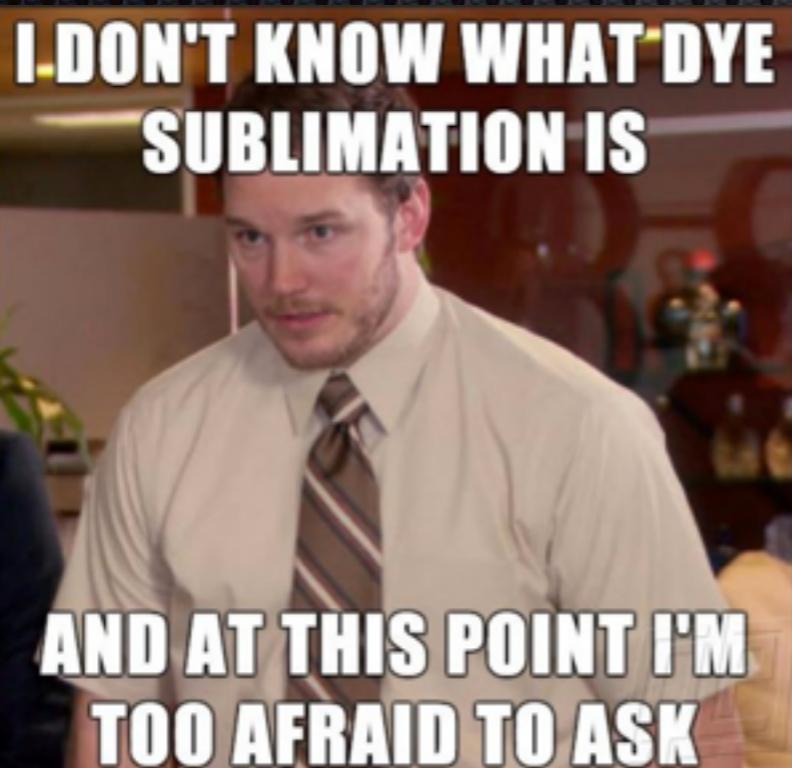
What is deployed is 1.6.2 + cherry-picked  
(e.g. PR 191) changes

# Acknowledgments

None of this work would be possible without the contributions of the strong team below

- Vidur Apparao
- Stephen Cattaneo
- Jon Chase
- Andrew Flury
- William Forrester
- Chris Haag
- Mike Jones
- Scot Kennedy
- Thede Loder
- Paul Lorence
- Kevin Mandich
- Gabriel Ortiz
- Jacob Rideout
- Josh Yang
- Julian Mehnle

# Questions? (@r39132)



Seen On FoDrizzle.com