



Running Apache Airflow Workflows as ETL Processes on Hadoop

By: Robert Sanders

CLAIRVOYANT

Agenda

- What is Apache Airflow?
 - Features
 - Architecture
 - Terminology
 - Operator Types
- ETL Best Practices
 - How they're supported in Apache Airflow
- Executing Airflow Workflows on Hadoop
- Use Cases
- Q&A

Robert Sanders



- Big Data Manager, Engineer, Architect, etc.
- Work for **Clairvoyant LLC**
- 5+ Years of Big Data Experience
- Email: robert.sanders@clairvoyantsoft.com
- LinkedIn: <https://www.linkedin.com/in/robert-sanders-61446732>
- Slide Share: <http://www.slideshare.net/RobertSanders49>

What's the problem?

- As a Big Data Engineer you work to create jobs that will perform various operations
 - Ingest data from external data sources
 - Transformation of Data
 - Run Predictions
 - Export data
 - Etc.
- You need to have some mechanism to schedule and run these jobs
 - Cron
 - Oozie
- Existing Scheduling Services have a number of limitations that make them difficult to work with and not usable in all instances

What is Apache Airflow?


- Airflow is an Open Source platform to programmatically author, schedule and monitor workflows
 - Workflows as Code
 - Schedules Jobs through Cron Expressions
 - Provides monitoring tools like alerts and a web interface
- Written in Python
 - As well as user defined Workflows and Plugins
- Was started in the fall of 2014 by Maxime Beauchemin at Airbnb
- Apache Incubator Project
 - Joined Apache Foundation in early 2016
 - <https://github.com/apache/incubator-airflow/>

Why use Apache Airflow?

- Define Workflows as Code
 - Makes workflows more maintainable, versionable, and testable
 - More flexible execution and workflow generation
- Lots of Features
- Feature Rich Web Interface
- Worker Processes Scale Horizontally and Vertically
 - Can be a cluster or single node setup
- Lightweight Workflow Platform


Apache Airflow Features (Some of them)


- Automatic Retries
- SLA monitoring/alerting
- Complex dependency rules: branching, joining, sub-workflows
- Defining ownership and versioning
- Resource Pools: limit concurrency + prioritization
- Plugins
 - Operators
 - Executors
 - New Views
- Built-in integration with other services
- Many more...



Airflow Architecture

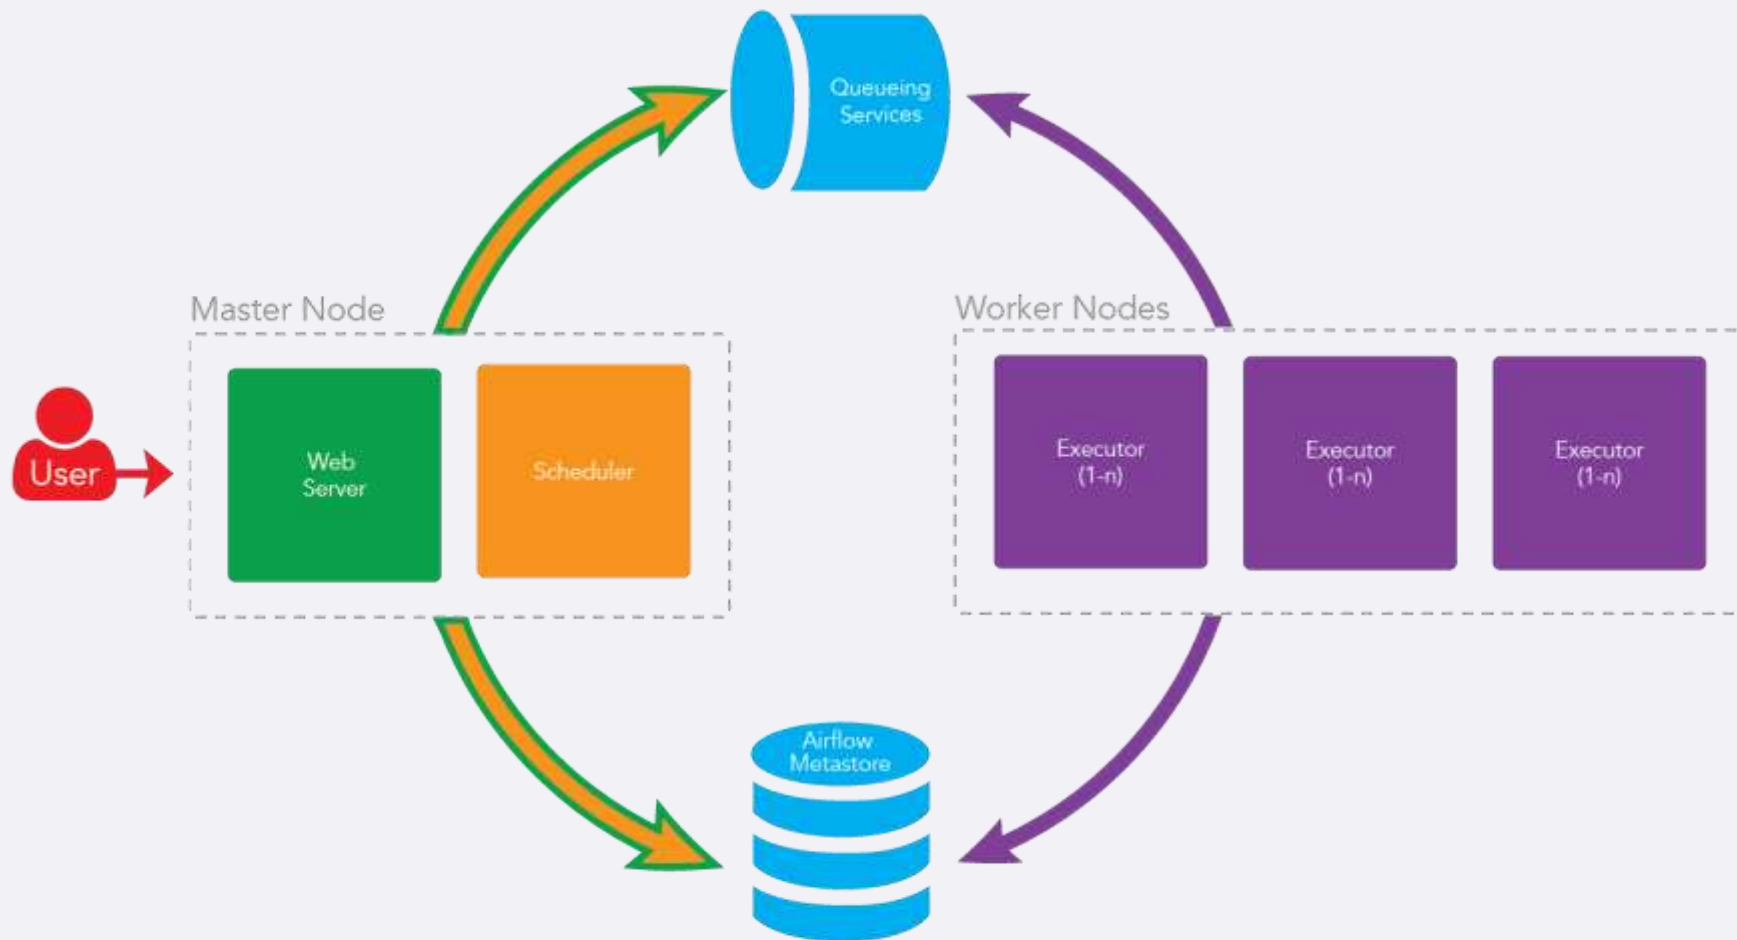
Single Node





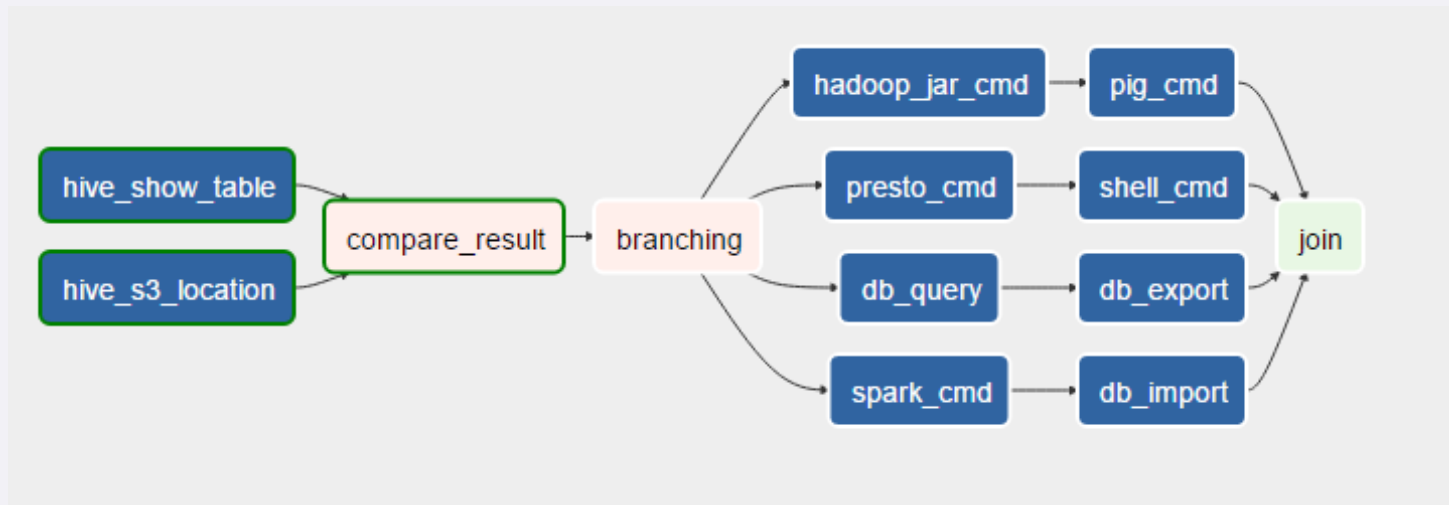
Airflow Architecture

Multi-node



What is a DAG?

- **Directed Acyclic Graph**
 - A finite directed graph that doesn't have any cycles
- A collection of tasks to run, organized in a way that reflects their relationships and dependencies
 - Defines your Workflow



What is an Operator?

- An operator describes a single task in a workflow
- Operators allow for generation of certain types of tasks that become nodes in the DAG when instantiated
- All operators derive from BaseOperator and inherit many attributes and methods that way

BashOperator DummyOperator EmailOperator PythonOperator

DummyOperator_Task → BashOperator_Task → PythonOperator_Task → EmailOperator_Task

Workflow Operators (Sensors)

- A type of operator that keeps running until a certain criteria is met
- Periodically pokes
- Parameterized poke interval and timeout
- Example
 - HdfsSensor
 - HivePartitionSensor
 - NamedHivePartitionSensor
 - S3KeyPartition
 - WebHdfsSensor
 - Many More...

Workflow Operators (Transfer)

- Operator that moves data from one system to another
- Data will be pulled from the source system, staged on the machine where the executor is running and then transferred to the target system
- Example:
 - HiveToMySqlTransfer
 - MySqlToHiveTransfer
 - HiveToMsSqlTransfer
 - MsSqlToHiveTransfer
 - S3ToHiveTransfer
 - Many More...

Defining a DAG


```
from airflow.models import DAG
from airflow.operators import ...
from datetime import datetime, timedelta

default_args = dict(
    'owner'='Airflow',
    'retries': 1,
    'retry_delay': timedelta(minutes=5),
)

# Define the DAG
dag = DAG('dag_id', default_args=default_args, schedule_interval='0 0 * * *')

# Define the Tasks
task1 = BashOperator(task_id='task1', bash_command="echo 'Task 1'", dag=dag)
task2 = BashOperator(task_id='task2', bash_command="echo 'Task 2'", dag=dag)
task3 = BashOperator(task_id='task3', bash_command="echo 'Task 3'", dag=dag)


# Define the task relationships
task1.set_downstream(task2)
task2.set_downstream(task3)
```



Defining a DAG (Dynamically)

```
dag = DAG('dag_id', default_args=default_args, schedule_interval='0 0 * * *')

last_task = None
for i in range(1, 3):
    task = BashOperator(
        task_id='task' + str(i),
        bash_command="echo 'Task' + str(i) + '",
        dag=dag)
    if last_task is None:
        last_task = task
    else:
        last_task.set_downstream(task)
        last_task = task
```



ETL Best Practices (Some of Them)

- Load Data Incrementally
 - Operators will receive an `execution_date` entry which you can use to pull in data since that date
- Process historic Data
 - Backfill operations are supported
- Enforce Idempotency (retry safe)
- Execute Conditionally
 - Branching, Joining
- Understand SLA's and Alerts
 - Alert if Failures
- Sense when to start a task
 - Sensor Operators
- Build Validation into your Workflows


Executing Airflow Workflows on Hadoop

- Airflow Workers should be installed on a edge/gateway nodes
 - Allows Airflow to interact with Hadoop related commands
 - Utilize the BashOperator to run command line functions and interact with Hadoop services
- Put all necessary scripts and Jars in HDFS and pull the files down from HDFS during the execution of the script
 - Avoids requiring you to keep copies of the scripts on every machine where the executors are running
- Support for Kerborized Clusters
 - Airflow can renew Kerberos tickets for itself and store it in the ticket cache

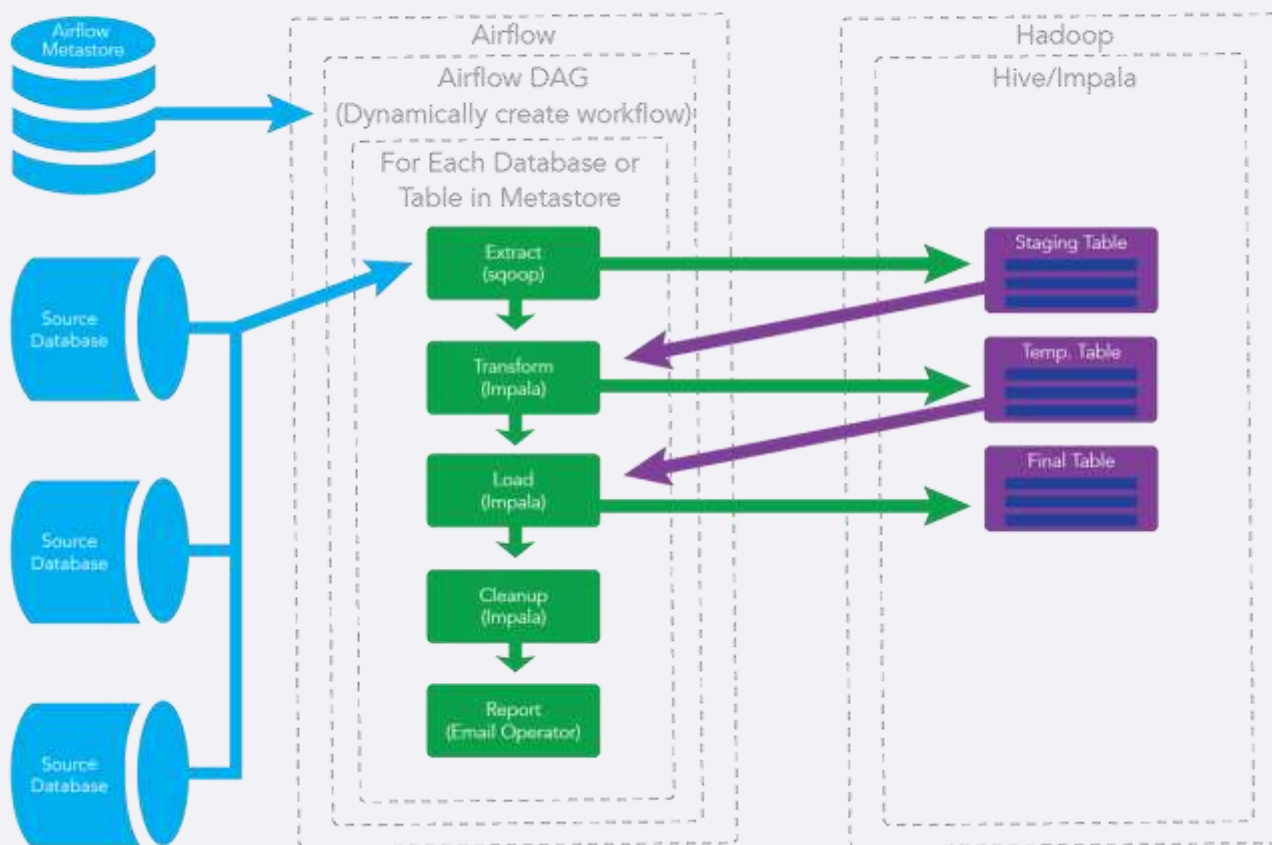
Use Case (BPV)

- Daily ETL Batch Process to Ingest data into Hadoop
 - Extract
 - 23 databases total
 - 1226 tables total
 - Transform
 - Impala scripts to join and transform data
 - Load
 - Impala scripts to load data into common final tables
- Other requirements
 - Make it extensible to allow the client to import more databases and tables in the future
 - Status emails to be sent out after daily job to report on success and failures
- Solution
 - Create a DAG that dynamically generates the workflow based off data

Use Case (BPV) (Architecture)




Airflow Architecture Use Case BPV



Use Case (BPV) (DAG)

100 foot view



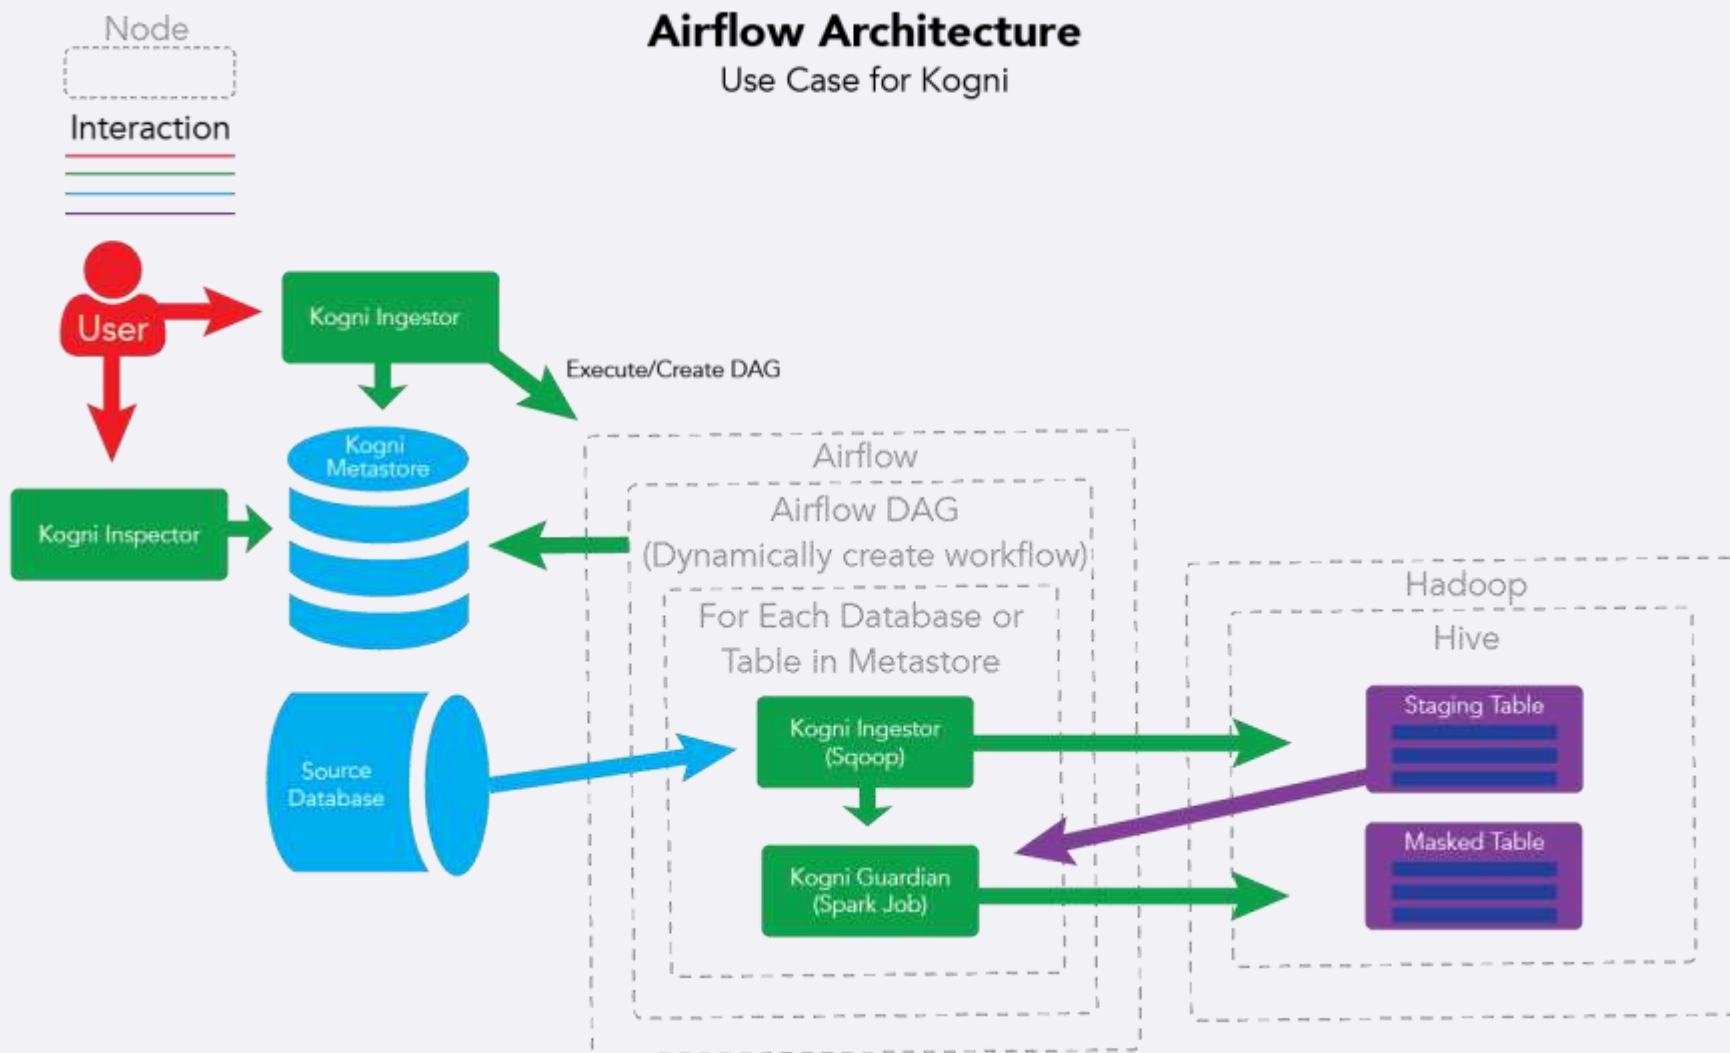
10,000 foot view



Use Case (Kogni)

- New Product being built by Clairvoyant to facilitate:
 - kogni-inspector – Sensitive Data Analyzer
 - kogni-ingestor – Ingests Data
 - kogni-guardian – Sensitive Data Masking (Encrypt and Tokenize)
 - Others components coming soon
- Utilizes Airflow for Data Ingestion and Masking
- Dynamically creates a workflow based off what is in the Metastore
- Learn More: <http://kogni.io/>

Use Case (Kogni) (Architecture)



References

- <https://pythonhosted.org/airflow/>
- <https://gtoonstra.github.io/etl-with-airflow/principles.html>
- <https://github.com/apache/incubator-airflow>
- <https://media.readthedocs.org/pdf/airflow/latest/airflow.pdf>



Q&A